Introduction to Database Systems IDBS – Spring 2024

- Week 1:
- Course Introduction
- DBMS Introduction
- Relational Data Model
- SQLDDL

Eleni Tzirita Zacharatou



PDBM 1, 6.1, 7.1-7.2

Course Responsible

Eleni Tzirita Zacharatou



Ph.D. in CS EPFL, Switzerland



MSc in Electrical & Computer Engineering NTUA, Greece



Visiting Researcher NYU, USA



2019 - 2022

2016

2013

Postdoctoral Researcher TUB, Germany

2013 - 2019



Assistant Professor ITU, Denmark

2022 - now

IT University of Copenhagen

Website: https://heltzi.github.io

Lecturer

Omar Shahbaz Khan

Postdoctoral Researcher

PhD in Computer Science

MSc in Computer Science

BSc in Software Development

2013

2016

2018

2022

Now











Teaching Assistants

Anders Arvesen

Study program: BSc. Software Development

Adam Hadou Temsamani

Study program: BSc. Software Development

Anne-Marie Rommerdahl

Study program: BSc. Software Development

Katrine Martos Sandø-Pedersen

Study program: BSc. Software Development

Oliver Flyckt Wilhjelm

Study program: BSc. Software Development

Erling Amundsen

Study program: BSc. Data Science

Frederik Rothe

Study program: MSc. Computer Science

Research Group

https://dasya.itu.dk

15+ People

Lab for Research and Education: 5th floor (5A56)

Entire Data Lifecycle:

- Collection
- Transfer
- Storage
- Curation
- Processing
- Analytics

Infrastructure for Data Science



Data-Intensive Systems and Applications

But This is About You!

- → ~ 145 Students

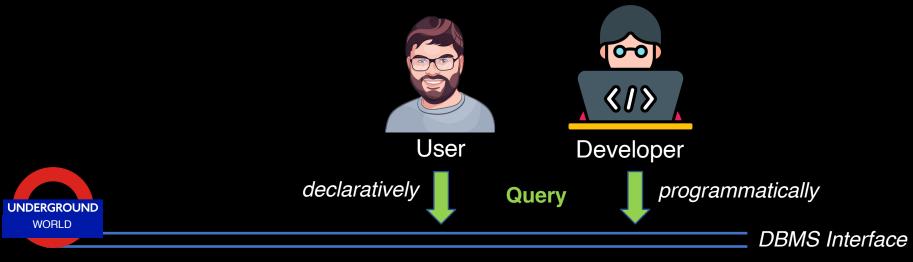
 Mostly from MSc. in Software Design
- Too many for a round table

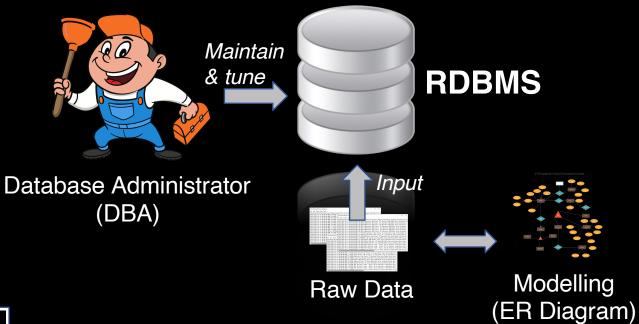
What do you expect from Introduction to Database Systems?

Mentimeter: https://www.menti.com/al8ndx4u3bpr



What You will Learn





Intented Learning Outcomes

- Write SQL queries: multiple relations; compound conditions; grouping; aggregation; and subqueries.
- Use relation
- Suggest a d in a suitable
- Analyze/pre indices.

- 1. Getting into Database Systems
- 2. Getting data using SQL and from your apps
- 3. Design a database
- 4. Tune a database
- 5. Advanced databases (internals and big data)

cure manner.

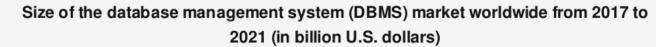
abase schema

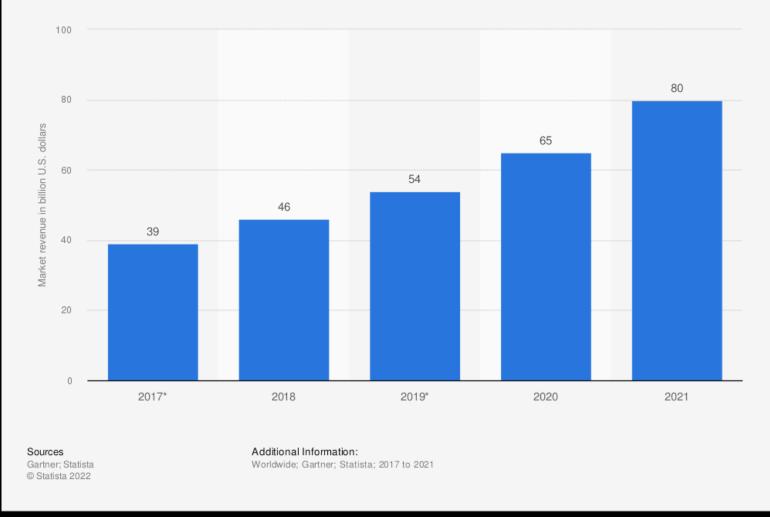
base using

- Reflect upon the evolution of the hardware and storage hierarchy and its impact on data management system design.
- Discuss the pros and cons of different classes of data systems for modern analytics and data science applications.

Why is It Important?

- Crucial to effectively manage and utilize data
- Help to maintain data integrity and security
- Ease app development





https://www.statista.com/statistics/724611/worldwide-database-market/

Course Schedule

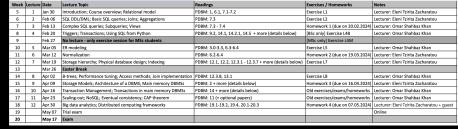
BSc vs MSc

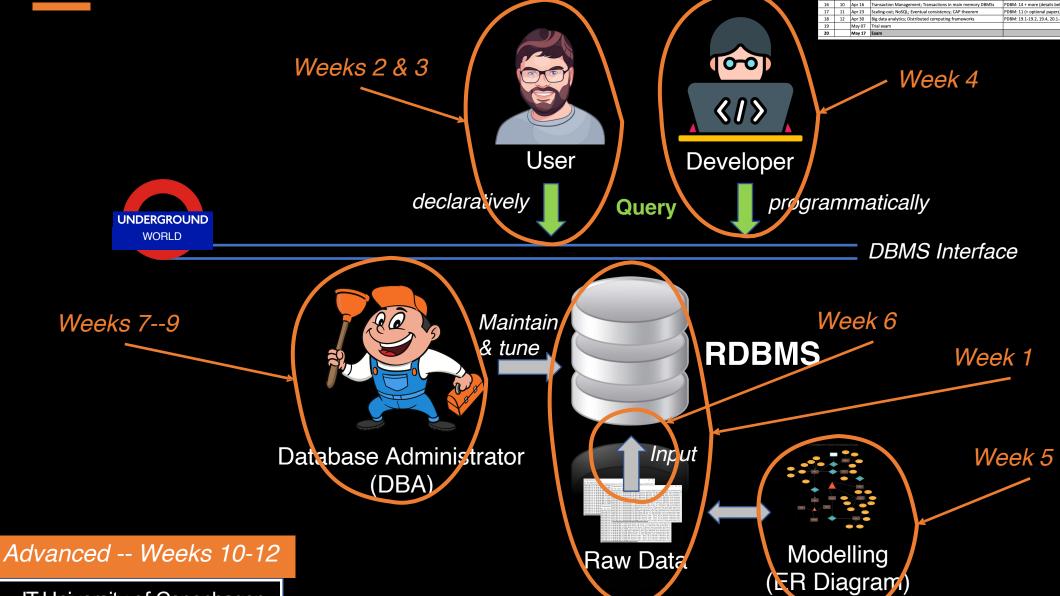
- Pretty much the same course
- One different question (5%) on exam

Week	Lecture	Date	Lecture Topic	Readings	Exercises / Homeworks	Notes
5	1	Jan 30	Introduction; Course overview; Relational model	PDBM: 1, 6.1, 7.1-7.2	Exercise L1	Lecturer: Eleni Tzirita Zacharatou
6	2	Feb 06	SQL DDL/DML; Basic SQL queries; Joins; Aggregations	PDBM: 7.3	Exercise L2	Lecturer: Eleni Tzirita Zacharatou
7	3	Feb 13	Complex SQL queries; Subqueries; Views	PDBM: 7.3 - 7.4	Homework 1 (due on 20.02.2024)	Lecturer: Omar Shahbaz Khan
8	4	Feb 20	Triggers; Transactions; Using SQL from Python	PDBM: 9.2, 14.1, 14.2.1, 14.5 + more (details below)	[BSc only] Exercise L4B	Lecturer: Omar Shahbaz Khan
9		Feb 27	No lecture - only exercise session for MSc students		[MSc only] Exercise L4M	
10	5	Mar 05	ER modeling	PDBM: 3.0-3.3, 6.3-6.4	Exercise L5	Lecturer: Omar Shahbaz Khan
11	6	Mar 12	Normalization	PDBM: 6.2-6.4	Homework 2 (due on 19.03.2024)	Lecturer: Eleni Tzirita Zacharatou
12	7	Mar 19	Storage hierarchy; Physical database design; Indexing	PDBM: 12.1, 12.2, 12.3.1 - 12.3.7 + more (details below)	Exercise L7	Lecturer: Eleni Tzirita Zacharatou
13		Mar 26	Easter Break			
14	8	Apr 02	B-trees; Performance tuning; Access methods; Join implementation	PDBM: 12.3.8, 13.1	Exercise L8	Lecturer: Omar Shahbaz Khan
15	9	Apr 09	Storage Models; Architecture of a DBMS; Main memory DBMSs	PDBM: 2 + more (details below)	Homework 3 (due on 16.04.2024)	Lecturer: Eleni Tzirita Zacharatou
16	10	Apr 16	Transaction Management; Transactions in main memory DBMSs	PDBM: 14 + more (details below)	Old exercises/exams/homeworks	Lecturer: Omar Shahbaz Khan
17	11	Apr 23	Scaling-out; NoSQL; Eventual consistency; CAP theorem	PDBM: 11 (+ optional papers)	Old exercises/exams/homeworks	Lecturer: Omar Shahbaz Khan
18	12	Apr 30	Big data analytics; Distributed computing frameworks	PDBM: 19.1-19.2, 19.4, 20.1-20.3	Homework 4 (due on 07.05.2024)	Lecturer: Eleni Tzirita Zacharatou + guest
19		May 07	Trial exam			Online
20		May 17	Exam			

Course Schedule (Illustrated)

IT University of Copenhagen





Course Structure

Lectures

- Tuesdays 8.15 10:00 at Aud 1
- Preparation required: reading material, watching videos

Exercises

- Tuesdays 10.15 12:00 at 3A54, 4A56, 4A58 (not this week)
- Preparation not required (related to previous lecture)

Homeworks

- 4 homeworks (deadlines published on learnIT)
- Mandatory (3 out of 4), yet easy to get accepted!
- Feedback in the following weeks (if submitted on time)
- LearnIT: course outline, materials, announcements, ...
- Piazza: Q&A, messages, and updates.
 - Ask consistently throughout the semester!
 - Help your peers!

Note

Exercise 1 Rooms per OS:

- Windows (3A54)
- Mac (4A56)
- Linux (4A56)
- Exceptionally this week: not using room 4A58

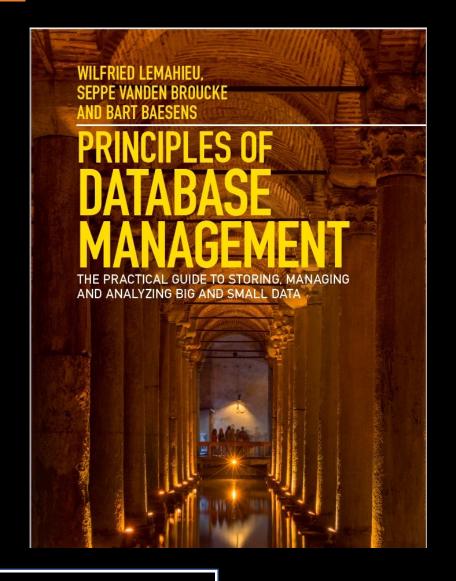
Course Methodology

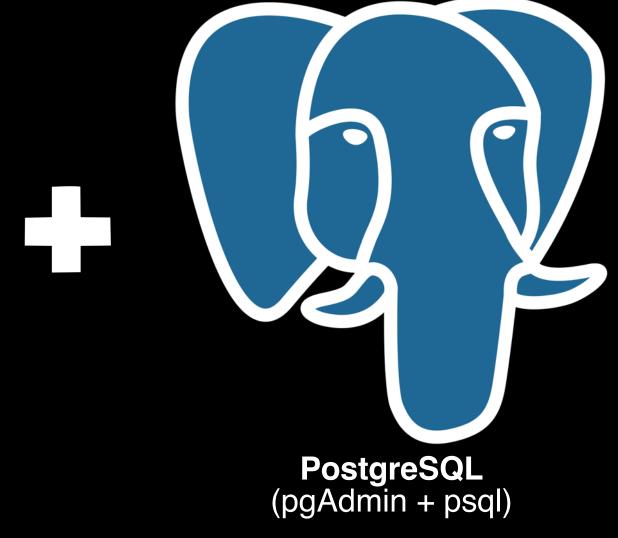
- You learn: We are here to help you!
- You need to read the book beforehand
 - Yes, we often assume you have done so
 - All readings are in the schedule on learnIT
 - In some weeks, there can be some video recording for you to watch before the lecture
- We work in a pull model fashion: ask questions!

Advice

Prepare and Ask Questions

Book and Database System





How will We Assess your Learning?

- 100% Exam (Quiz on LearnIT)
 - Restricted: no Internet access!
 - All course materials are allowed offline
 - Communication is not allowed
- Exercises and homeworks will help you prepare!

Advice

Study the material weekly



Edgar F. Codd

Father of Databases (Relational Model)

1923: Born 23/8, Isle of Portland, England

1965: PhD in CS from University of Michigan

1967: Moved to IBM Almaden Research Center

1969: Invented the relational model

1976: IBM Fellow

1981: Turing Award

1994: ACM Fellow

Information Retrieval

A Relational Model of Data for Large Shared Data Banks

E. F. CODD IBM Research Laboratory, San Jose, California P. BAXENDALE, Editor

The relational view (or model) of data described in Section 1 appears to be superior in several respects to the graph or network model [3, 4] presently in vogue for noninferential systems. It provides a means of describing data with its natural structure only—that is, without superimposing any additional structure for machine representation purposes. Accordingly, it provides a basis for a high level data language which will yield maximal independence between programs on the one hand and machine representa-

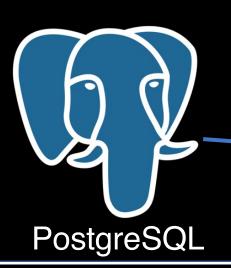
IT University of Copenhagen





Three-Layer Applications





IT University of Copenhagen

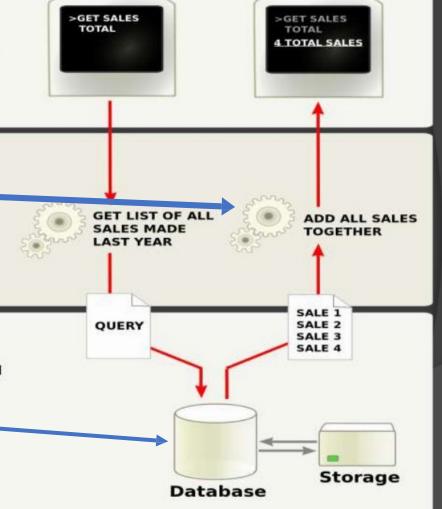
Presentation tier The top-most level of the application is the user interface. The main function of the interface is to translate tasks and results to something the user can understand.

Logic tier

application, processes commands, makes logical decisions and evaluations, and performs calculations. It also moves and processes data between the two surrounding layers.

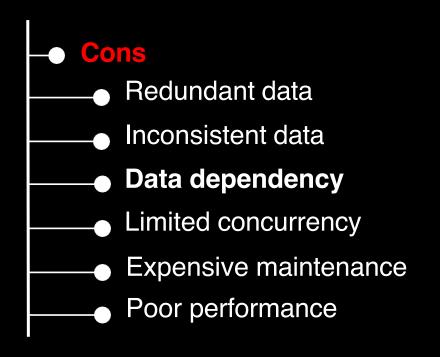
Data tier

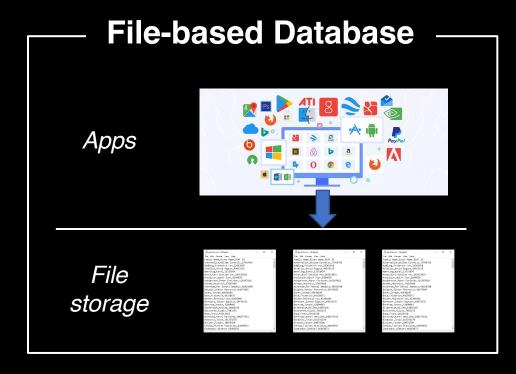
Here information is stored and retrieved from a database or file system. The information is then passed back to the logic tier for processing, and then eventually back to the user.



Database Definition

- A database is a collection of related data items within a specific business process or problem setting
- A database system provides a way to systematically organize, store, retrieve, and query a database

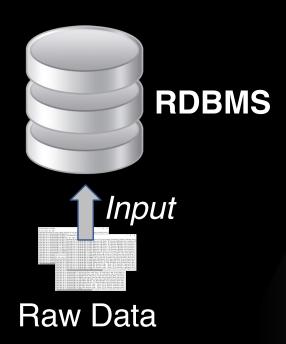




Relational Database

- A relational database is a type of database that is based on the relational model
 - stores data in a set of tables with rows and columns (a.k.a. relations)
 - uses relationships between these tables to manage the data.
- → A relational database system (RDBMS) implements and manages relational databases.

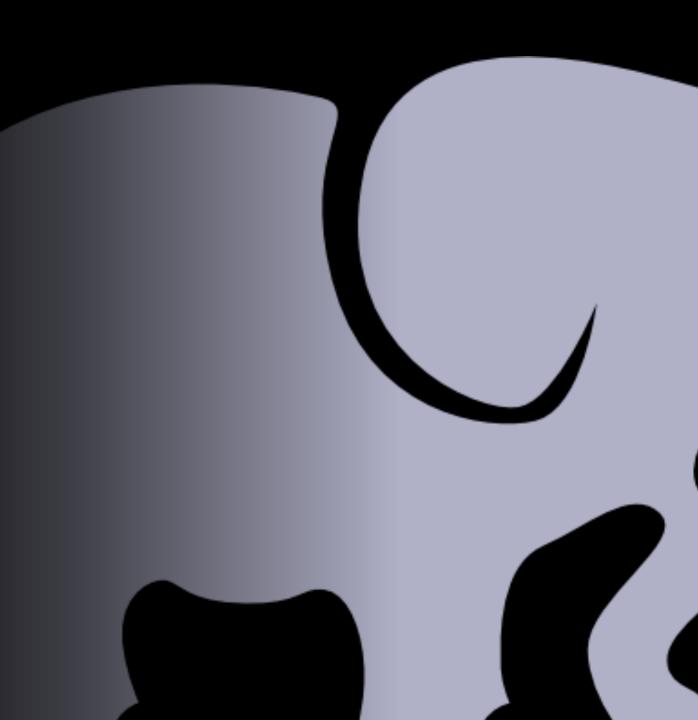




Relational Model

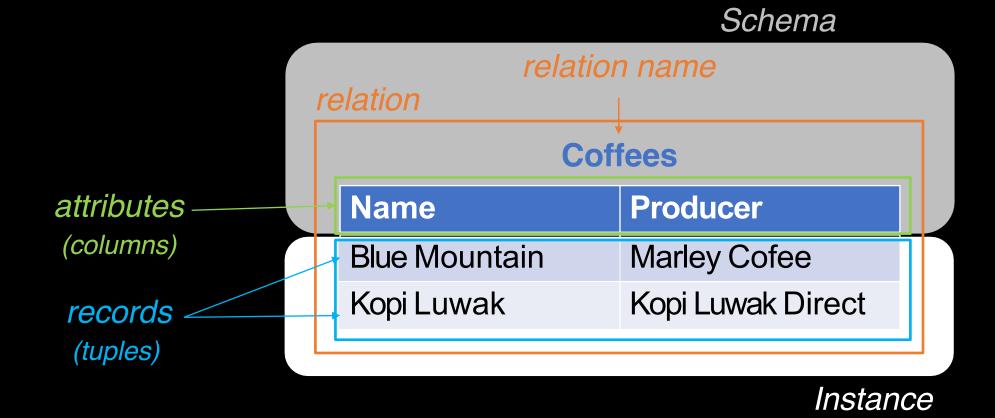


PDBM 1, 6.1



Basic Concepts

Relation



IT University of Copenhagen

Basic Concepts

Schema vs Instance vs Database

- Relation
 - Schema = name + list of attributes
 - Optional: attribute types
 - Coffees (Name, Producer)
 - Coffees (Name:STRING, Producer:String) -- specifies the domain
 - Instance
 - Records in a relation
 - E.g., (Blue Mountain, Marley Coffee)
- Database = collection of relations
 - Database schema = set of all relations names in the database
 - Database instance = set of all relations instances in the database

Basic Concepts

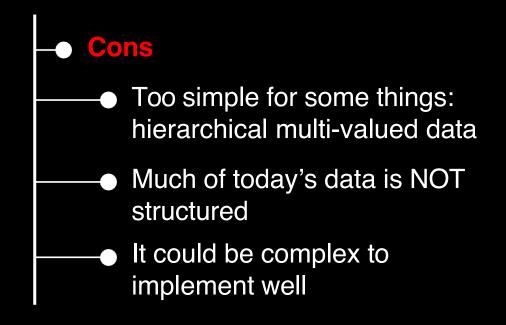
Example of a Database Schema

- Students (Sld:INT, Name:STRING, Email:STRING, Semester:INT)
- Faculty (Fld:INT, Name:STRING, Dld:INT)
- Courses (Cld:STRING, Name:STRING, Dld:INT)
- Departments (Dld:INT, Name:STRING)
- Transcripts (Cld:STRING, Sld:INT, Grade:STRING, Comment:STRING)

IT University of Copenhagen

Why Relations?

Pros
 Very simple model
 How we typically think about structured data
 Conceptual model behind SQL, which is the most important query language today



Identifiers

identifier

- Students (SId:INT, Name:STRING, Email:STRING, Semester:INT)
- Faculty (Fld:INT, Name:STRING, Dld:INT)
- Courses (Cld:STRING, Name:STRING, Dld:INT)
- Departments (Dld:INT, Name:STRING)
- Transcripts (Cld:STRING, Sld:INT, Grade:STRING, Comment:STRING)

IT University of Copenhagen

Keys and Superkeys

What is a key?

- Defines unique records (instances)
- Helps in setting relationships between relations
- Ensures the mathematical definition of a relation (set of records)

Superkeys

- Is a set of attributes that uniquely identify records: Uniqueness property
- The entire set of attributes of a relation is a superkey
- Minimal superkey: Minimality property
 - No attribute can be removed from a superkey without violating the uniqueness property

superkey		key	Students	
	Sld	Name	Emai	Semester
	01785	Bob Brown	bobr@itu.dk	2
	01615	Lucas White	luwh@itu.dk	5

Candidate Keys

- Attributes that satisfies the uniqueness and minimality properties
 - Minimal superkey = (candidate) key
 - Superkeys contains at least one (candidate) key
 - A relation can have many (candidate) keys

superkey	key Students		Students	
	Sld	Name	Emai	Semester
	01785	Bob Brown	bobr@itu.dk	2
	01615	Lucas White	luwh@itu.dk	5

Primary Keys

- A key to identify records in a relation
 - Important to define indexes and for storage purposes (later in the course)
 - Cannot be NULL
 - Also used to establish relationships with other relations
 - From all candidate keys only one can be primary key
 - The remaining ones are known as Alternative Keys

superkey	key		Students	
	Sld	Name	Emai	Semester
	01785	Bob Brown	bobr@itu.dk	2
	01615	Lucas White	luwh@itu.dk	5

Live Exercise **Students**

What are superkeys and keys?

(SId)

-- (Candidate) Key

(Email)

-- (Candidate) Key

(Sld, Name)

-- Superkey

(Semester)

-- None

(Email, Semester) -- Superkey

(Name)

-- None

(Name, Semester) -- (Candidate) Key

Sld	Name	Email	Semester
01785	Bob Brown	bobr@itu.dk	2
01615	Lucas White	luwh@itu.dk	5
01436	Olga Marx	olma@itu.dk	6
01875	Jens Schuh	jesc@itu.dk	1
01803	Olga Marx	olmr@itu.dk	2
01567	Peter Pitt	pepi@itu.dk	1

- What is the best key for being the primary key? (Sld) is ideal for being the Primary Key
- Which of these keys (does) not make sense in practice? (Name, Semester)



Relationships

- Students (Sld:INT, Name:STRING, Email:STRING, Semester:INT)
- Faculty (Fld:INT, Name:STRING, Dld:INT)
- Courses (Cld:STRING, Name:STRING, Dld:INT)
- Departments (Dld:INT, Name:STRING)

Transcripts (Cld:STRING, Sld:INT, Grade:STRING, Comment:STRING)

IT University of Copenhagen

relationship

Foreign Keys

- Defines the relationship between relations
- A key FK in a relation R is a foreign key iff:
 - The attributes in FK matches a primary key PK of a relation S and they are of the same type
 - Any record i in R has a value in FK that either
 - occurs as a value of PK for some tuple j in S, or
 - is null
 - I.e., FK = PK (domain and values)
- A relation can have several foreign keys

Foreign Keys -- Example

PK Departments



)FK

PK	Call	4000
	Cou	1565

Cld		Name	Dld
MATI	4	Mathematics	2
12DB	3	Intro. To DB Syst.	2
SWE		Soft. Engineering	4

PK		Students		ey Relation
S	ld	Name	Email	Semester
01	785	Bob Brown	bobr@itu.dk	2
01	615	Lucas White	luwh@itu.dk	5
01	436	Olga Marx	olma@itu.dk	6
닖			•••	

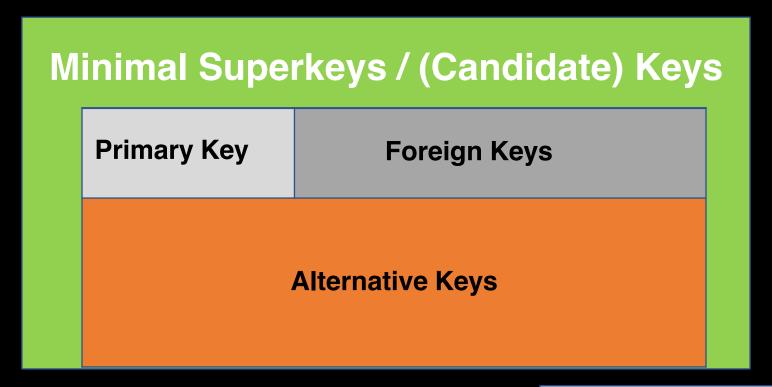
Transcripts

Foreign Relation

Cld	Sld	Grade	Comment
I2DBS	01785	7	The student didn't
MATH	01436	12	She was

All Keys in a Relation

Superkeys

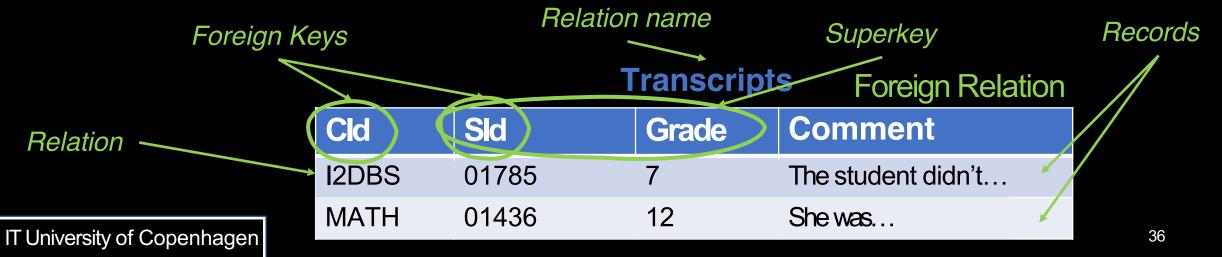


Note

Keys are part of the schema

So Far, All Together

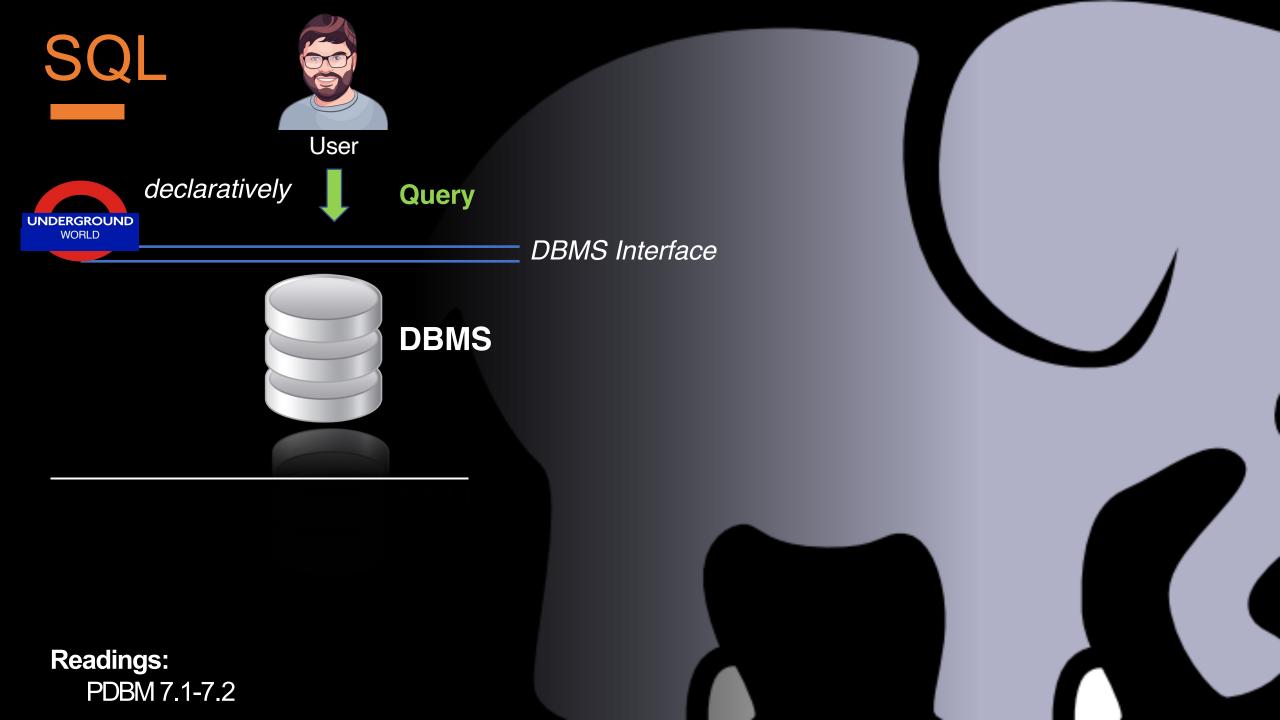




Integrity Constraints

- An integrity constraint (IC) is a limitation of the allowed content (or development) of a database
 - Ensures that the data are always correct and consistent
 - There exist various ICs
- It is the RDBMS that takes care of ensuring the ICs in a database
- ICs already seen so far
 - Domain constraint
 - Key constraint
 - Entity constraint (PK) -- NOT NULL
 - Referential constraint (FK) -- PK = FK
- -- attribute type and format (e.g., DATE)
- -- uniqueness & minimality

- More advanced ICs
 - Functional dependencies
 - Temporal constraint...



Structured Query Language (SQL)

- SEQUEL if you worked for IBM in the 80s
- SQL is primarily a query language, for getting information from a database (DML)
 - also includes a data-definition component for describing database schemas (DDL)
- Invented in the 70s by IBM
- The three most common commands in SQL queries
 - SELECT, FROM, WHERE

```
SELECT * FROM Students WHERE Name = 'Lucas White';
```

History [edit]

SQL was initially developed at IBM by Donald D. Chamberlin and Raymond F. Boyce in the early 1970s.^[14] This version, initially called *SEQUEL* (*Structured English Query Language*), was designed to manipulate and retrieve data stored in IBM's original quasi-relational database management system, System R, which a group at IBM San Jose Research Laboratory had developed during the 1970s.^[14] The acronym SEQUEL was later changed to SQL because "SEQUEL" was a trademark of the UK-based Hawker Siddeley aircraft company.^[15]



Two in One Language

- Data Definition Language (DDL)
 - Used by the database administrator (DBA) to define the database's data model
 - Three common commands:
 - CREATE TABLE, ALTER TABLE, and DROP TABLE

Today's focus

- Data Manipulation Language (DML)
 - Used by applications and users to retrieve, insert, modify, and delete records
 - Four statements:
 - SELECT, INSERT, UPDATE, and DELETE



First Normal Form (1NF)

- Each attribute in a relation has:
 - a primitive type (atomic values), and;
 - a unique name
- The main goal is to eliminate redundant data in a relation
- Benefits:
 - Data integrity
 - Data consistency
 - Easy data manipulation
 - Better data organization

SQL -- DDL

Data Types

Type	Description
CHAR(n)	Fixed-length string of size n
VARCHAR(n)	Variable-length string of maximum size n
SMALLINT	Small integer (-32,768 and 32,767)
INT	Integer (-2,147,483,648 and 2,147,483,647)
FLOTAT(n, d)	Small number with a floating decimal point: n = max digits and d = max decimals
DOUBLE(n, d)	Large number with a floating decimal point: n = max digits and d = max decimals
DATE	Date in format YYYY-MM-DD
DATETIME	Date and time in format YYYY-MM-DD HH:MI:SS
TIME	Time in format HH:MI:SS
BOOLEAN	True or false
BLOB	Binary large object (typically unstructured)

Note

Check types in PostgreSQL



Create Relation Students

```
CREATE TABLE Students (
   SId INT,
   Name VARCHAR(255),
   Email CHAR(11),
   Semester INT
);
```

Students

Sld	Name	Email	Semester

Advice

Start playing with PostgreSQL ASAP!



Define a Primary Key

```
sql
                                               ALTER TABLE Students
                                               ADD PRIMARY KEY (SId);
                                               CREATE TABLE Students (
                                                  SId INT PRIMARY KEY
                                                  Name VARCHAR(255),
                                                  Email CHAR(11),
                                                  Semester INT
                  Students
                                               );
Sld
        Name
                        Email
                                             Semester
```

IT University of Copenhagen



Define a Foreign Key

Students

SId	Name	Email	Semester

```
CREATE TABLE Departments (
DId INT PRIMARY KEY,
Name VARCHAR(255)
);
Departments

DId Name
```

```
CREATE TABLE Courses (
   CId VARCHAR(25) PRIMARY KEY,
   Name VARCHAR(255),
   DId INT,
   FOREIGN KEY (DId) REFERENCES Departments (DId) Courses
);

Cld Name
```



Multiple-Attribute Primary Keys

Students

Name	Email	Semester

Cld	Name	Dld

Courses

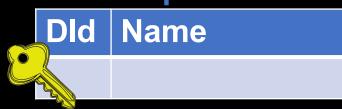
Cld Grade Comment

CREATE TABLE Transcripts (

CId VARCHAR(25),

SId INT,

Departments



Grade VARCHAR(10), Comment VARCHAR(255) PRIMARY KEY (CId, SId), FOREIGN KEY (CId) REFERENCES Courses (CId), FOREIGN KEY (SId) REFERENCES Students (SId) 47);

Transcripts

SQL -- DDL

NULL

- What if a value is missing?
 - Does not exist?
 - Unknown?
 - Secret?
- SQL solution: NULL
 - NULL = no value
 - More next week...
- By default, attributes can be NULL
 - Except: PRIMARY KEY attributes
 - Except: NOT NULL attributes
- Allowing NULL values is a design decision!

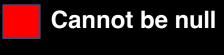
SQL -- DDL

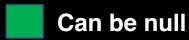
Relation Students with Not Nulls

```
sal
CREATE TABLE Transcripts (
  CId VARCHAR(255) NOT NULL,
  SId INT NOT NULL,
  Grade VARCHAR(10) NOT NULL,
  Comment VARCHAR(255),
  PRIMARY KEY (CId, SId),
  FOREIGN KEY (CId) REFERENCES Courses(CId),
  FOREIGN KEY (SId) REFERENCES Students(SId)
);
```











Drop All Created Relations

```
DROP TABLE Transcripts;
DROP TABLE Courses;
DROP TABLE Students;
DROP TABLE Departments;
```

IT University of Copenhagen



Takeaways

- Relational model
 - Relations, attributes, keys, primary & foreign keys, ...
- SQL DDL = Data Definition Language
 - CREATE TABLE, DROP TABLE, ALTER TABLE, ...
 - Allows to create complex schemas and maintain them
- SQL DML = Data Manipulation Language
 - INSERT, DELETE, UPDATE, SELECT
 - Simple set of commands for complicated actions
 - (we will dive into it next week)



What is next?

- Next week: SQL DML
 - SQL DML basics, joins, aggregates, grouping, ...
- Install PostgreSQL
 - ... if you have not already done so
 - Problems: get help during the exercise session!
- Exercises today
 - 1. Scripts to start playing on LearnIT (DB install script and queries)
 - 2. Create a sample "Coffee" database
 - Use commands from slides provided on LearnIT
 - Write some INSERT and SELECT statements
 - Play with constraints
 - 3. Consider databases without a DBMS