

Introduction to Database Systems

IDBS - Spring 2024

Week 3 - Advanced SQL

Division

JOINS and NULL

Set Operations

Subqueries

Views

Readings: PDBM 7.3-7.4

Omar Shahbaz Khan

WHOAMI

- Name: Omar Shahbaz Khan
- Background:
 - 2013 - 2016: BSc, ITU (SWU)
 - 2016 - 2018: MSc, ITU (CS)
 - 2018 - 2022: PhD, ITU, Large Scale Interactive Learning for Multimedia
 - 2023 - Now: Postdoc, Reykjavik University, Multimedia Exploration in VR
- Research Interests:
 - Multimedia Analytics, Interactive Learning, Scalability, Databases, Big Data Management and Analytics

Running Example 1

Same schema as in week 2

Script on LearnIT: [03-coffees-schema/queries.sql](#)

Coffees (name, manufacturer)

Coffeehouses (name, address, license)

Drinkers (name, address, phone)

Likes (drinker, coffee)

Sells (coffeehouses, coffee, price)

Frequents (drinker, coffeehouse)

Running Example 2

Sports Database same schema as in week 2 exercises

Script on LearnIT: [03-sports-schema/queries.sql](#)

People (ID, name, gender, height)

Gender (gender, description)

Sports (ID, name, record)

Competitions (ID, place, held)

Results (peopleID, competitionID, sportID, result)

Last Time in IDBS...

Last Time...

- SELECT - FROM - WHERE
- JOIN
- ORDER BY, DISTINCT, Aggregations
- GROUP BY + HAVING

Last Time...

- SELECT - FROM - WHERE

```
SELECT name  
FROM Coffees;
```

- JOIN

```
SELECT name AS coffee  
FROM Coffees  
WHERE manufacturer LIKE 'Kopi%';
```

- ORDER BY, DISTINCT, Aggregations

- GROUP BY + HAVING

Last Time...

- SELECT - FROM - WHERE

- JOIN

- ORDER BY, DISTINCT, Aggregations

- GROUP BY + HAVING

```
SELECT f.drinker,  
       f.coffeehouse,  
       ch.address  
FROM Frequents f  
JOIN Coffeehouses ch ON  
       f.coffeehouse = ch.name;
```


Last Time...

- SELECT - FROM - WHERE

```
SELECT *  
  FROM Sells  
 ORDER BY price ASC;
```

- JOIN

```
SELECT COUNT(*)  
  FROM Sells;
```

- ORDER BY, DISTINCT, Aggregations

```
SELECT COUNT(DISTINCT price)  
  FROM Sells;
```

- GROUP BY + HAVING

Last Time...

- SELECT - FROM - WHERE

```
SELECT coffee, MAX(price)
  FROM Sells
 GROUP BY coffee;
```

- JOIN

- ORDER BY, DISTINCT, Aggregations

```
SELECT coffee, AVG(price)
  FROM Sells
 GROUP BY coffee
HAVING COUNT(coffeehouse) > 1;
```

- GROUP BY + HAVING

Wake Up Task!

Sports DB

People (ID, name, gender, height)

Gender (gender, description)

Sports (ID, name, record)

Competitions (ID, place, held)

Results (peopleID, competitionID,
sportID, result)

How many sports have the word
“Jump” in them?

How many people have participated
in more than 25 competitions?

This Time...

-- TODO

- Division
- JOINS & NULL
 - Natural Joins, Cross Joins, Self-Joins
 - NULL
 - Three-valued logic
 - Operators, Outer Joins
- Set Operations
 - UNION, INTERSECT, EXCEPT
- Subqueries (Nested Queries)
 - =, IN, EXISTS, ALL, ANY
- Views: Queries as subroutines



What is Division?

- $R1 / R2$ = tuples of $R1$ associated with all tuples of $R2$
- Find the students who have taken all courses in a program
- Find the airlines who land at all airports in a country/continent/the world
- Find the coffeehouses that sell all existing coffees

Division with Counting

- We can write *division* using GROUP BY, HAVING and a COUNT subquery
- Find the coffeehouses that sell all existing coffees

Coffees (name, manufacturer)

Coffeehouses (name, address, license)

Drinkers (name, address, phone)

Likes (drinker, coffee)

Sells (coffeehouses, coffee, price)

Frequents (drinker, coffeehouse)

```
SELECT coffeehouse
FROM Sells
GROUP BY coffeehouse
HAVING COUNT(coffee) = (
```

2. For each
coffeehouse,
return it only
if it sells that
many

```
SELECT COUNT(*)
FROM Coffees);
```

1. Count the number of coffees

Practice

- Names of drinkers who frequent all coffeehouses

Coffees (name, manufacturer)

Coffeehouses (name, address, license)

Drinkers (name, address, phone)

Likes (drinker, coffee)

Sells (coffeehouses, coffee, price)

Frequents (drinker, coffeehouse)

```
SELECT drinker
  FROM Frequents
 GROUP BY drinker
HAVING COUNT(coffeehouse) = (
      SELECT COUNT(*)
      FROM Coffeehouses);
```

More Complete/Complex Example

- Exercise 2, #20

Show the ID, name, record and worst results of all sports that have at least one result from every place where a competition has ever been held

People (ID, name, gender, height)

Gender (gender, description)

Sports (ID, name, record)

Competitions (ID, place, held)

Results (peopleID, competitionID, sportID, result)

```
SELECT s.id, s.name, s.record, MIN(r.result)
FROM Sports s
JOIN Results r ON s.id = r.sportId
JOIN Competitions c ON
    r.competitionId = c.id
GROUP BY s.id, s.name, s.record
HAVING COUNT(DISTINCT c.place) = (
    SELECT COUNT(DISTINCT c.place)
    FROM Competitions c);
```

-- TODO

✓ Division

- JOINS & NULL

- Natural Joins, Cross Joins, Self-Joins
- NULL
- Three-valued logic
- Operators, Outer Joins


- Set Operations

- UNION, INTERSECT, EXCEPT

- Subqueries (Nested Queries)


- =, IN, EXISTS, ALL, ANY

- Views: Queries as subroutines



INNER JOIN

- Most joins are inner joins
- Combines the data of two relations based on a common column
- This is what we did last week
 - Following are some variants of INNER Joins



```
SELECT *  
FROM Drinkers d  
JOIN Frequent f ON  
d.name = f.drinker;
```

Natural Joins and Cross Joins (Product)

A NATURAL JOIN B:

- Assumes “=” on columns of same name
 - Removes duplicates
- Can be dangerous, not recommended!

```
SELECT *  
FROM Frequent  
NATURAL JOIN Likes;
```

A CROSS JOIN B:

- Removes no columns, has no join condition
 - Can simulate with an always true condition (ON 1 = 1)
 - Rarely (but not never) the right thing to do!

```
SELECT *  
FROM Likes l  
CROSS JOIN Drinkers d;
```

Self-Join

- From **Coffees(name, manufacturer)**, find all pairs of coffees by the same manufacturer
 - Need a double loop to compare two coffee records
 - WHERE clause ensures no pairs with the same coffees, and alphabetical order

```
SELECT c1.name, c2.name,  
       c1.manufacturer  
FROM Coffees c1  
JOIN Coffees c2 ON  
       c1.manufacturer =  
       c2.manufacturer  
WHERE c1.name < c2.name;
```

Self-Join: Example

- Show all coffees that are more expensive than some other coffee sold at the same coffeehouse

Coffees (name, manufacturer)
Coffeehouses (name, address, license)
Drinkers (name, address, phone)
Likes (drinker, coffee)
Sells (coffeehouses, coffee, price)
Frequents (drinker, coffeehouse)

1. Get all pairs:

```
SELECT *  
FROM Sells s1  
JOIN Sells s2  
ON s1.coffeehouse =  
s2.coffeehouse;
```

2. Remove unwanted pairs:

```
SELECT *  
FROM Sells s1  
JOIN Sells s2  
ON s1.coffeehouse =  
s2.coffeehouse  
WHERE s1.price > s2.price;
```

3. Remove duplicates
and show only coffees:

```
SELECT DISTINCT s1.coffee  
FROM Sells s1  
JOIN Sells s2  
ON s1.coffeehouse =  
s2.coffeehouse  
WHERE s1.price > s2.price;
```


NULL Values

- Tuples in SQL relations can have NULL as a value for one or more components
- Meaning depends on context. Two Common cases:
 - **Missing Value** - e.g., we know Joe's Bar has some address, but we don't know what it is
 - “secret”, “figure not available”, “TBA”, “impossible to calculate”, “partly unknown”, “uncertain”, “pending”
 - **Inapplicable** - e.g., the value attribute spouse for an unmarried person
 - “undefined”, “moot”, “irrelevant”, “none”, “N/A”

Comparing NULL Values

- The logic conditions in SQL is 3-valued
 - TRUE, FALSE, UNKNOWN
- Comparing any value (including NULL itself) with NULL yields UNKNOWN
 - `NULL != NULL`
 - Must use `IS NULL` or `IS NOT NULL`
- A tuple is in a query answer iff the `WHERE` clause is TRUE (not FALSE or UNKNOWN)

Example

- Show the coffeehouses that sell a coffee at an unknown price
- Show the name of all coffees that are sold at a known price

Coffees (name, manufacturer)

Coffeehouses (name, address, license)

Drinkers (name, address, phone)

Likes (drinker, coffee)

Sells (coffeehouses, coffee, price)

Frequents (drinker, coffeehouse)

```
SELECT DISTINCT coffeehouse  
  FROM Sells  
 WHERE price IS NULL;
```

```
SELECT DISTINCT coffee  
  FROM Sells  
 WHERE price IS NOT NULL;
```

Three-Valued Truth Tables

AND	TRUE	FALSE	NULL
TRUE	TRUE	FALSE	NULL
FALSE	FALSE	FALSE	FALSE
NULL	NULL	FALSE	NULL

NOT	TRUE	FALSE	NULL
	FALSE	TRUE	NULL

OR	TRUE	FALSE	NULL
TRUE	TRUE	TRUE	TRUE
FALSE	TRUE	FALSE	NULL
NULL	TRUE	NULL	NULL

=	TRUE	FALSE	NULL
TRUE	TRUE	FALSE	NULL
FALSE	FALSE	TRUE	NULL
NULL	NULL	NULL	NULL

NULL vs Aggregation


- NULL never contributes to a sum, average, or count, and can never be the minimum or maximum of a column
- But if there are no non-NULL values in a column then the result of the aggregation is NULL
 - **Exception:** COUNT of an empty set is 0

Outer Joins


- Sometimes a coffee is sold nowhere, but we want it in our results
- **A OUTER JOIN B** is the core of an outer join expression
- It is modified by LEFT, RIGHT, or FULL before OUTER
 - LEFT = pad dangling tuples of A with NULL
 - RIGHT = pad dangling tuples of B with NULL
 - FULL = pad both;

JOINS in a nutshell


LEFT OUTER
JOIN




RIGHT OUTER
JOIN



FULL OUTER
JOIN



INNER JOIN



Example of an Outer Join

Suppose we add a new coffee:

```
INSERT INTO Coffees VALUES  
( 'Bragakaffi', 'Ó.J. & Kaaber' );
```

- Show coffeehouses along with the manufacturers of the coffees they are selling. Also include manufacturers that do not sell any coffees.

Coffees (name, manufacturer)
Coffeehouses (name, address, license)
Drinkers (name, address, phone)
Likes (drinker, coffee)
Sells (coffeehouses, coffee, price)
Frequents (drinker, coffeehouse)

```
SELECT s.coffeehouse,  
       c.manufacturer  
FROM Sells s  
RIGHT OUTER JOIN Coffees c  
ON s.coffee = c.name;
```


Practice

- Show all drinkers and the coffees they like, but include drinkers that do not like any coffees
- Can you use OUTER JOIN to show ONLY drinkers who do not like any coffees?

Coffees (name, manufacturer)

Coffeehouses (name, address, license)

Drinkers (name, address, phone)

Likes (drinker, coffee)

Sells (coffeehouses, coffee, price)

Frequents (drinker, coffeehouse)

Practice

- Show all drinkers and the coffees they like, but include drinkers that do not like any coffees
- Can you use OUTER JOIN to show ONLY drinkers who do not like any coffees?

Coffees (name, manufacturer)

Coffeehouses (name, address, license)

Drinkers (name, address, phone)

Likes (drinker, coffee)

Sells (coffeehouses, coffee, price)

Frequents (drinker, coffeehouse)

```
SELECT *  
  FROM Likes l  
  JOIN Drinkers d ON d.name = l.drinker;
```

WRONG!

```
SELECT *  
  FROM Likes l  
 RIGHT OUTER JOIN Drinkers d  
   ON d.name = l.drinker;
```

Correct!

```
SELECT *  
  FROM Drinkers d  
 LEFT OUTER JOIN Likes l  
   ON l.drinker = d.name;
```

Correct!

Practice

- Show all drinkers and the coffees they like, but include drinkers that do not like any coffees
- Can you use OUTER JOIN to show ONLY drinkers who do not like any coffees?

```
SELECT d.name
FROM Drinkers d
LEFT JOIN Likes l
ON d.name = l.drinker
WHERE l.drinker IS NULL;
```

Coffees (name, manufacturer)

Coffeehouses (name, address, license)


Drinkers (name, address, phone)

Likes (drinker, coffee)


Sells (coffeehouses, coffee, price)

Frequents (drinker, coffeehouse)


SQL JOINS




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```




```
SELECT <select_list>
FROM TableA A
INNER JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
OR B.Key IS NULL
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
```



-- TODO

✓ Division

✓ JOIN & NULL

- ✓ Natural Joins, Cross Joins, Self-Joins
- ✓ NULL
- ✓ Three-valued logic
- ✓ Operators, Outer Joins

• Set Operations

- UNION, INTERSECT, EXCEPT

• Subqueries (Nested Queries)

- =, IN, EXISTS, ALL, ANY


• Views: Queries as subroutines

Set Queries


- Syntax:
 <Query 1>
 UNION/INTERSECT/EXCEPT
 <Query 2>
- Queries must be “union compatible” = results have matching schema:
 - Same number of attributes
 - Attributes i of both tables have same (matching type)

Set Operators in a nutshell


QA EXCEPT QB




QB EXCEPT QA



QA UNION QB
QB UNION QA



QA INTERSECT QB
QB INTERSECT QA



Operators

- **UNION**
Show all drinkers that like “Kopi luwak” or live in “Amager”
- **INTERSECT**
Show all coffees that are manufactured by “Marley Coffee” and sold at an unknown price
- **EXCEPT**
Show all drinkers that frequents “Analog” but not those that like the coffee “Kopi luwak”.

Operators

- **UNION**

Show all drinkers that like “Kopi luwak” or live in “Amager”

- **INTERSECT**

Show all coffees that are manufactured by “Marley Coffee” and sold at an unknown price

- **EXCEPT**

Show all drinkers that frequents “Analog” but not those that like the coffee “Kopi luwak”.

```
SELECT l.drinker
      FROM Likes l
     WHERE l.coffee = 'Kopi luwak'
UNION
SELECT d.name
      FROM Drinkers d
     WHERE d.address = 'Amager';
```

Operators

- **UNION**

Show all drinkers that like “Kopi luwak” or live in “Amager”

- **INTERSECT**

Show all coffees that are manufactured by “Marley Coffee” and sold at an unknown price

- **EXCEPT**

Show all drinkers that frequents “Analog” but not those that like the coffee “Kopi luwak”.

```
SELECT c.name
      FROM Coffees c
     WHERE c.manufacturer = 'Marley Coffee'
INTERSECT
SELECT s.coffee
      FROM Sells s
     WHERE s.price IS NULL;
```

Operators

- **UNION**

Show all drinkers that like “Kopi luwak” or live in “Amager”

- **INTERSECT**

Show all coffees that are manufactured by “Marley Coffee” and sold at an unknown price

- **EXCEPT**

Show all drinkers that frequents “Analog” but not those that like the coffee “Kopi luwak”.

```
SELECT f.drinker
      FROM Frequents f
      WHERE f.coffeehouse = 'Analog'
EXCEPT
SELECT l.drinker
      FROM Likes l
      WHERE l.coffee = 'Kopi luwak';
```

Duplicates

- Set operators remove duplicates
- To keep duplicates use:
 - $R \text{ UNION ALL } S$ - (0, 1, 2x5, 3x3, 4, 5)
 - $R \text{ INTERSECT ALL } S$ - (2x2, 1x3)
 - $R \text{ EXCEPT ALL } S$ - (1, 0x2, 1x3, 4)
- Most practical use:
UNION ALL when R and S are known to be disjoint!

<u>R</u>	<u>S</u>
1	0
2	2
2	2
3	2
3	3
4	5

-- TODO

✓ Division

✓ JOIN & NULL

✓ Natural Joins, Cross Joins,
Self-Joins

✓ NULL

✓ Three-valued logic

✓ Operators, Outer Joins

✓ Set Operations

✓ UNION, INTERSECT, EXCEPT

• Subqueries (Nested Queries)

◦ =, IN, EXISTS, ALL, ANY

• Views: Queries as subroutines

QUERY
QUERY
QUERY
QUERY
QUERY
QUERY

Subqueries

In SQL, a parenthesized SELECT-FROM-WHERE statement can be used as a value in many places, including WHERE and FROM clauses

```
SELECT *  
  FROM R  
 WHERE R.a <op>  
       (SELECT S.b  
        FROM S);
```

```
SELECT T.a  
  FROM (SELECT R.a, S.c  
        FROM R  
        JOIN S  
        ON R.a = S.b)  
 AS T;
```

```
SELECT R.a  
  FROM R  
 JOIN (SELECT S.b, S.c  
       FROM S) AS T  
 ON R.a = T.b;
```


Single-Tuple (=)

- Subqueries return relations, but...
If a subquery is guaranteed to produce one tuple, then it can be used as a value
 - Usually, the tuple has only one attribute
 - Implicit type cast
 - Run-time Error when >1 tuples or no tuples are returned
- Practical for:
 - Finding records with highest or lowest values
 - Division queries

```
SELECT coffee
FROM Sells
WHERE price =
  (SELECT MIN(price)
   FROM Sells);
```

```
SELECT s.coffeehouse
FROM Sells s
GROUP BY s.coffeehouse
HAVING COUNT(s.coffee) =
  (SELECT COUNT(*)
   FROM Coffees);
```

Practice

- Find the coffeehouses that serve some coffee for the same price
Mocha charges for Blue Mountain

Coffees (name, manufacturer)

Coffeehouses (name, address, license)

Drinkers (name, address, phone)

Likes (drinker, coffee)

Sells (coffeehouses, coffee, price)

Frequents (drinker, coffeehouse)

Practice

- Find the coffeehouses that serve some coffee for the same price Mocha charges for Blue Mountain
- Using the `Sells(coffeehouse, coffee, price)` relation, this can be done in two queries:
 - Find the price Mocha charges for Blue Mountain
 - Find the coffeehouses that serve a coffee at that price
- Can it be done in a single query?

```
SELECT price
  FROM Sells
 WHERE coffeehouse = 'Mocha'
    AND coffee = 'Blue Mountain';
```

```
SELECT coffeehouse
  FROM Sells
 WHERE price = 300;
```

Practice

- Find the coffeehouses that serve some coffee for the same price Mocha charges for Blue Mountain
- Using the `Sells(coffeehouse, coffee, price)` relation, this can be done in two queries:
 - Find the price Mocha charges for Blue Mountain
 - Find the coffeehouses that serve a coffee at that price
- Can it be done in a single query?

```
SELECT coffeehouse
FROM Sells
WHERE price =
  (SELECT price
   FROM Sells
   WHERE coffeehouse = 'Mocha'
   AND coffee = 'Blue Mountain');
```

IN Operator

- Used to check if a tuple is IN a subquery's produced relation
 - The opposite can be achieved using NOT IN
- IN expressions can appear in WHERE clauses
- Example:
Using Coffees(name, manufacturer) and Likes(drinker, coffee), find the name and manufacturer of each coffee that Johan likes

```
SELECT *  
  FROM Coffees  
 WHERE name IN (  
     SELECT coffee  
     FROM Likes  
     WHERE drinker LIKE 'Johan%');
```

Example: NOT IN

- Show the name of all drinkers that does not frequent any coffeehouses

```
SELECT d.name
  FROM Drinkers d
 WHERE d.name NOT IN (
      SELECT f.drinker
      FROM Frequents f);
```

Set Operators: UNION

- Show all drinkers that like “Kopi luwak” or live in “Amager”

```
SELECT l.drinker
  FROM Likes l
 WHERE l.coffee = 'Kopi Luwak'
UNION
SELECT d.name
  FROM Drinkers d
 WHERE d.address = 'Amager';
```

```
SELECT d.name
  FROM Drinkers d
 WHERE d.address = 'Amager'
    OR d.name IN (
    SELECT l.drinker
      FROM Likes l
     WHERE l.coffee = 'Kopi Luwak');
```

Set Operators: INTERSECT

- Show all coffees that are manufactured by “Marley Coffee” and sold at an unknown price

```
SELECT c.name
  FROM Coffees c
 WHERE c.manufacturer = 'Marley Coffee'
INTERSECT
SELECT s.coffee
  FROM Sells s
 WHERE s.price IS NULL;
```

```
SELECT c.name
  FROM Coffees c
 WHERE c.manufacturer = 'Marley Coffee'
    AND c.name IN (
    SELECT s.coffee
      FROM Sells s
     WHERE s.price IS NULL);
```


Set Operators: EXCEPT

- Show all drinkers that frequents “Analog” but not those that like the coffee “Kopi luwak”

```
SELECT f.drinker
  FROM Frequents f
 WHERE f.coffeehouse = 'Analog'
EXCEPT
SELECT l.drinker
  FROM Likes l
 WHERE l.coffee = 'Kopi luwak';
```

```
SELECT f.drinker
  FROM Frequents f
 WHERE f.coffeehouse = 'Analog'
 AND f.drinker NOT IN (
  SELECT l.drinker
    FROM Likes l
   WHERE l.coffee = 'Kopi luwak');
```

EXISTS Operator

- **EXISTS** (<subquery>)
 - Returns true if and only if a query result is not empty
- Example:
Show the names of coffees that are liked by someone
 - Then show the opposite

```
SELECT c.name
  FROM Coffees c
 WHERE EXISTS (
    SELECT *
    FROM Likes l
    WHERE l.coffee = c.name);
```

```
SELECT c.name
  FROM Coffees c
 WHERE NOT EXISTS (
    SELECT *
    FROM Likes l
    WHERE l.coffee = c.name);
```

EXISTS Operator

- Example:
From **Coffees(name, manufacturer)**,
find coffees of manufacturers that
only sell a single coffee

```
SELECT c1.name  
  FROM Coffees c1  
 WHERE NOT EXISTS (  
     SELECT *  
     FROM Coffees c2  
     WHERE c2.manufacturer =  
           c1.manufacturer  
     AND c2.name <> c1.name);
```

Set of coffees with the same
manufacturer as c1, but not
the same name

Practice

Insert the following:

```
INSERT INTO Coffeehouses VALUES  
('Kaffitar', 'Stroget', 'No question');
```

- Using NOT EXISTS, show the name of all coffeehouses that no one frequents

Practice

Insert the following:

```
INSERT INTO Coffeehouses VALUES  
('Kaffitar', 'Stroget', 'No question');
```

- Using NOT EXISTS, show the name of all coffeehouses that no one frequents

```
SELECT ch.name  
FROM Coffeehouses ch  
WHERE NOT EXISTS (  
    SELECT *  
    FROM Frequents f  
    WHERE f.coffeehouse = ch.name);
```

ANY Operator

- $x <op> \text{ANY} (<subquery>)$
 - Returns true iff at least one tuple in the subquery satisfies the comparison operator $=, >, <, \dots$
- Example:
 - For each coffeehouse, show all coffees that are more expensive than some other coffee sold at that coffeehouse

```
SELECT coffee
FROM Sells
WHERE price > ANY (
    SELECT price
    FROM Sells);
```

ALL Operator

- $x \text{ <op> ALL } (<\text{subquery}>)$
 - Returns true iff for every tuple in the relation satisfies the comparison operator $=, <, >, \dots$
- Example:
 - From Sells(coffeehouse, coffee, price), find the coffee(s) sold for the highest price

```
SELECT coffee
FROM Sells
WHERE price >= ALL (
    SELECT price
    FROM Sells
    WHERE price IS NOT NULL);
```


FROM

- Subqueries can also be used in FROM clauses
- Must use a tuple-variable to name tuples of the results
- Example:
 - Find the coffees liked by at least one person who frequents Mocha

```
SELECT DISTINCT l.coffee
FROM Likes l
JOIN (SELECT drinker
      FROM Frequents
      WHERE coffeehouse =
        'Mocha') md
ON l.drinker = md.drinker;
```


-- TODO

- ✓ Division
- ✓ JOIN & NULL
 - ✓ Natural Joins, Cross Joins, Self-Joins
 - ✓ NULL
 - ✓ Three-valued logic
 - ✓ Operators, Outer Joins
- ✓ Set Operations
 - ✓ UNION, INTERSECT, EXCEPT
- ✓ Subqueries (Nested Queries)
 - ✓ =, IN, EXISTS, ALL, ANY
- Views: Queries as subroutines



What is a View?

- A virtual table constructed from actual tables on the fly
 - Accessed in queries like any other table
 - Not materialized, constructed when accessed
 - Similar to a subroutine in ordinary programming
 - Is a schema element

```
CREATE VIEW Name (Columns,...)  
AS  
<SQL QUERY>
```

Example of View (Sports DB)

- For each result, we want:
 - the name of the athlete
 - the name of the sport
 - the percentage of the record achieved by the result (a result that is a record should therefore appear as 100; this column should be named “percentage”).

People (ID, name, gender, height)

Gender (gender, description)

Sports (ID, name, record)

Competitions (ID, place, held)

Results (peopleID, competitionID, sportID, result)

```
SELECT p.name, s.name,  
ROUND(100*CAST((r.result/s.record) AS  
numeric),0) AS percentage  
FROM People p  
JOIN Results r ON p.id = r.peopleid  
JOIN Sports s ON s.id = r.sportsid;
```

Example of View (Sports DB)

- For each result, we want:
 - the name of the athlete
 - the name of the sport
 - the percentage of the record achieved by the result (a result that is a record should therefore appear as 100; this column should be named “percentage”).

People (ID, name, gender, height)

Gender (gender, description)

Sports (ID, name, record)

Competitions (ID, place, held)

Results (peopleID, competitionID, sportID, result)

```
CREATE VIEW
    E2Q10 (pid, pname, sid, sname, percentage)
AS
SELECT p.id, p.name, s.id, s.name,
    ROUND(100*CAST((r.result/s.record) AS
numeric),0) AS percentage
    FROM People p
    JOIN Results r ON p.id = r.peopleid
    JOIN Sports s ON s.id = r.sportid;

SELECT pname, sname, percentage
    FROM E2Q10;
```

Practice (Sports DB)

Create views for these queries:

- The ID, name and gender of all athletes who participated in the competition held in Hvide Sande in 2009.
- The name and gender of all people with last names that starts with a “J” and ends with a “sen”

People (ID, name, gender, height)

Gender (gender, description)

Sports (ID, name, record)

Competitions (ID, place, held)

Results (peopleID, competitionID, sportID, result)

```
SELECT DISTINCT p.id, p.name, g.description
FROM People p
JOIN Gender g ON p.gender = g.gender
JOIN Results r ON p.id = r.peopleid
JOIN Competitions c
  ON c.id = r.competitionid
WHERE c.place = 'Hvide Sande'
AND EXTRACT(YEAR FROM c.held) = 2009;
```

```
SELECT p.name, g.description
FROM People p
JOIN Gender g ON p.gender = g.gender
WHERE p.name LIKE '% J%sen';
```

-- TODO -> DONE

✓ Division

✓ JOIN & NULL

✓ Natural Joins, Cross Joins,
Self-Joins

✓ NULL

✓ Three-valued logic

✓ Operators, Outer Joins

✓ Set Operations

✓ UNION, INTERSECT, EXCEPT

✓ Subqueries (Nested Queries)

✓ =, IN, EXISTS, ALL, ANY

✓ Views: Queries as subroutines

Takeaways

SQL
SEQUEL
SQUEAL
SQUIRREL

One Block

- SELECT, FROM, WHERE, GROUP BY, HAVING, ORDER BY

Many Blocks

- Multi Query
 - Connected by: UNION, INTERCEPT, EXCEPT
- Subqueries in FROM and WHERE
 - Blocks connected by: =, IN, EXISTS, ANY, ALL
 - Opposites (usually) with NOT
 - BUT: Prefer to use Joins when possible
- Division is one important type of queries with >1 query block!

Next Time in IDBS...

Introduction to Database Systems

IDBS - Spring 2024

Week 4 - SQL and Python

Functions

Triggers

SQL and DBMS in Python

Transactions

Readings: PDBM 9.2, 14.1, 14.2.1 and 14.5

Omar Shahbaz Khan