

Introduction to Database Systems

IDBS - Spring 2024

Lecture 4 - SQL Programming and Python

Functions

Triggers & Constraints

SQL and DBMS in Python

Transactions

Readings: PDBM 9.2, 14.1, 14.2.1 and 14.5

Omar Shahbaz Khan

Last Time in IDBS...

-- TODO -> DONE

✓ Division

✓ JOIN & NULL

✓ Natural Joins, Cross Joins,
Self-Joins

✓ NULL

✓ Outer Joins (LEFT, RIGHT, FULL)

✓ Set Operations

✓ UNION, INTERSECT, EXCEPT

✓ Subqueries (Nested Queries)

✓ =, IN, EXISTS, ALL, ANY

✓ Views: Queries as subroutines

Wake Up Task!

Bills DB

People (PID, pName, pGender, pHeight)

Accounts (AID, PID, aDate,
aBalance, aOver)

AccountRecords (RID, AID, rDate, rType,
rAmount, rBalance)

Bills (BID, PID, bDueDate,
bAmount, bIsPaid)

- How many accounts have never been used before?
- How many customers have a negative account balance and also have bills due?

This Time...

-- TODO

- Database Functions
- Database Triggers & Constraints
- DBMS Programming with Python
- Transactions

$$f(x)$$

Functions in SQL

- What is a function?
- Similar to other programming languages methods/functions
- How is it useful in SQL?

Create a new Person

- In the Sports database create a new person
 - Insert a row
- Can be cumbersome to write `INSERT INTO ... VALUES ...` if we want to insert many new rows
- How about a function?

```
INSERT INTO  
People (name, gender, height)  
VALUES ('Terry', 'M', 1.77);
```

Create a new Person

- In the Sports database create a new person
 - Insert a row
- Can be cumbersome to write `INSERT INTO ... VALUES ...` if we want to insert many new rows
- How about a function?

```
DROP FUNCTION IF EXISTS NewPerson;
```

```
CREATE FUNCTION NewPerson (
```

Function
Parameters

```
IN pname VARCHAR(50),  
IN pgender CHAR(1),  
IN pheight FLOAT
```

```
)
```

```
RETURNS INTEGER
```

Return Type

```
AS $$
```

```
BEGIN
```

Insert
Statement

```
INSERT INTO  
Person (name, gender, height)  
VALUES (pname, pgender, pheight);  
RETURN lastval();
```

Return the last value of
the table's SEQUENCE

```
END;
```

```
$$ LANGUAGE plpgsql;
```

Using Functions

- From an SQL Script
- Typically used from an ODBC (Java, Python, ...)

```
SELECT NewPerson('Terry', 'M', 1.77);
```

```
SELECT * FROM NewPerson('Terry', 'M', 1.77);
```

```
DO $$  
    BEGIN  
        PERFORM NewPerson('Terry', 'M', 1.77);  
    END  
$$;
```

Let's make a function (I)

- **Function: BiggestRecordJump**
 - Input: ID of sport
 - Output: The largest record increase of that sport
- **Assume we have the RecordLog table with old and new records of sports**
 - Later: maintain this table automatically

```
RecordLog (  
    peopleID INT,  
    competitionID INT,  
    sportID INT,  
    oldrecord FLOAT,  
    newrecord FLOAT,  
    seton DATE  
);
```

Let's make a function (II)

- Function: BiggestRecordJump
 - Input: ID of sport
 - Output: The largest increase of that sport
- Assume RecordLog table:

```
RecordLog (  
    peopleID INT,  
    competitionID INT,  
    sportID INT,  
    oldrecord FLOAT,  
    newrecord FLOAT,  
    seton DATE  
);
```

```
DROP FUNCTION IF EXISTS BiggestRecordJump;  
  
CREATE FUNCTION BiggestRecordJump (  
    ...  
)  
RETURNS FLOAT  
AS $$  
DECLARE r FLOAT; Variable  
BEGIN  
    SELECT ... INTO r  
    FROM RecordLog  
    WHERE ...  
    RETURN r;  
END;  
$$ LANGUAGE plpgsql;
```

Let's make a function (III)

- Function: BiggestRecordJump
 - Input: ID of sport
 - Output: The largest increase of that sport
- Assume RecordLog table:

```
RecordLog (  
    peopleID INT,  
    competitionID INT,  
    sportID INT,  
    oldrecord FLOAT,  
    newrecord FLOAT,  
    seton DATE  
);
```

```
DROP FUNCTION IF EXISTS BiggestRecordJump();  
  
CREATE FUNCTION BiggestRecordJump(  
    IN sid INT  
)  
RETURNS FLOAT  
AS $$  
DECLARE r FLOAT;  
BEGIN  
    SELECT MAX(newrecord - oldrecord) INTO r  
    FROM RecordLog  
    WHERE sportID = sid  
    RETURN r;  
END;  
$$ LANGUAGE plpgsql;  
  
SELECT BiggestRecordJump(1);
```

Are Functions Faster?

- May be faster than executing from a client
 - Why?
- Code may be pre-compiled and optimized
 - Do not need to invoke optimizer again
 - May occur with well written queries using plan caching
- The code runs at the server
 - The server may be more efficient
 - No need to move data

Pros & Cons

- Code shared across ALL applications
- May be used for access control
- May give performance benefits
- Very system-specific
- Code maintenance requires care
 - Versioning is difficult

IDENTITY Columns

- Run 04-sports-schema.sql
- Try inserting a new person
- Due to the ID column not being an IDENTITY column it will fail inserting the row
- Quickfix: Add parameter pid in the NewPerson Function
- Best to avoid this by using `GENERATED ALWAYS AS IDENTITY` when creating the table
- Fixed tables with ID columns correctly set to `GENERATED ALWAYS AS IDENTITY` (run 04-sports-schema-fixed.sql)
- In case you want to insert a row with a specific ID, you can do so by `OVERRIDING SYSTEM VALUE`. See example in the file.

-- TODO

✓ Database Functions

- Database Triggers & Constraints
- DBMS Programming with Python
- Transactions



Triggers in SQL

- Automatically executed function in response to certain events on a table (or view) in a database
 - Events are typically INSERT, UPDATE, DELETE
 - Triggers can execute BEFORE or AFTER the event
- Useful for maintaining the integrity of the information on the Database.
 - For example, in a company database when a new record (representing a new worker) is added to the employees table, new records could also be created in the tables of taxes, vacations and salaries.
- Logging historical data, for example to keep track of employees' previous salaries.

Triggers in PostgreSQL

- Multiple triggers per table per event
 - INSERT / UPDATE / DELETE
 - Run in alphabetical order
 - Per row OR Per statement – focus here on per row
- New data is in the **NEW** record
 - For INSERT / UPDATE
 - Same schema as modified relation
 - Can refer to **NEW.ID**, **NEW.name**, ...
- Old data is in the **OLD** record
 - For UPDATE / DELETE
 - Same schema, refer to **OLD.ID**, **OLD.name**, ...
- The variable TG_OP says which operation it is

Trigger: Checking Values

- Use triggers to check value semantics
- A result in the Sports DB must not be negative

```
CREATE FUNCTION CheckResult()  
RETURNS TRIGGER  
AS $$ BEGIN  
    IF (NEW.result < 0.0) THEN  
        RAISE EXCEPTION  
        'CheckResult: Result must be a  
        positive'  
        USING ERRCODE = '45000';  
    END IF;  
    RETURN NEW;  
END; $$ LANGUAGE plpgsql;
```

Create
Trigger based
on function

```
CREATE TRIGGER CheckResult  
BEFORE INSERT OR UPDATE  
ON Results  
FOR EACH ROW EXECUTE PROCEDURE CheckResult();
```

Trigger: Ban Updates/Deletes


- Use triggers to not allow updating or deleting rows

```
CREATE FUNCTION BanChanges()  
RETURNS TRIGGER  
AS $$  
BEGIN  
    RAISE EXCEPTION  
        'BanChanges: Cannot change  
        results!'  
    USING ERRCODE = '45000';  
END;  
$$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER BanChanges  
BEFORE UPDATE OR DELETE  
ON Results  
FOR EACH ROW EXECUTE PROCEDURE  
BanChanges();
```

Trigger: Update Records

- Use triggers to update the values after an event
- If a result in Sports DB is a new record, the row for that sport in the sports table need to be updated

```
CREATE FUNCTION UpdateRecord()  
RETURNS TRIGGER  
AS $$ BEGIN  
    IF NEW.result > (  
        SELECT s.record  
        FROM Sports s  
        WHERE s.id = NEW.sportID  
    ) THEN  
        UPDATE Sports  
        SET record = NEW.result  
        WHERE s.id = NEW.sportID;  
    END IF;  
    RETURN NEW;   
END; $$ LANGUAGE plpgsql;
```

Could also be **RETURN NULL** since result of AFTER trigger is ignored

```
CREATE TRIGGER UpdateRecord  
AFTER INSERT OR UPDATE  
ON Results  
FOR EACH ROW EXECUTE PROCEDURE UpdateRecord();
```


Trigger: Log Changes

```
CREATE TABLE RecordLog (  
    peopleID INT,  
    competitionID INT,  
    sportID INT,  
    oldrecord FLOAT,  
    newrecord FLOAT,  
    seton DATE  
    PRIMARY KEY  
    (peopleID, competitionID, sportID)  
    FOREIGN KEY  
    (peopleID, competitionID, sportID)  
    REFERENCES Results  
    (peopleID, competitionID, sportID)  
);
```

```
CREATE FUNCTION LogRecord()  
RETURNS TRIGGER  
DECLARE oldRecord FLOAT;  
AS $$ BEGIN  
    IF NEW.result > (  
        SELECT s.record  
        FROM Sports s  
        WHERE s.id = NEW.sportID  
    ) THEN  
        SELECT s.record INTO oldRecord  
        FROM Sports s  
        WHERE s.id = NEW.sportID;  
        INSERT INTO RecordLog  
        VALUES (NEW.peopleID, NEW.competitionID,  
            NEW.sportID, oldRecord,  
            NEW.result);  
    END IF;  
    RETURN NEW;  
END; $$ LANGUAGE plpgsql;  
  
CREATE TRIGGER LogRecord  
AFTER INSERT ON Results  
FOR EACH ROW EXECUTE PROCEDURE LogRecord();
```

Questions

- Why not create the LogRecord on Sports?
- What about Trigger order?
 - LogRecord vs UpdateRecord
 - Order is alphabetical = good in this case
 - Try with different order
 - May be better to merge similar triggers
- What about BEFORE or AFTER?
 - LogRecord and UpdateRecord are AFTER
 - CheckResult and BanUpdates are BEFORE
 - BEFORE and AFTER what? When does each apply? Why?

Merged Trigger (BanUpdates + Check Updates)

```
CREATE FUNCTION MergedTrigger()  
RETURNS TRIGGER  
AS $$ BEGIN  
    IF (TG_OP = 'DELETE' OR TG_OP = 'UPDATE')  
    THEN  
        RAISE EXCEPTION  
        'MergedTrigger: Cannot change results!'  
        USING ERRCODE = '45000';  
    END IF;  
    IF (NEW.result < 0.0) THEN  
        RAISE EXCEPTION  
        'CheckResult: Result must be a positive!'  
        USING ERRCODE = '45000';  
    END IF;  
END; $$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER MergedTrigger  
BEFORE INSERT OR UPDATE OR DELETE  
ON Results  
FOR EACH ROW EXECUTE PROCEDURE  
MergedTrigger();  
  
INSERT INTO Results  
VALUES (1,1,3,-1.0);  
DELETE FROM Results WHERE sportID = 3;
```

BEFORE vs AFTER

1. Are you only checking the newly inserted/updated entry? -> **BEFORE (or AFTER)**
 - Checking happens earlier with BEFORE, so less work
 - But if other triggers might modify the values, then prefer AFTER
2. Are you inserting a row to another table with a foreign key constraint to the NEW record? -> **AFTER**
 - Otherwise, the NEW record is NOT in the database, so your insertion will fail!
3. Are you modifying the NEW record? -> **BEFORE**
 - Otherwise, the record is already in the database and will not be changed
4. Are you doing both 2 and 3? -> **BEFORE and AFTER**
 - Two different triggers!

Are Triggers Faster?

- May be faster than executing the code from a client
 - Why?
- Code may be pre-compiled and optimized
 - Do not need to invoke optimizer again
 - But this may happen with well written queries using plan caching
- The code runs at the server
 - The server may be more efficient
 - No need to move data to client

Pros & Cons

- Code often runs faster
 - No context switch
 - Compiled code
 - No data transfers
- Useful for security
 - Wrap data and functionality
 - Same access from all clients
- Same code for all applications
- Code is “hidden”
 - Only visible via system tables
 - Easily forgotten
 - Versioning is hard!
 - Schema may be edited using a GUI = no-no!
- Generally not portable

Exercises on SQL Programming

- Exercises 4 are out: Views, procedures, triggers
 - Use test script to verify your work
 - Can run parts in pgAdmin – best via command prompt
 - You can extend the script with your own tests!
 - Also has a very nice database cleaning script

```
# once
psql -q Bills < bills-schema.sql

#repeatedly
psql Bills < universal-cleanup.sql
psql Bills < your-solution.sql
psql Bills < test-script.sql > output 2>&1
less output
```

NOTE: Windows uses '-f' instead of '<' and requires the DB name at the end of the command

NOTE 2: You may also be required to add user, -U <username>

Exercises on SQL Programming (Catalog)

- All information about the database is stored in tables!
 - This is the catalog!

<https://www.postgresql.org/docs/current/catalogs.html>

- See also: System information functions

<https://www.postgresql.org/docs/current/functions-info.html>

```
SELECT *  
  FROM information_schema.tables  
 WHERE table_schema = 'public';
```

```
SELECT current_database();
```


```
SELECT current_user;
```

```
SELECT lastval();
```




-- TODO

- ✓ Database Functions
- ✓ Database Triggers & Constraints
 - DBMS Programming with Python
 - Transactions



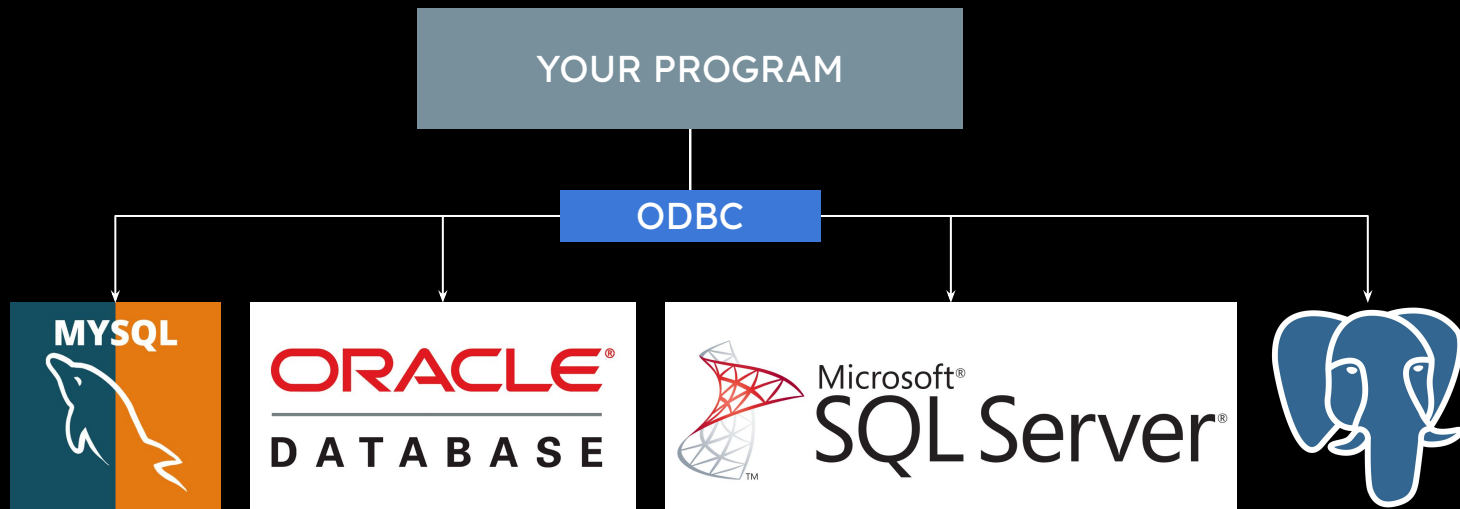
Problem: Vendor Lock-in

- They're all compatible with standard SQL...
- ...and mutually incompatible due to vendor-specific functionality!



Open Database Connectivity (ODBC)

- In-between translation layer
- Vendors not so happy as lock-in is profitable!



Python and PostgreSQL

- Database adapter: psycopg package
 - Reliable, documented, references
 - To install: `python -m pip install psycopg`
 - `import psycopg`
- A simple database API
 - `connect()`: Establish connection to specified database
 - `cursor()`: Object to manage context of an SQL operation
 - `execute()`: Execute SQL queries
 - `fetchone()`, `fetchmany()`, `fetchall()`: Get results from queries
 - Much more, but these are the most used functions

Connecting to a Database

- All code for Python is in 04-slides-python.py
- Using a connection string

```
import psycopg2 as pg

connString = "host=localhost
dbname=sports user=postgres
password=_____"

conn = pg.connect(connString)

with pg.connect(connString) as conn:
```

Querying the DB

- INSERT query example
- Explicitly declare a cursor
 - Requires call to close()
- Use `with` block
 - Closes automatically at end of block

```
cur = conn.cursor()
```

```
cur.execute("""  
    INSERT INTO  
    People (name,gender,height)  
    VALUES ('Terry', 'M', 1.77)  
    """)
```

```
cur.close()
```

```
with conn.cursor() as cur:  
    cur.execute("""  
        INSERT INTO  
        People (name,gender,height)  
        VALUES ('Terry', 'M', 1.77)  
        """)
```


Querying the DB

- INSERT query example
- By default changes are not committed to the Database
 - Can set autocommit=True in connect(...)

```
cur = conn.cursor()

cur.execute("""
    INSERT INTO
    People (name,gender,height)
    VALUES ('Terry', 'M', 1.77)
    """)

conn.commit()
```

Querying the DB

- **SELECT** query example
 - Could also use `fetchone()` here

```
cur = conn.cursor()

cur.execute("""
    SELECT *
    FROM People
    WHERE name = 'Terry'
""")

people = cur.fetchall()

print(people)
```

Querying with Variables

- Store the name in a variable
- Inject the variable into the query string
- Is this good or bad?

Prone to SQL Injection!

```
name = "'Terry'"
cur = conn.cursor()

cur.execute("""
    SELECT *
    FROM People
    WHERE name = %s
    """ % name)

people = cur.fetchall()

print(people)
```

SQL Injection

- Use a variable to escape the current query and inject own query
- How can we avoid this?

```
name = ''; DELETE FROM People WHERE  
name='Terry'; --"  
cur = conn.cursor()  
  
cur.execute("""  
    SELECT *  
    FROM People  
    WHERE name = %s  
    """ % name)  
  
people = cur.fetchall()  
  
print(people)
```

Prepared Statements

- Separate the query string and variables
- Prepared Statements
 - Pre-compiled SQL statement at DBMS
 - Client supplied data is treated as content of a parameter and not an SQL statement

```
PREPARE my_query AS
SELECT *
  FROM People
 WHERE name = "$1"

EXECUTE my_query ("Terry")
```

Prepared Statements with Python

- Separate query string and variable in execute
- Automatically determines type

```
name = "Terry"
cur = conn.cursor()

cur.execute("""
    SELECT *
    FROM People
    WHERE name = %s
""", [name])

people = cur.fetchall()

print(people)
```

Prepared Statements with Python

- You can be even more specific
- General notion, never trust the user, always check input

```
name = "Terry"
cur = conn.cursor()

cur.execute("""
    SELECT *
    FROM People
    WHERE name = %(name)s
""", { 'name' : name })

people = cur.fetchall()

print(people)
```

Tips

- Always remember to close the cursor after use and the connection when done
- Use either the `with` block or `try-except-finally`

```
try:
    # QUERY
except Exception as error:
    print(error)
    conn.rollback() # If autocommit is off
finally:
    cur.close()
```

```
with conn.cursor() as cur:
    # QUERY
```


-- TODO

- ✓ Database Functions
- ✓ Database Triggers & Constraints
- ✓ DBMS Programming with Python
 - Transactions



Transfer 50 DKK from Bob to Alice

Step 1: Get balance of Bob

```
SELECT Balance INTO x
FROM accounts
WHERE AccountName = 'Bob'
```

Account Name	Balance
Bob	200
Alice	0

Account Name	Balance
Bob	150
Alice	0

Step 3: Get balance of Alice

```
SELECT Balance INTO y
FROM accounts
WHERE AccountName = 'Alice'
```

WHAT IF SOMETHING CRASHES?

Step 2: Reduce balance of Bob

```
UPDATE accounts
SET Balance = x-50
WHERE AccountName = 'Bob'
```

Account Name	Balance
Bob	150
Alice	0

Account Name	Balance
Bob	150
Alice	50

Step 4: Increase balance of Alice

```
UPDATE accounts
SET Balance = y+50
WHERE AccountName = 'Alice'
```

What is a Transaction?

- A group of related operations to the database
- Wish for “all or nothing” execution
- Wish for isolation from other transactions

Basic SQL Syntax:

```
BEGIN;
```

```
COMMIT;
```

```
ROLLBACK;
```

Savepoints:

```
SAVEPOINT <name>;
```

```
ROLLBACK TO SAVEPOINT <name>;
```

Transactions in PostgreSQL

- By default, every statement is a transaction
 - To override this behaviour:
 - `BEGIN; ... COMMIT; / ROLLBACK;`
 - Some (DDL) statements implicitly `COMMIT` transactions
- Calling a function starts a transaction
 - Can assume that the function has transactional properties!
 - Errors abort the transaction, erase all previous operations!
- Errors inside functions:
 - Cannot simply say: `ROLLBACK`
 - Need to raise and handle exceptions!
 - See example in exercise code

A C I D

Transaction Properties

Recovery

Constraints, Triggers

Concurrency
Control

Recovery

A TOMICITY

A Transaction is
“one operation”

CONSISTENCY

Each Transaction
moves the DB from one
consistent state to
another

ISOLATION

Each Transaction
is alone in the
world

DURABILITY

Persistence of
successful
transactions even
through system failure

Transaction Implementation Methods (Lecture 10)

- Consistency
 - PRIMARY and FOREIGN KEY = limited DBMS support
- Isolation
 - Historically locking = (strict/rigorous) two-phase locking
 - Recently multi-version concurrency control
- Atomicity / Durability
 - Logging all changes to disk
 - Write Ahead Logging protocol (WAL)

Transactions and Testing

- Transactions are useful in test scripts
 - Multiple examples in Exercise 4
- Simple pattern to test changing the database
 - ... without actually changing the database

```
BEGIN
```

```
-- Make changes
```

```
-- Run test queries
```

```
ROLLBACK
```

Transactions in psycopg (Auto Commit)

- By default the connection to the database starts with autocommit off
 - If a `conn.commit()` is not done, no permanent changes will occur to the database
 - `with conn.transaction()` block is a new way to run transactions, example in exercise code
- With auto commit on, every execute command will be persistent after execution

```
conn = pg.connect(connString)
conn.autocommit=True
```

```
conn = pg.connect(connString,
                  autocommit=True)
```

Rollback (typically in the except block)

```
conn.rollback()
```

-- TODO -> DONE

- ✓ Database Functions
- ✓ Database Triggers & Constraints
- ✓ DBMS Programming with Python
- ✓ Transactions

Takeaways

Transaction failed...
Rolling back...

Functions

Set of database operations, performance benefits

Triggers

Execute a function in response to a database event

ODBC (Python **psycopg**, Java **JDBC**)

API for database operations

Transactions

Atomicity, Consistency, Isolation, Durability

Next Time in IDBS...

Introduction to Database Systems

IDBS - Spring 2024

The Void

Live Q&A Session on Piazza 09:00-10:00

Time to catch up...

Exercises to do...

To read or not to read...

Next Next Time in IDBS...

Introduction to Database Systems

IDBS - Spring 2024

Lecture 5 - Designing Databases

ER Diagrams

Translation to SQL DDL

Readings: PDBM 3.0-3.3, 6.3-6.4

Omar Shahbaz Khan