

# Introduction to Database Systems

# Intro to Big Data Management

Eleni Tzirita Zacharatou

Based on slides from Björn Þór Jónsson

Week 12



# First Things First: Rest of Class

- Trial exam next Week 13 (May 7)
  - 08:15 – 12:15: Quiz on LearnIT
    - Piazza/email ([elza@itu.dk](mailto:elza@itu.dk)) if you have questions
- Homework 4 is out (deadline: May 7)
  - Individual work
- Questions: Piazza/email
- Course Evaluation Survey is open

# Lecture Outline

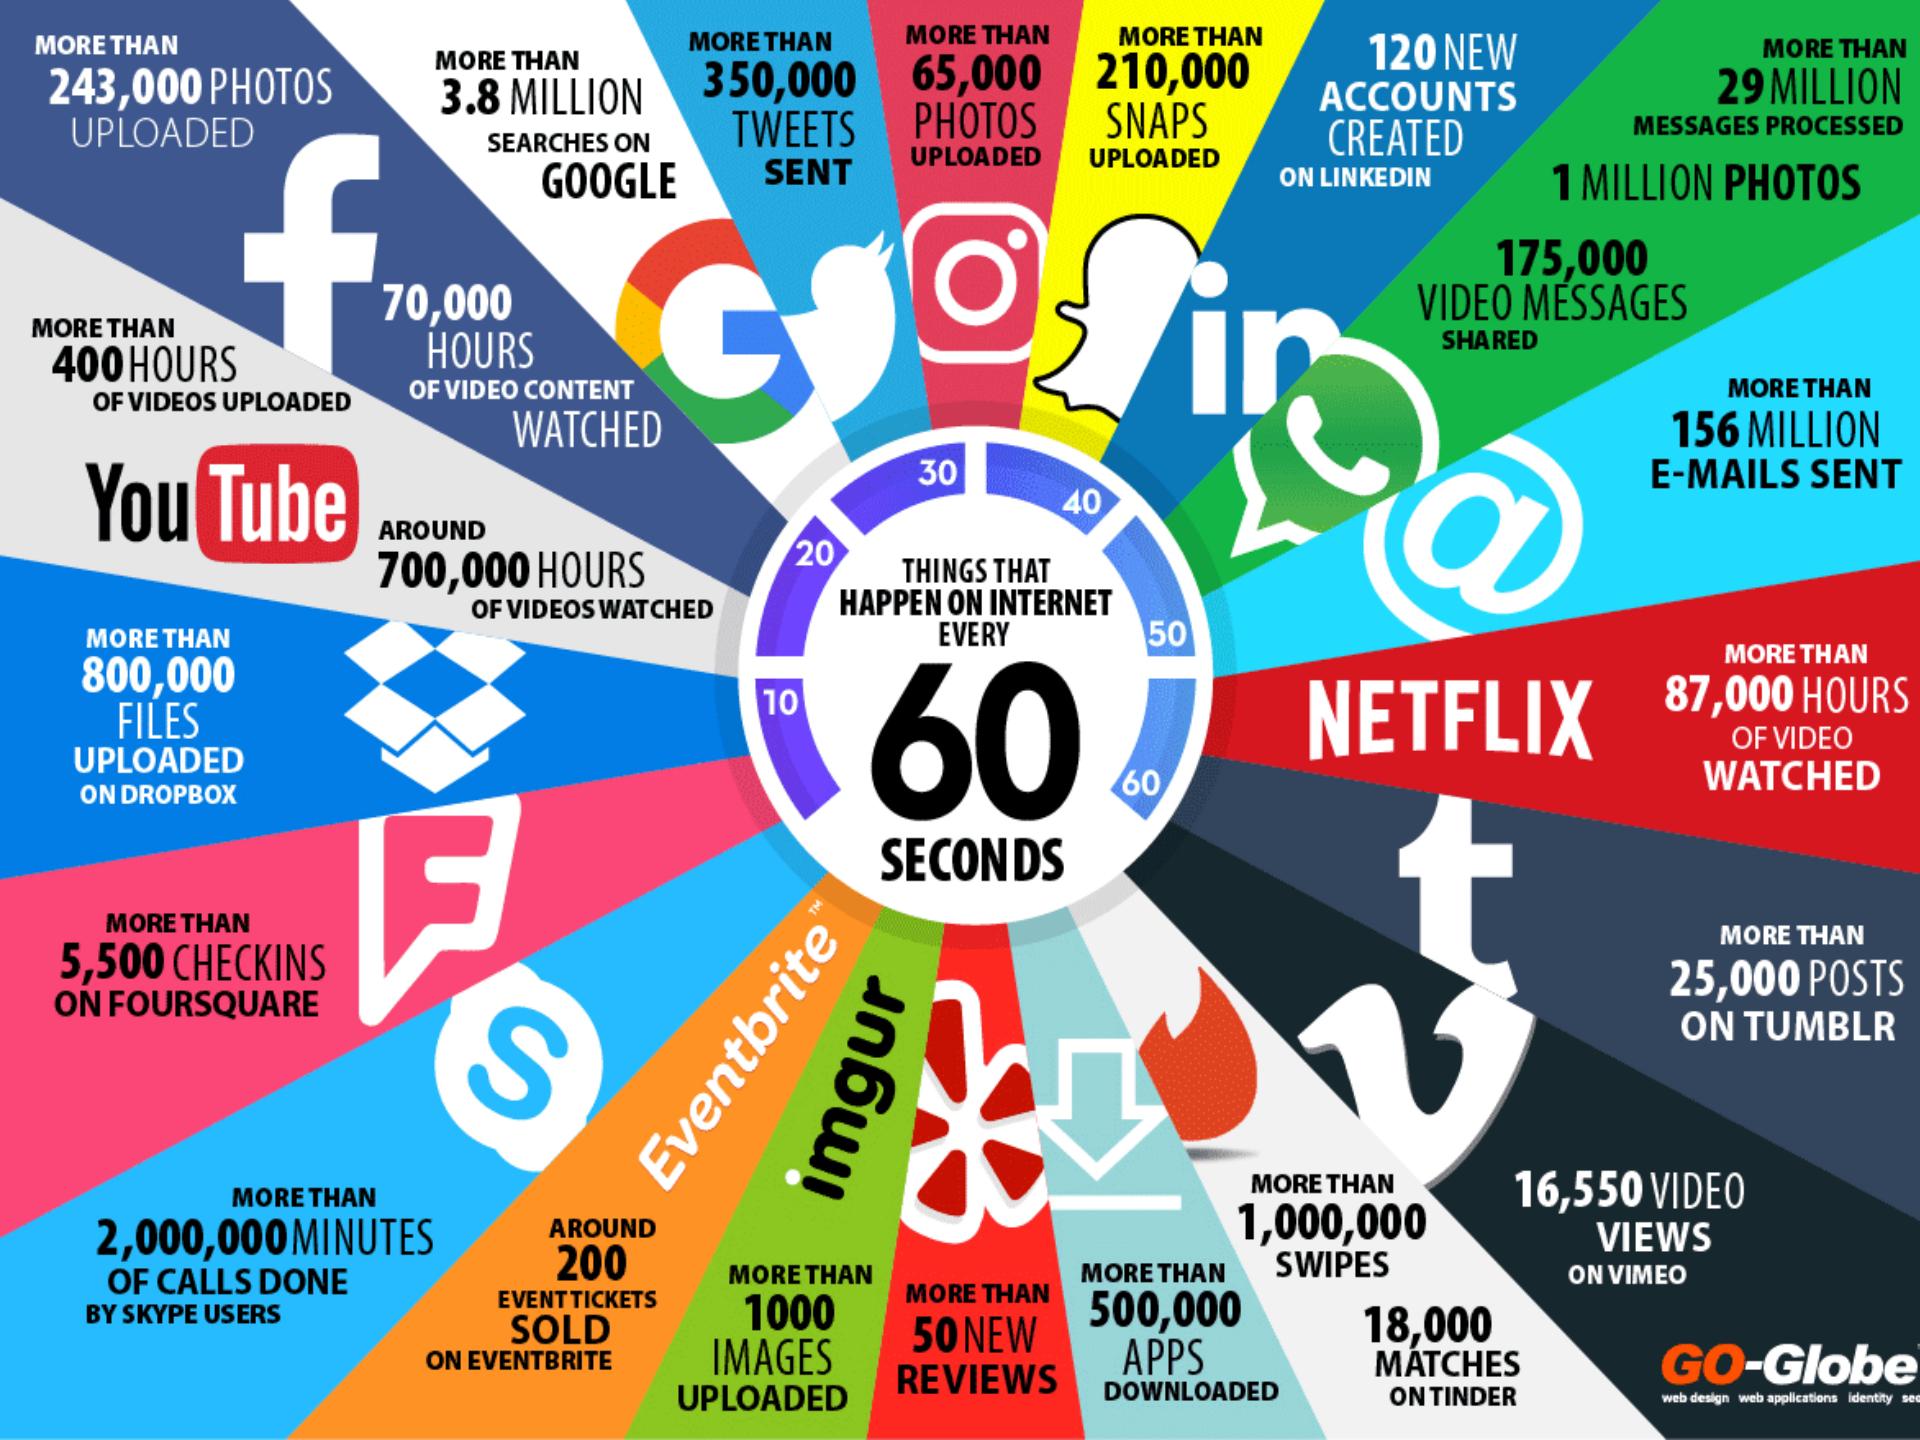
**Learning Outcome 6:** Discuss the pros and cons of different classes of data systems for modern analytics and data science applications.

- What are modern analytics and data science applications?
- What are their access patterns and processing requirements?
- How well/badly do the classes of data systems covered so far support these access patterns?
- MapReduce/Spark: Their properties and their support for these access patterns

# Lecture Outline

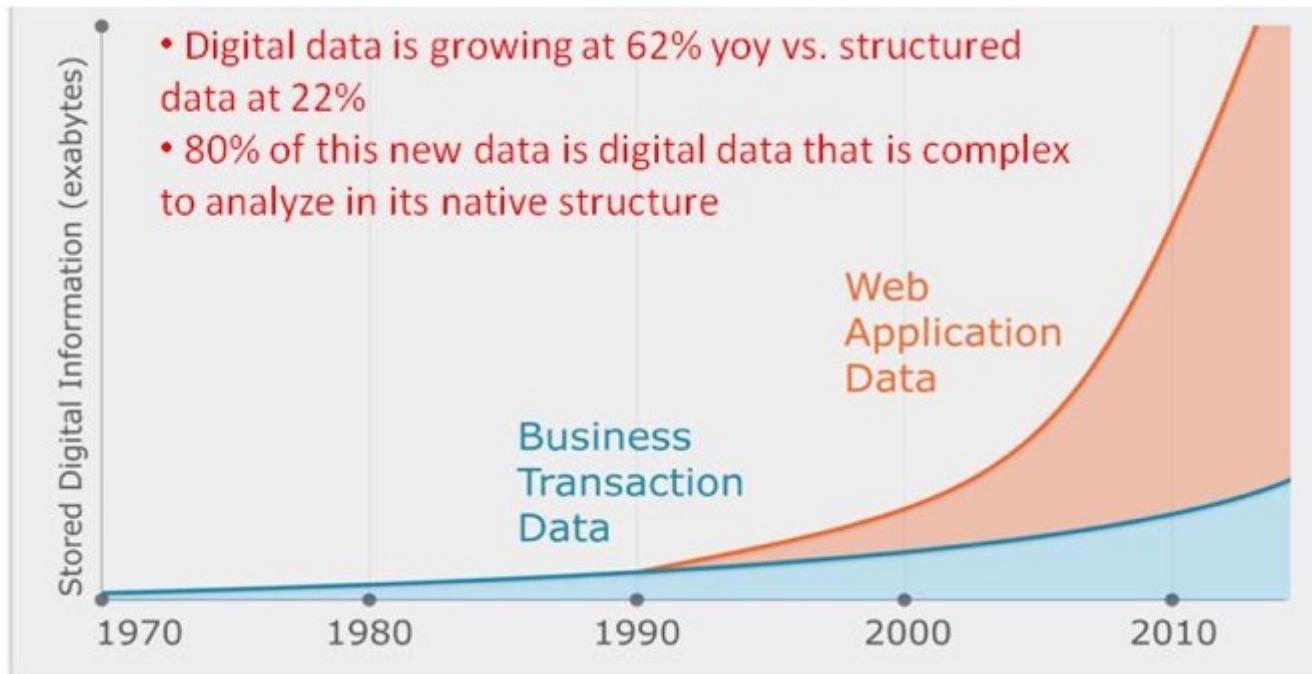
**Learning Outcome 6:** Discuss the pros and cons of different classes of data systems for modern analytics and data science applications.

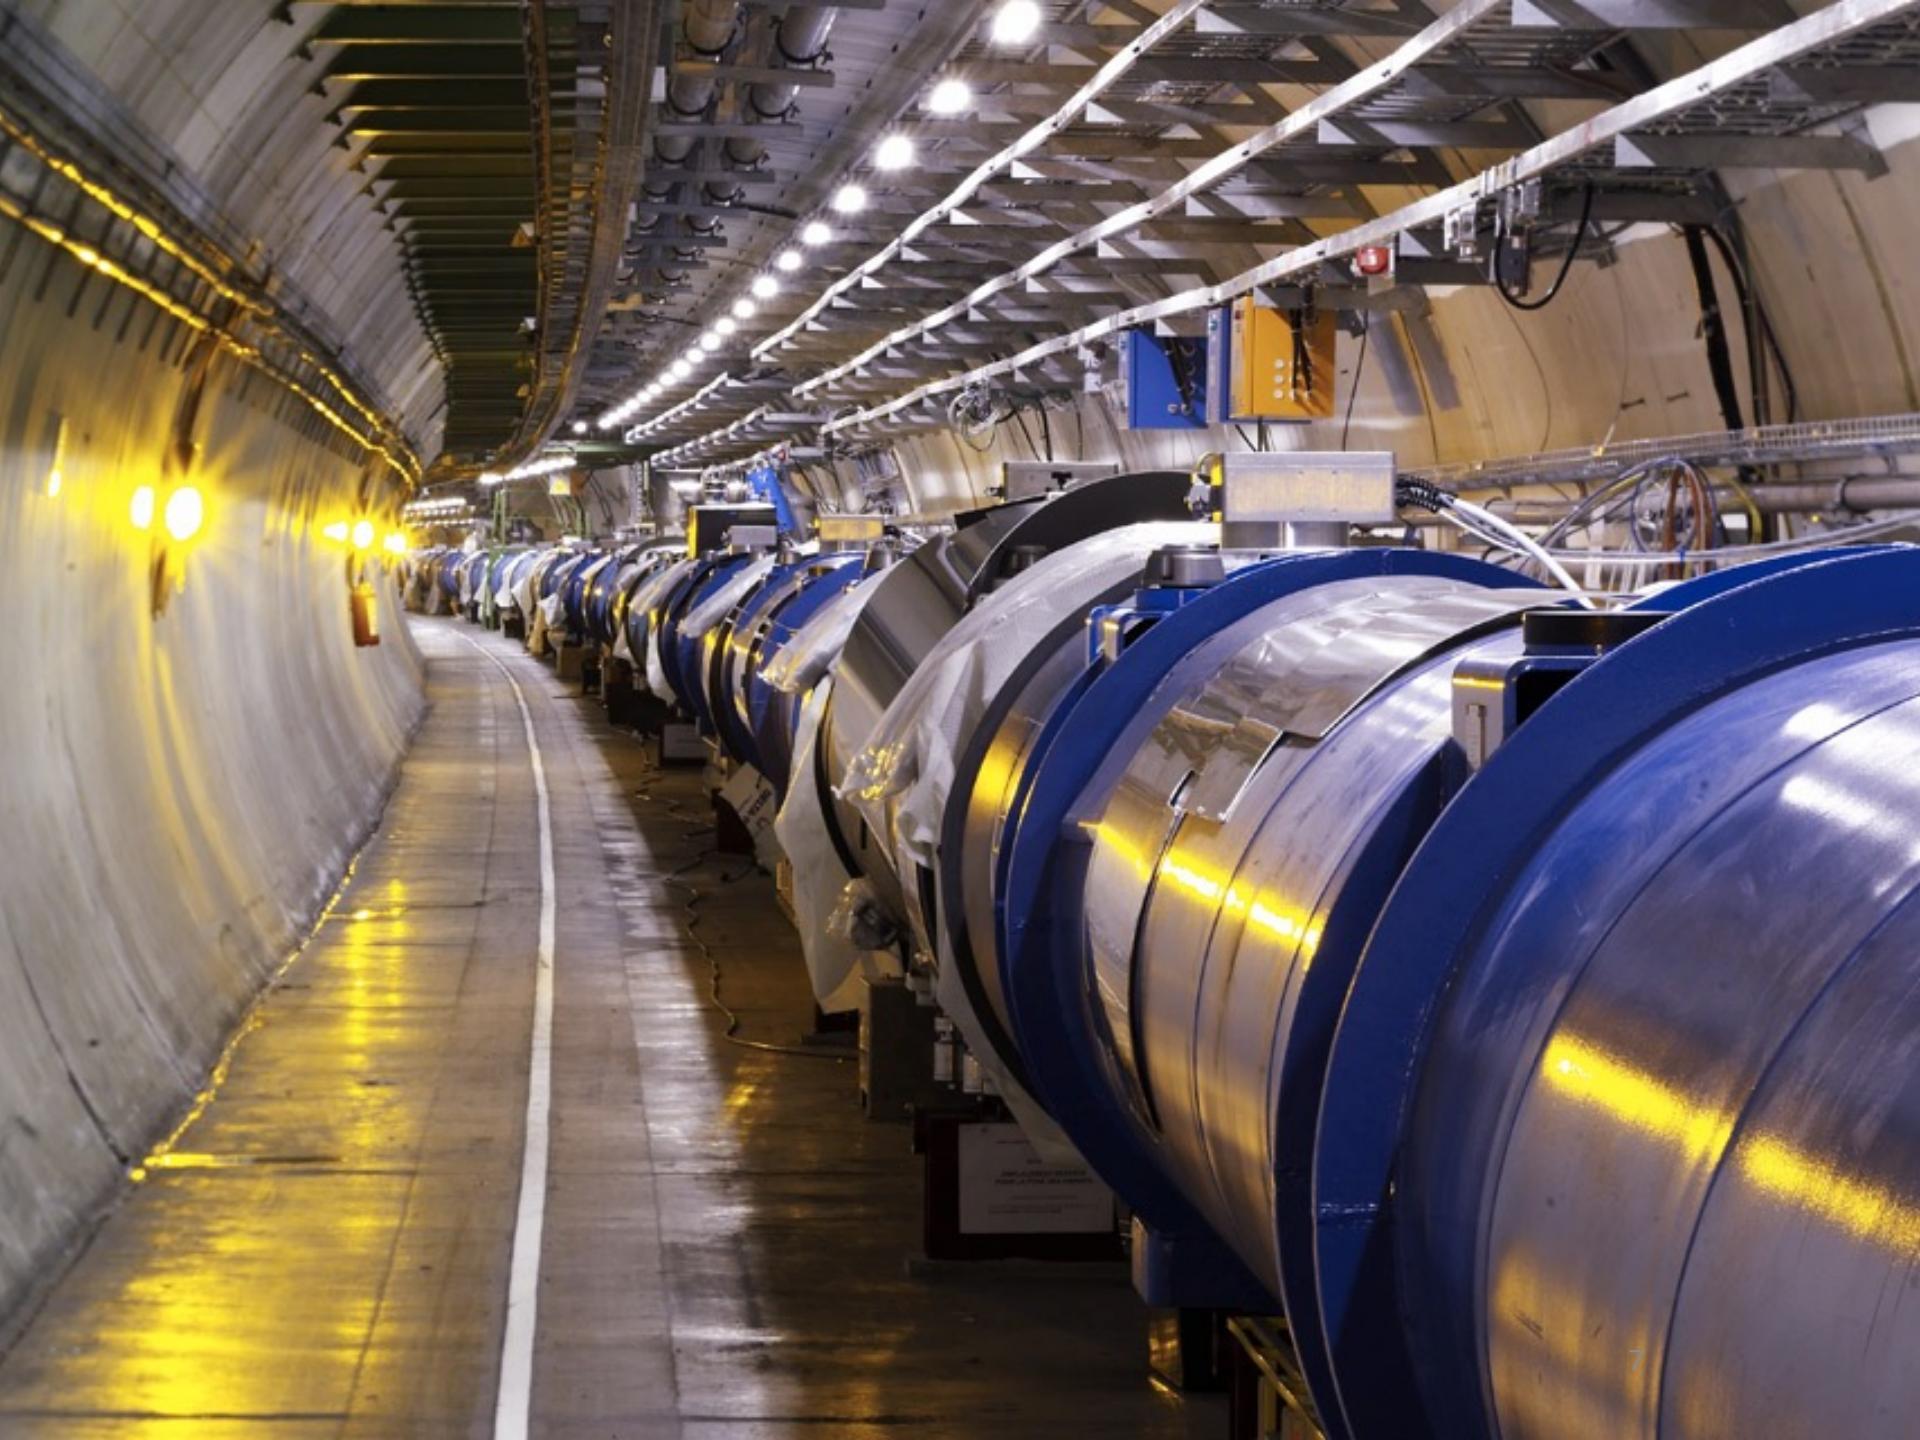
- What are modern analytics and data science applications?
- What are their access patterns and processing requirements?
- How well/badly do the classes of data systems covered so far support these access patterns?
- MapReduce/Spark: Their properties and their support for these access patterns



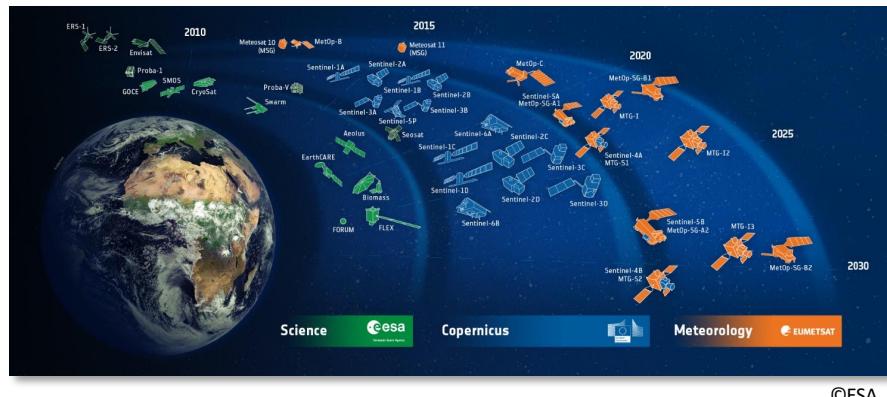
# Beyond Relational Data!

- Keep track of **all** the history
- Keep track of **all** interactions, also low-level
- Keep track of **all** data: media, text data, logs, ...

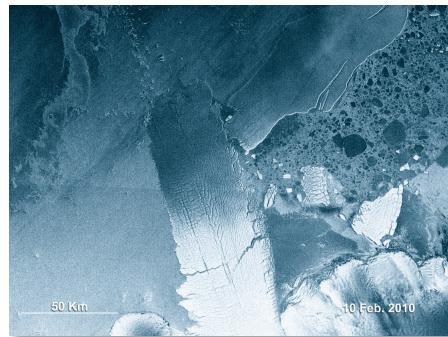




# Earth Observation Big Data Analytics

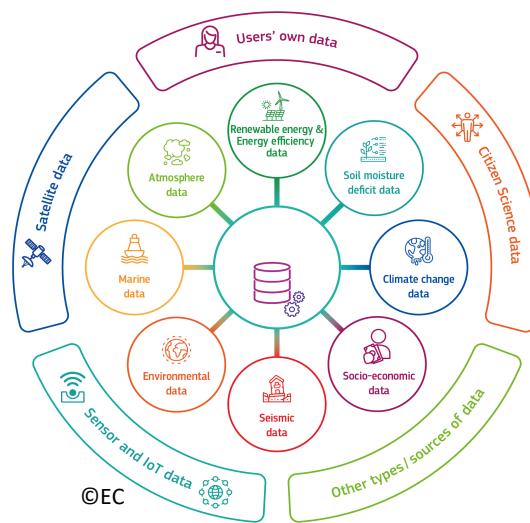


# Agricultural monitoring

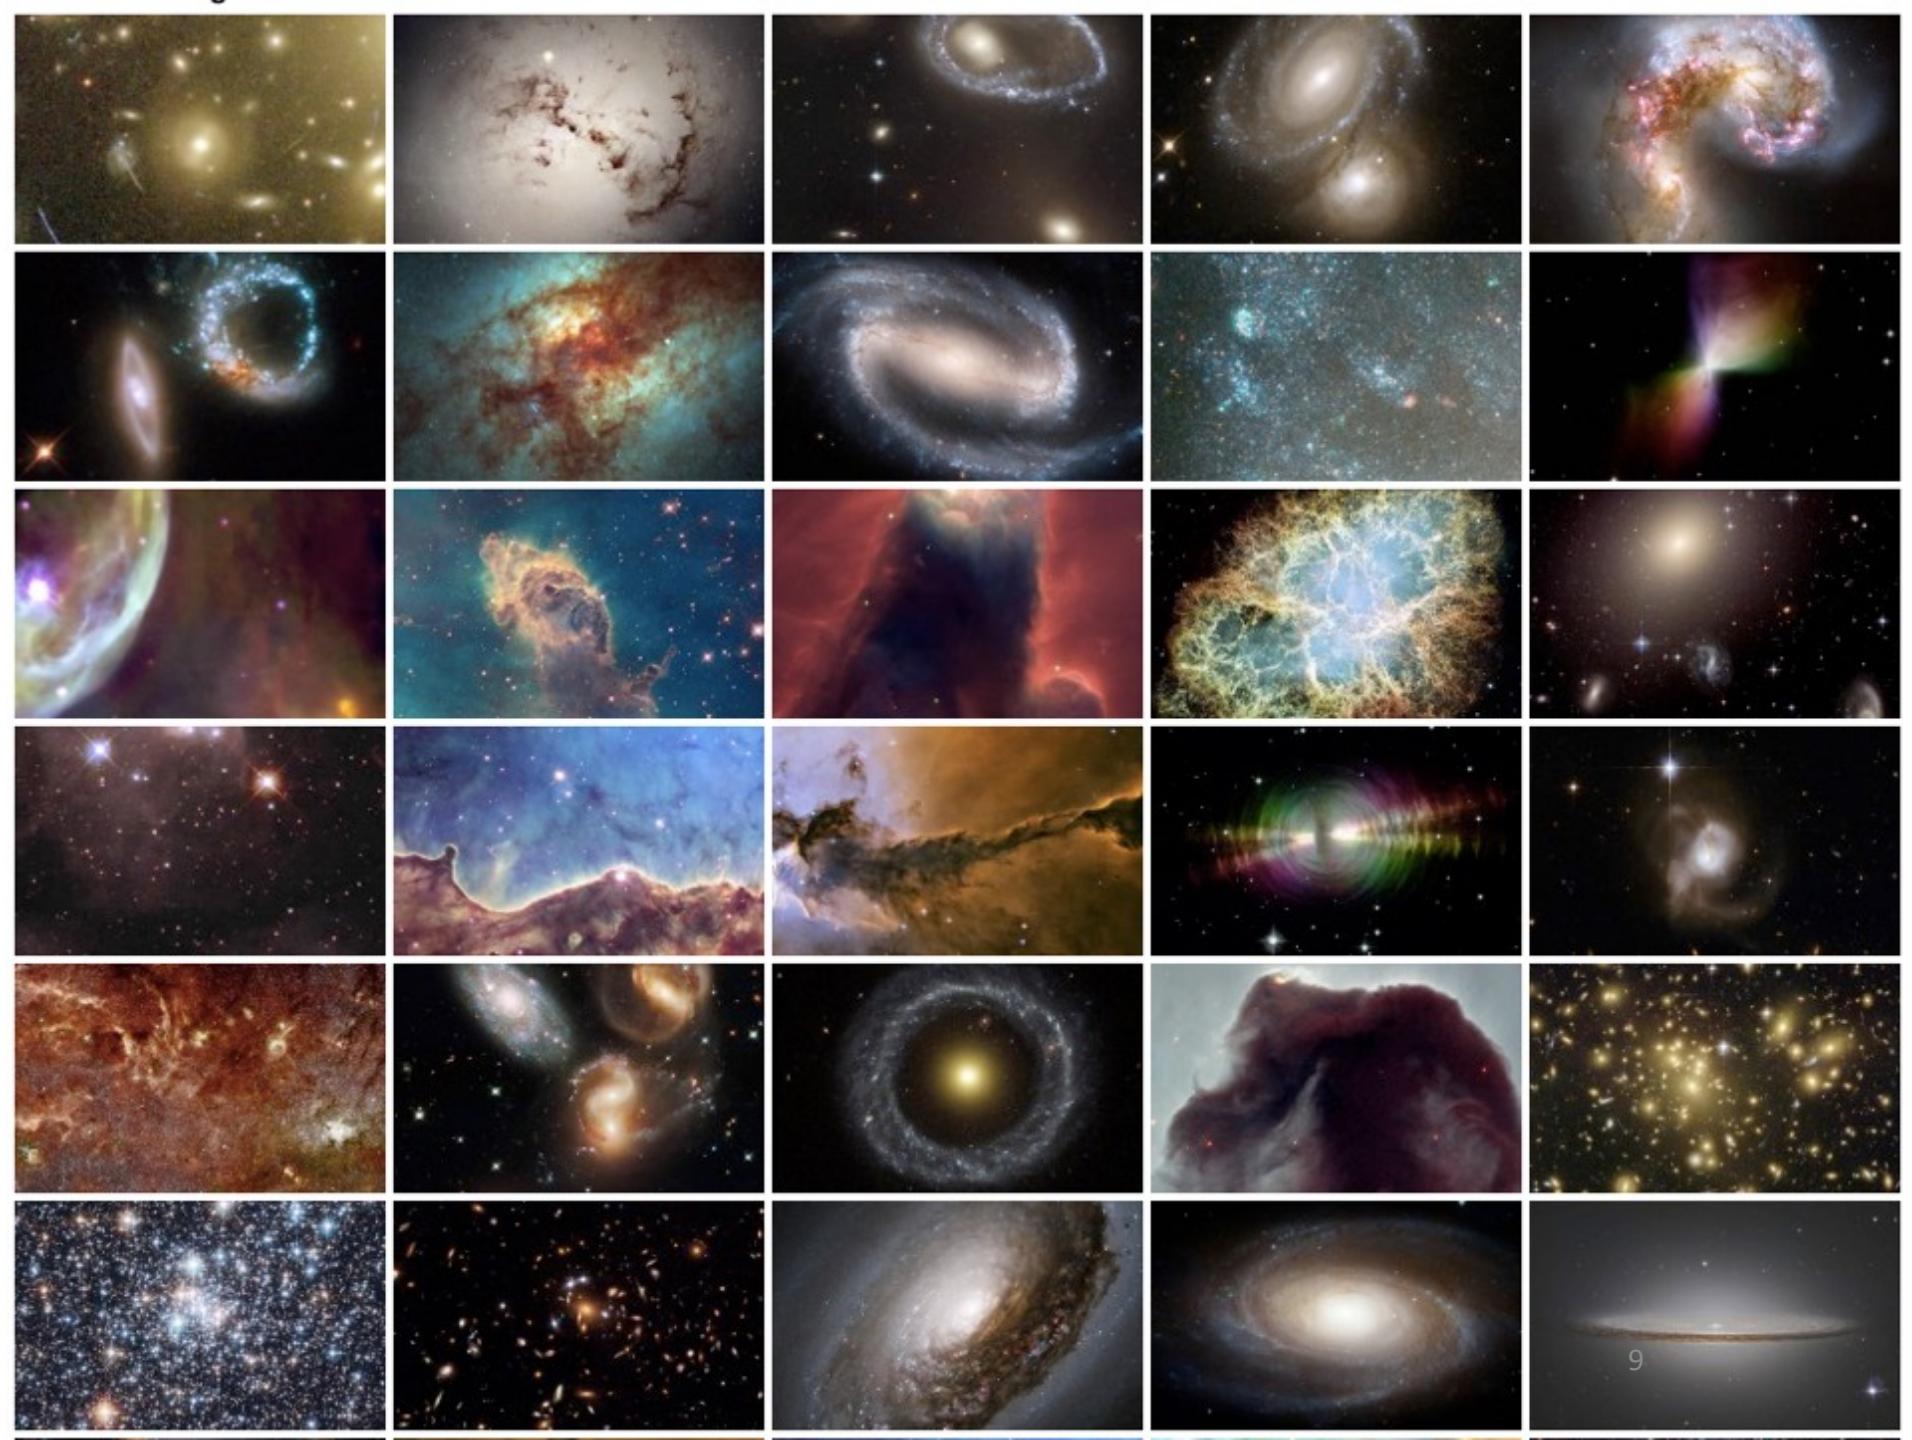


## Climate change monitoring

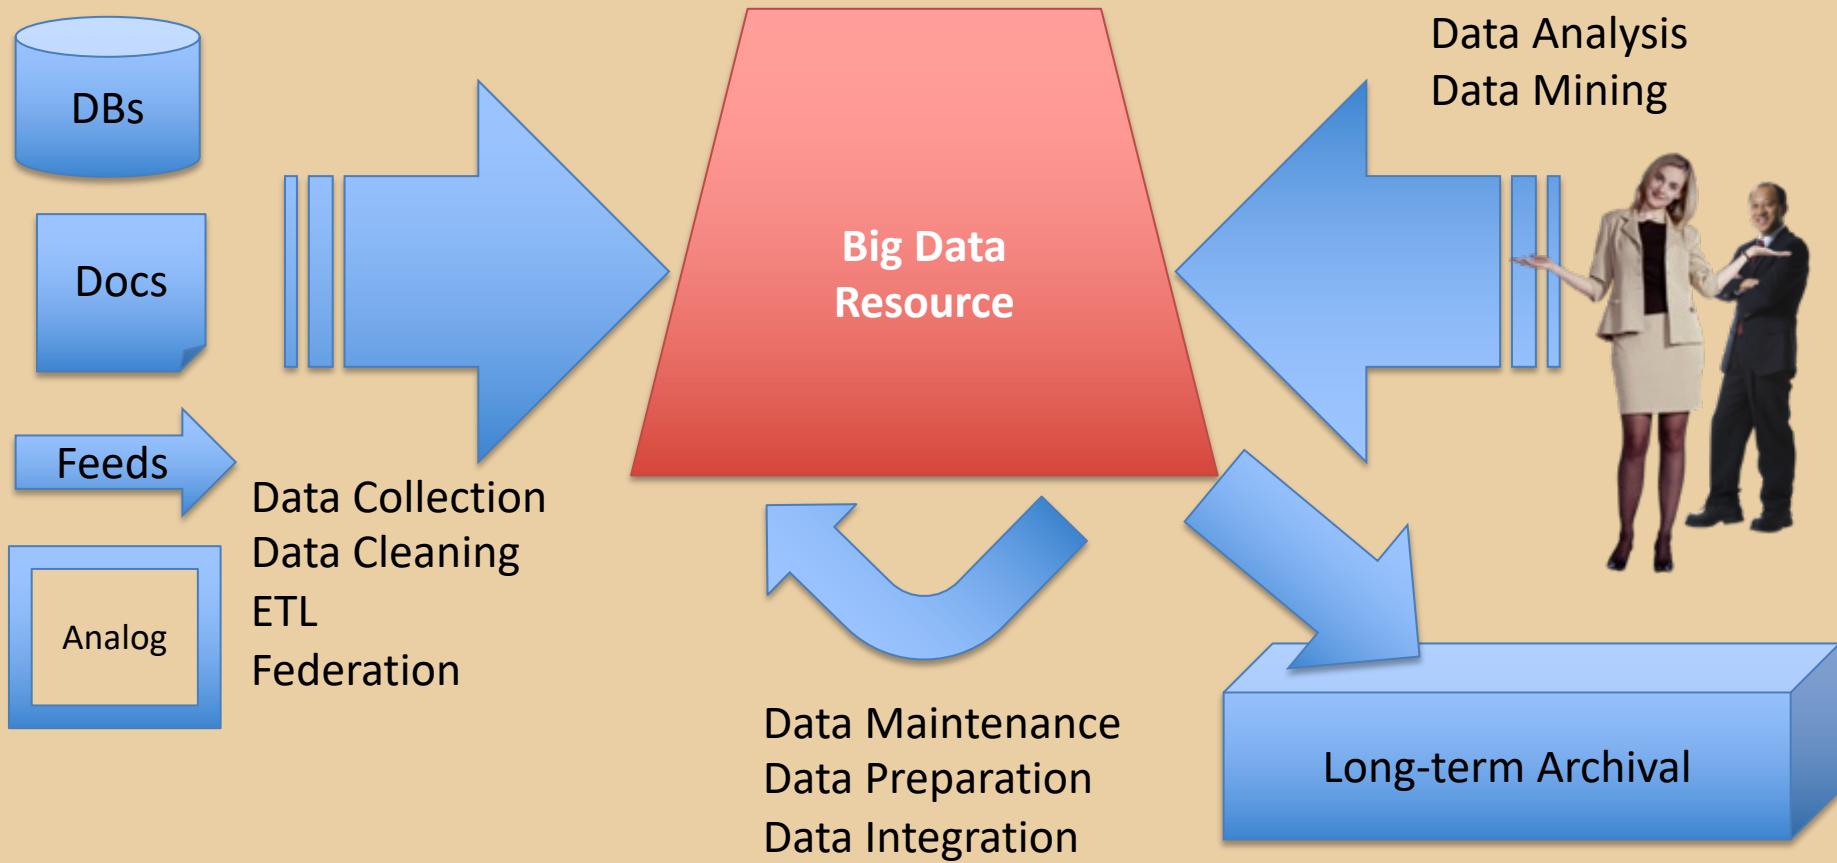
- Growing EO data (petabytes!)
  - Diverse EO applications
  - Growing heterogeneity



# High Data Heterogeneity







Big data is not a product, but a collection of **processes**

# Big Data and ~~Three~~ **Five** “V”s

**Volume**



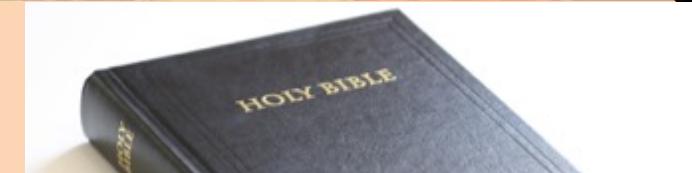
**Velocity**



**Variety**



**Veracity**



**Value**



Space is big. You just won't believe how vastly, hugely, mind-bogglingly big it is.

I mean, you may think it's a long way down the road to the chemist's, but that's just peanuts to space.

*Douglas Adams  
The Hitch Hiker's Guide to the Galaxy*

# Volume





100 million photos

---

Volume



Image credit: [https://www.123rf.com/photo\\_25981580/a-collection-of-travel-photos-from-around-the-world.html](https://www.123rf.com/photo_25981580/a-collection-of-travel-photos-from-around-the-world.html)

“Yesterday”:  
10K images



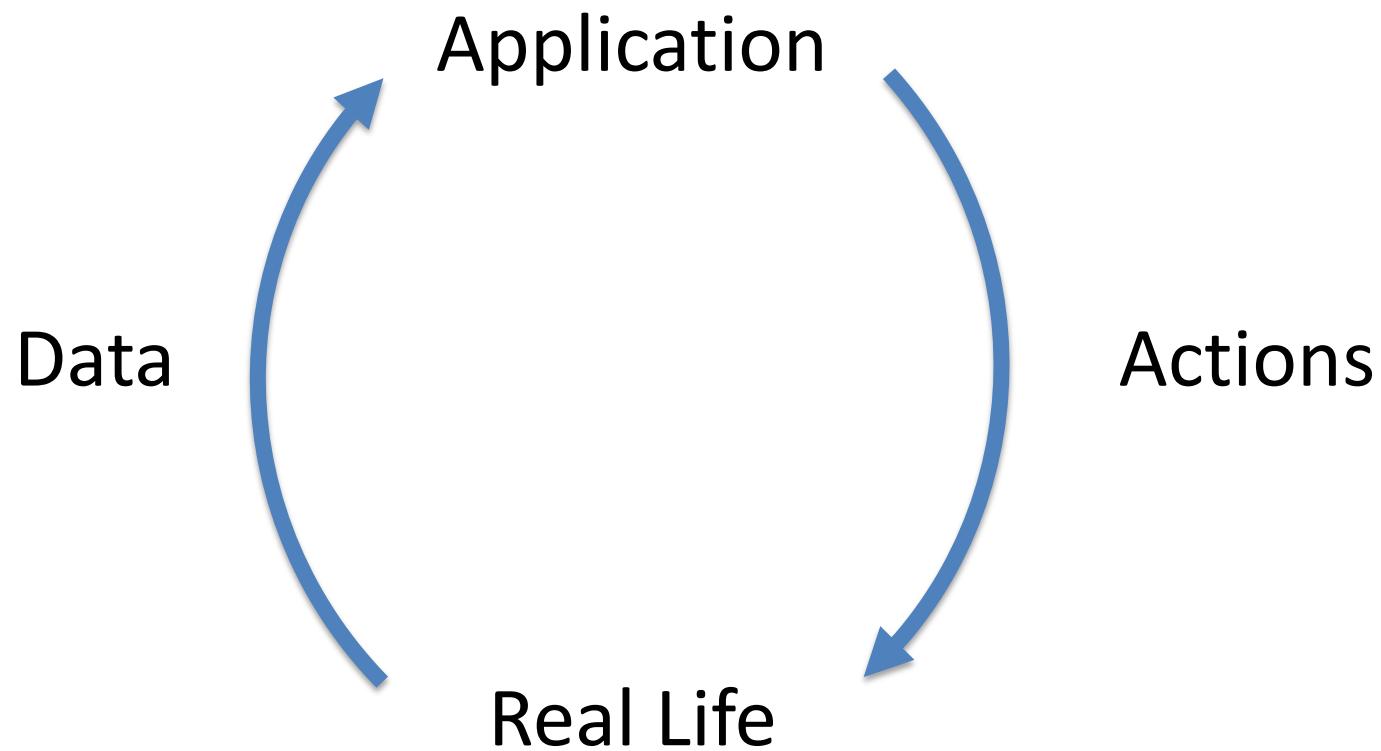
Today:  
**100M images**



Volume



# Two Sides of Velocity



Velocity



- Structured Data
  - Business data in RDBMSs, metadata (often)
  - Highly regular and repeating patterns
- **Semi-Structured Data**
  - Log data, literature (books), comment threads (?)
  - Some structure, but also much data w/o structure
- Unstructured Data
  - Multimedia content, free text (?), comment threads (?)
  - No (or little) structure
  - Try to extract some structure from this data to make it semi-structured!

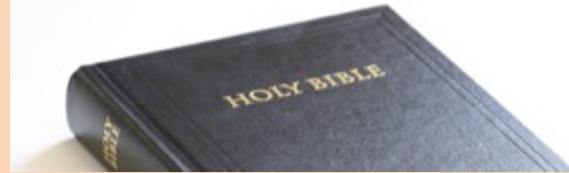
# Variety



# The Blind Men and the Elephant

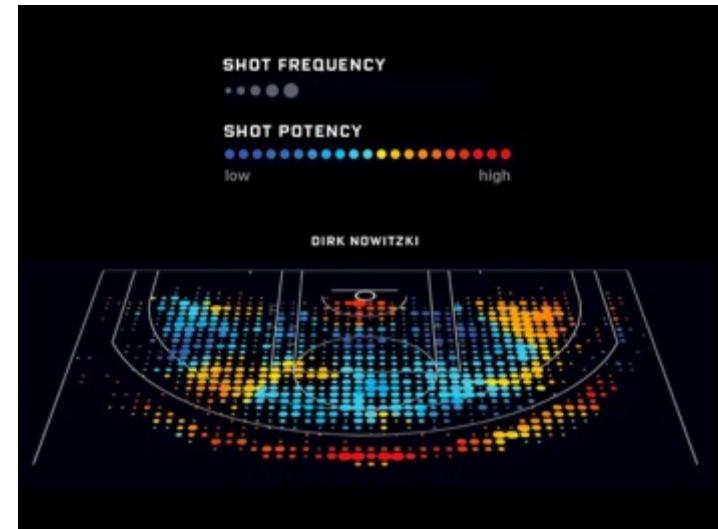
Six blind men were asked to determine what an elephant looked like by feeling different parts of the elephant's body. The blind man who feels a leg says the elephant is like a **pillar**; the one who feels the tail says the elephant is like a **rope**; the one who feels the trunk says the elephant is like a **tree branch**; the one who feels the ear says the elephant is like a **hand fan**; the one who feels the belly says the elephant is like a **wall**; and the one who feels the tusk says the elephant is like a **solid pipe**.

[http://en.wikipedia.org/wiki/Blind\\_men\\_and\\_an\\_elephant](http://en.wikipedia.org/wiki/Blind_men_and_an_elephant)

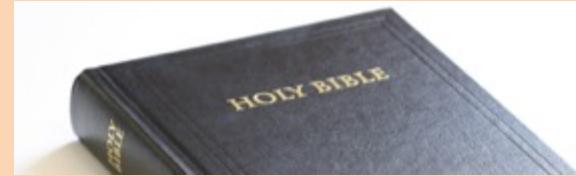


# Know Your Data

- Take time to understand the **whole** data
  - Value ranges, distributions
  - Plots, descriptive statistics, ...
- Estimate the outcome of your analysis
  - Best way to know whether the final outcome is correct



Veracity



# Data is of Low Quality!

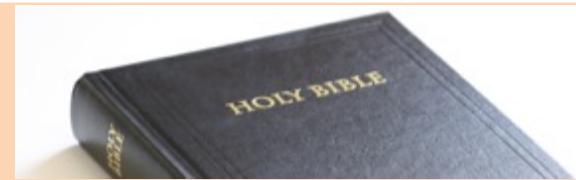
## Impact of poor data quality?

- Erroneous data costs US businesses \$600 billion/year
- In analytics projects, data cleaning takes 30-80% of time and budget
- Data quality tools market is growing at 16% annually, way over the 7% average for other IT segments

## How much data is erroneous?

- Enterprise data error rates: average of 1-5%, some > 30%
- Only 1/3<sup>rd</sup> of XML Web documents with XSD/DTD are valid, 14% even lack well-formedness

Veracity



# Two Sides of Value



Value



# Lecture Outline

**Learning Outcome 6:** Discuss the pros and cons of different classes of data systems for modern analytics and data science applications.

- What are modern analytics and data science applications?
- What are their access patterns and processing requirements?
- How well/badly do the classes of data systems covered so far support these access patterns?
- MapReduce/Spark: Their properties and their support for these access patterns

# Access Patterns to Data on Disk

- Data collections typically consist of large files of semi-structured data
- Too large for a single server
  - **distributed storage**
- Read the whole collection
  - **large sequential scans**
    - Typically in any order
    - Sometimes repeatedly (e.g., clustering)

# Processing Requirements

- Apply filters to reduce data quantity
  - May be similar to SQL or based on regular expressions
- Run complex processing pipelines
  - Filtering + algorithms, sometimes loops
- Too large for a single server  
→ **distributed processing**

# Lecture Outline

**Learning Outcome 6:** Discuss the pros and cons of different classes of data systems for modern analytics and data science applications.

- What are modern analytics and data science applications?
- What are their access patterns and processing requirements?
- How well/badly do the classes of data systems covered so far support these access patterns?
- MapReduce/Spark: Their properties and their support for these access patterns

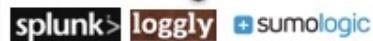
# The Big Data Landscape

## Apps

### Vertical



### Operational Intelligence



### Ad/Media



### Business Intelligence



### Analytics and Visualization



### Data As A Service



## Infrastructure

### Analytics



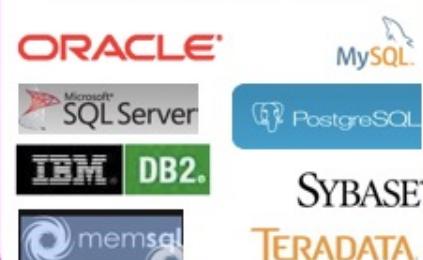
### Operational



### As A Service



### Structured DB



hadoop

hadoop  
mapReduce

### Technologies

mahout

HBASE

Cassandra

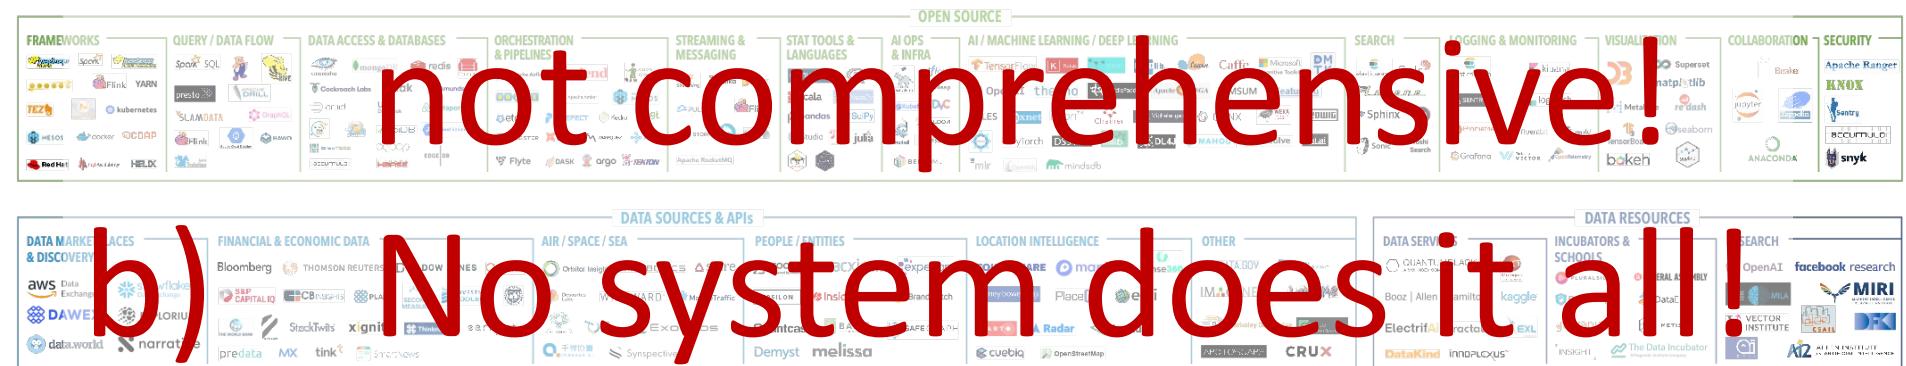
# Big Data Landscape 2016





a) This discussion is not comprehensive!

b) No system does it all



# RDBMSs (OLTP and OLAP)

## Strengths

- Transactions
- Complex query processing
  - Especially join processing
  - ... but also filtering data

## Weaknesses

- Neither strength is relevant to big data
  - ... except for filtering
- Little support for unstructured data
- Little support for distributed processing

# NoSQL: Key-Value Stores

## Strengths

- Distributed storage
- Short latency for single items
- Handles any data equally well/badly

## Weaknesses

- No means to scan the entire collection
- No support for filtering data
- No support for distributed processing pipelines

# NoSQL: Document Stores

## Strengths

- Distributed storage
- Short latency for single documents
- Deals with semi-structured data
- Support for filtering documents

## Weaknesses

- Not designed for scanning the entire collection
- No support for distributed processing pipelines

# NoSQL: Graph Stores

## Strengths

- Good for certain graph-related applications
- May be strongly relational or more NoSQL-like and inherit strengths of each...

## Weaknesses

- May be strongly relational or more NoSQL-like and inherit weaknesses of each...

# Lecture Outline

**Learning Outcome 6:** Discuss the pros and cons of different classes of data systems for modern analytics and data science applications.

- What are modern analytics and data science applications?
- What are their access patterns and processing requirements?
- How well/badly do the classes of data systems covered so far support these access patterns?
- MapReduce/Spark: Their properties and their support for these access patterns

# MapReduce and Hadoop

- ♦ **MapReduce:** A programming model

- ♦ Functional style
- ♦ Users specify Map and Reduce functions
- ♦ Expressive

- ♦ **Hadoop:** An associated implementation

- ♦ Automatically parallelized
- ♦ Automatic partitioning, execution, failure handling, load balancing, communication
- ♦ Can handle very large datasets using clusters of commodity computers

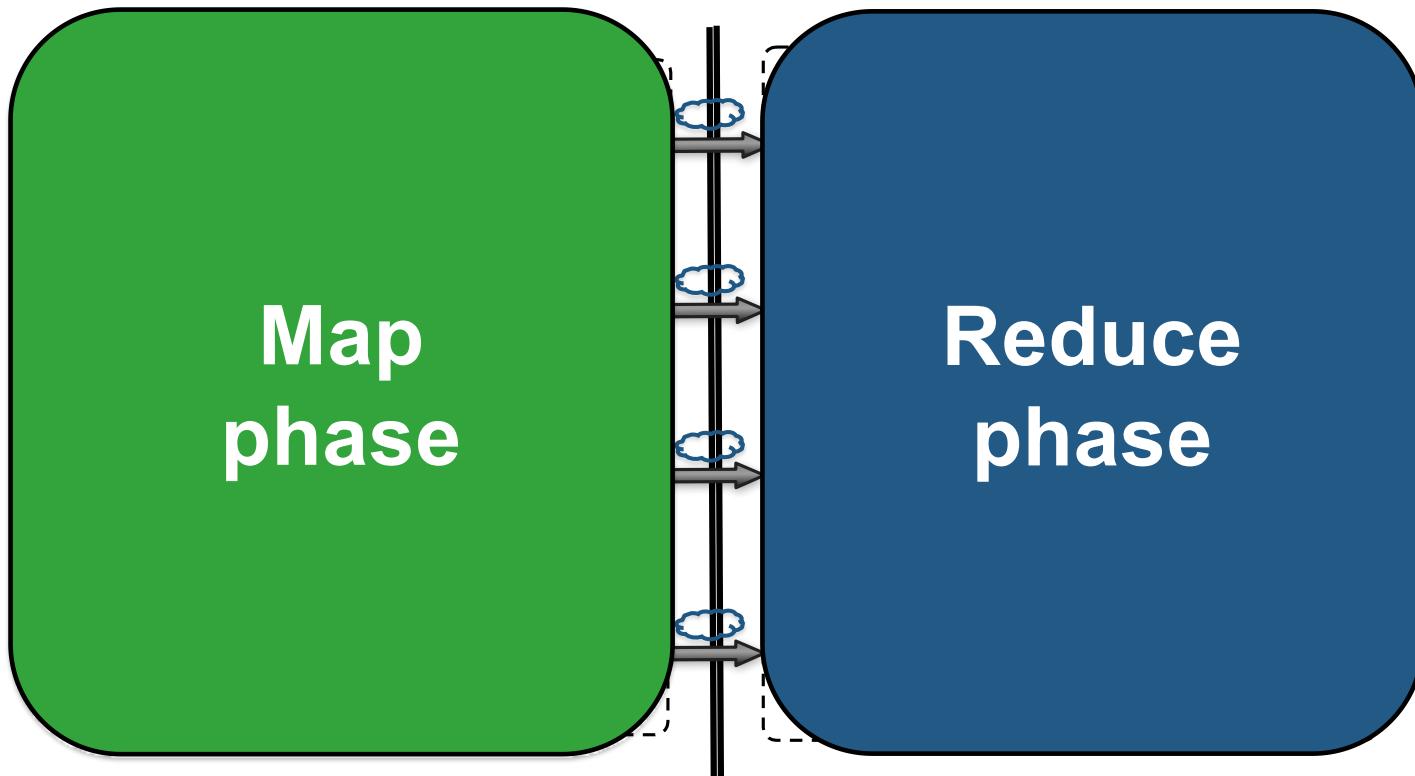
# MapReduce in a Nutshell (I)

- ♦ Framework
  - ♦ Read lots of data
  - ♦ **Map**: process a data item/record
  - ♦ **Sort and shuffle**
  - ♦ **Reduce**: aggregate, summarize, filter, transform
  - ♦ Write results
- ♦ Map and Reduce are **user-specified** → used to model given problem

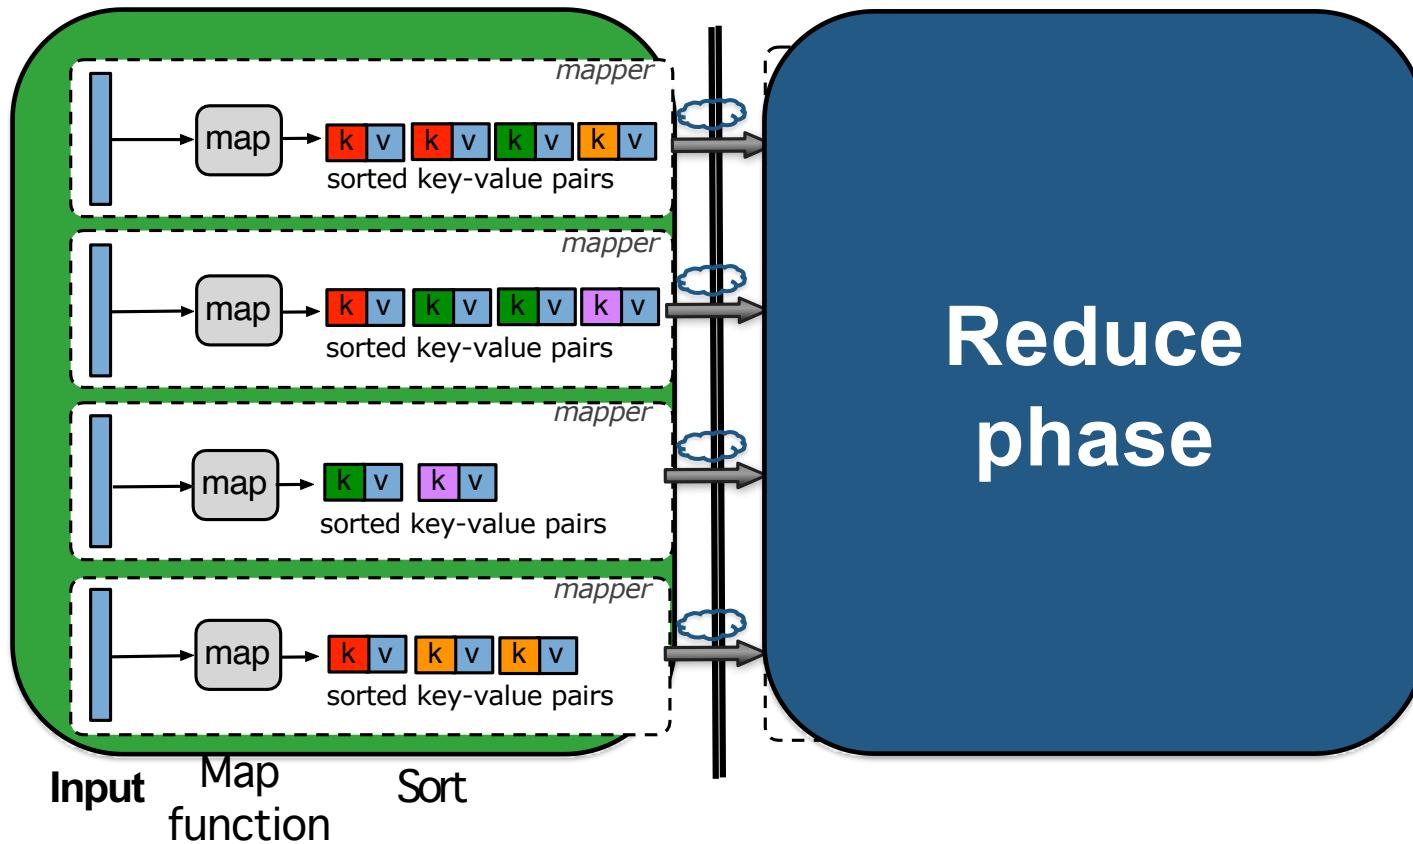
# MapReduce in a Nutshell (II)

- ♦ Simple API
  - ♦ **Map** (*key, value*) —> {*ikey, ivalue*}
  - ♦ **Reduce** (*ikey, {ivalue}*) —> (*key', value'*)
- ♦ **Map phase**: independent processes (mappers) which run **in parallel**
  - ♦ operate on the chunks (blocks) of input data
  - ♦ output intermediate results
- ♦ **Shuffle phase**: Intermediate results are shuffled through the network
- ♦ **Reduce phase**: independent processes (reducers) which run **in parallel**
  - ♦ group intermediate results of the map phase
  - ♦ operate on the groups
  - ♦ output final results

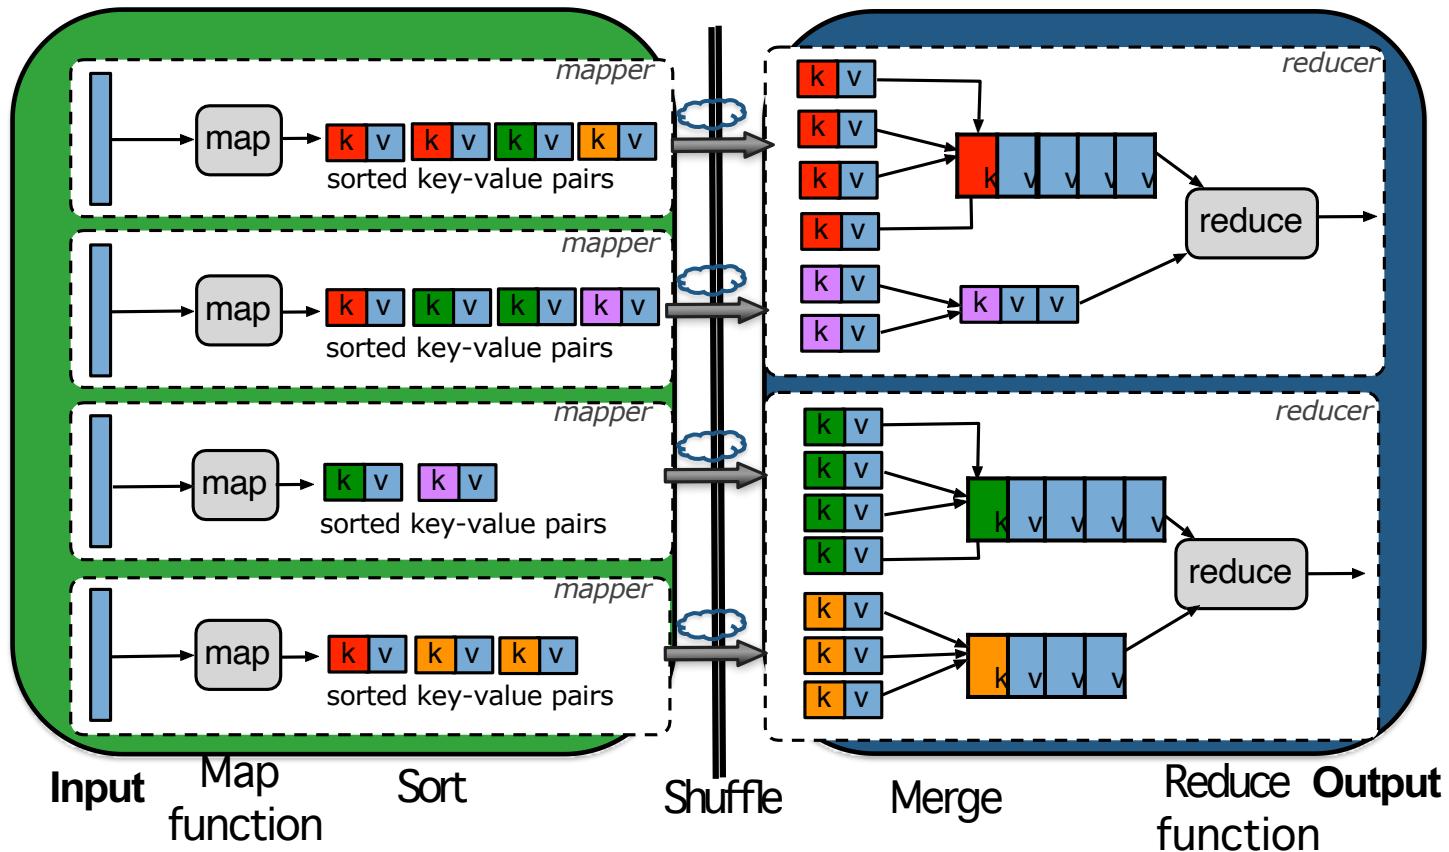
# MapReduce Illustration



# Map Reduce Illustration



# Map Reduce Illustration



# Wordcount Example

**Map** (*key, value*)  $\rightarrow \{i\text{key}, i\text{value}\}$

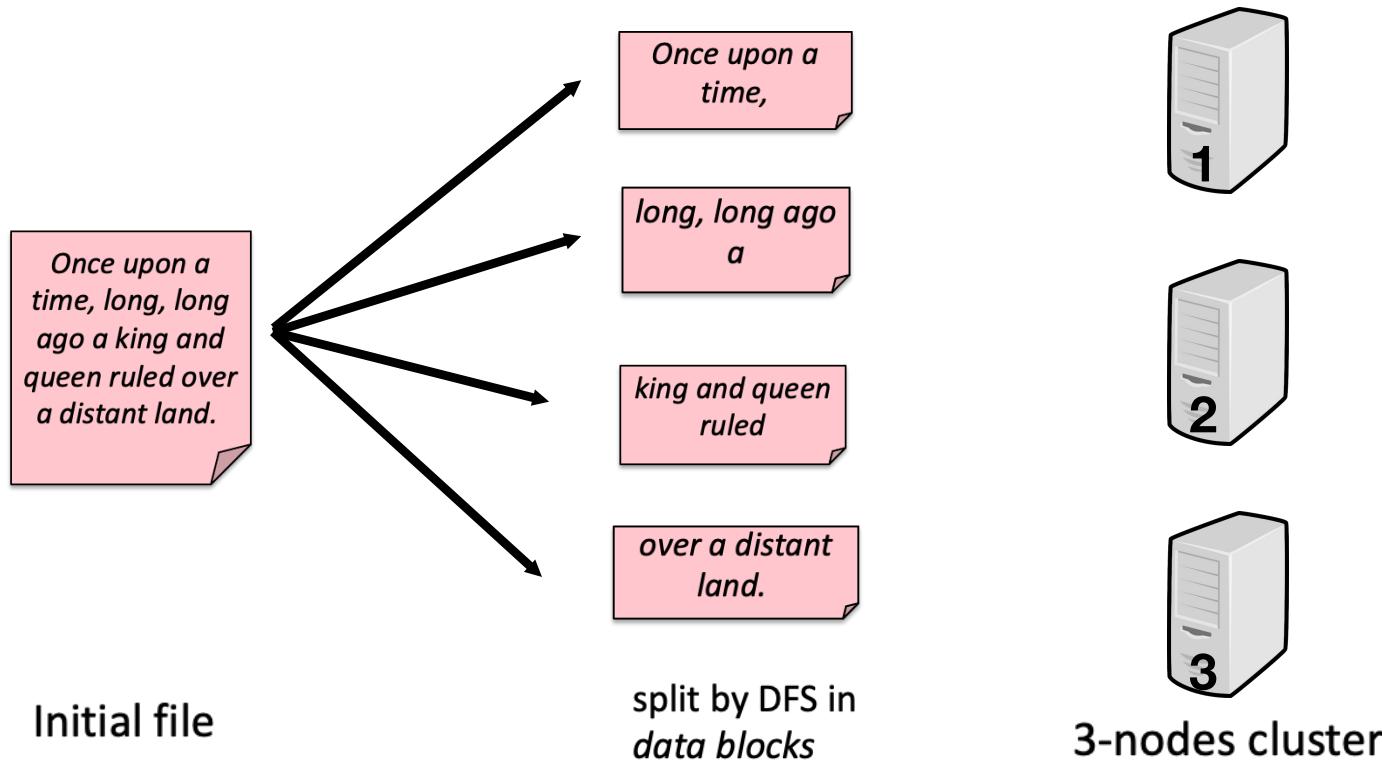
**Reduce** (*ikey, {ivalue}*)  $\rightarrow (\text{key}', \text{value}')$



**Map** (*docID, text*)  $\rightarrow \{\text{word}, 1\}$

**Reduce** (*word, {1, 1, ...}*)  $\rightarrow (\text{word}, \text{count})$

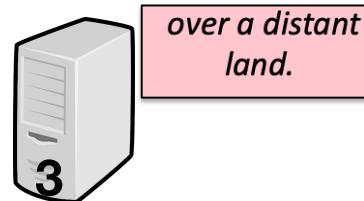
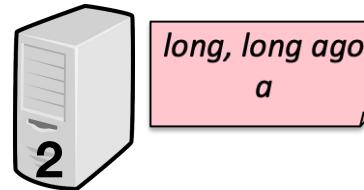
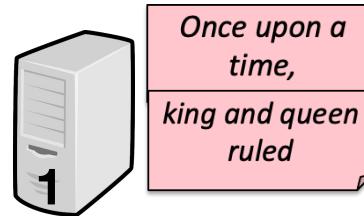
# Wordcount Example



# Wordcount Example

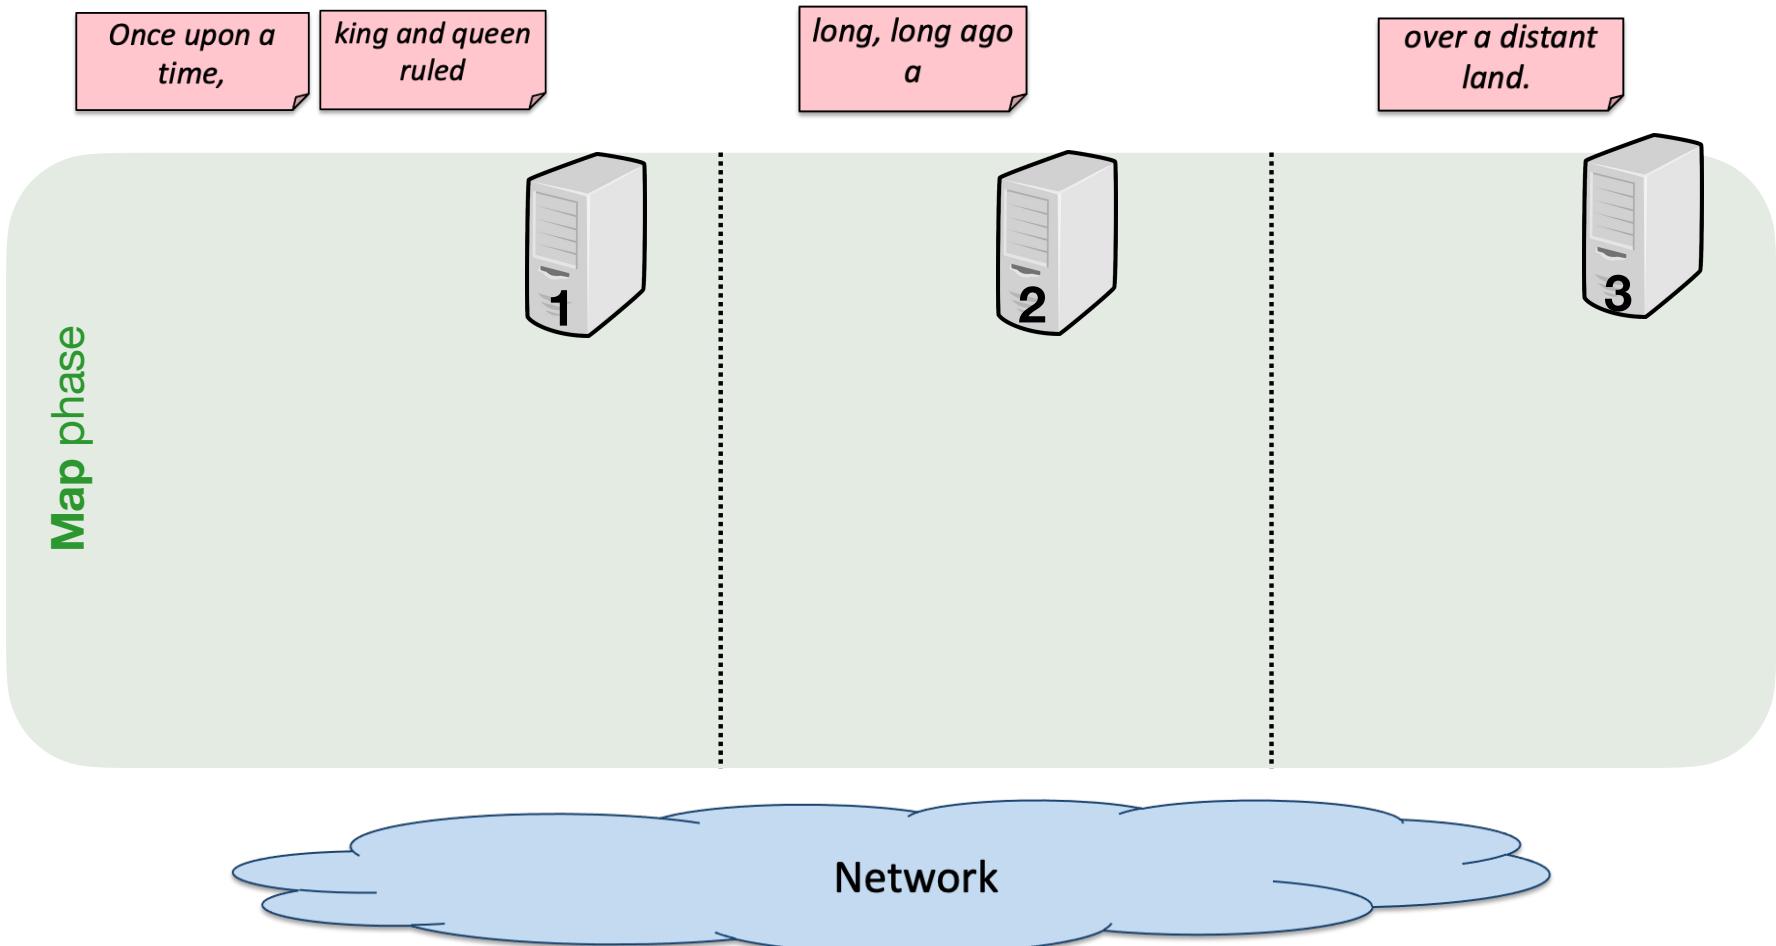
*Once upon a time, long, long ago a king and queen ruled over a distant land.*

Initial file

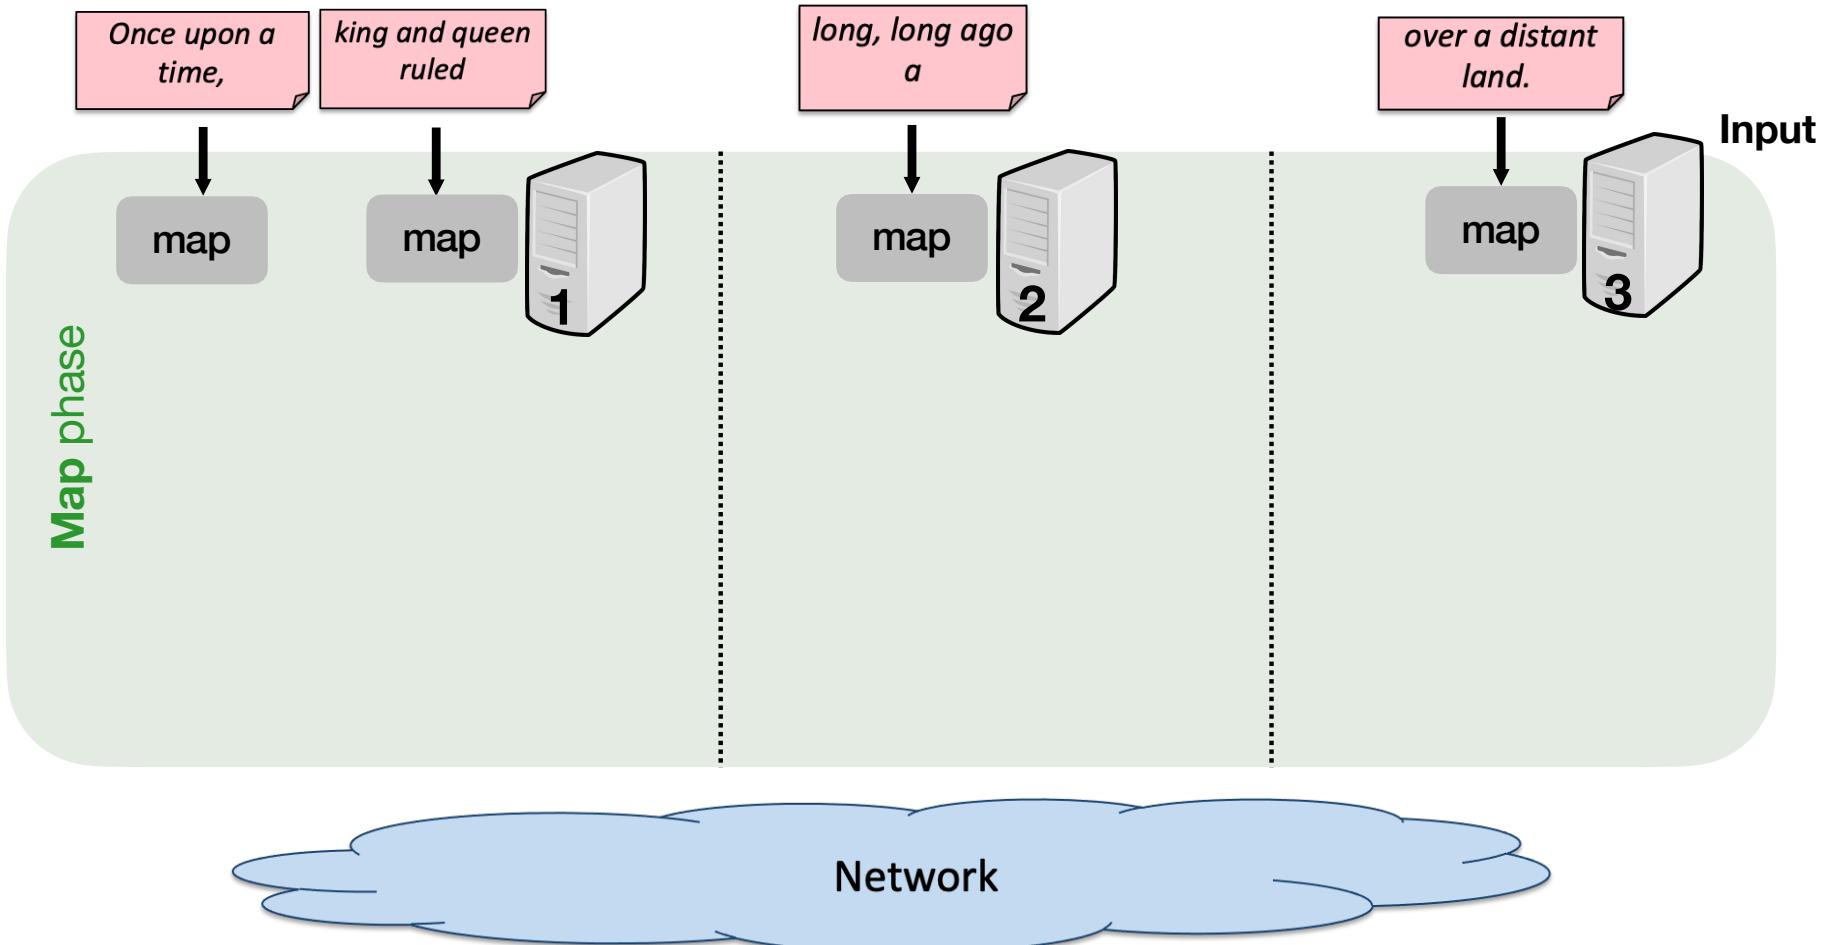


3-nodes cluster

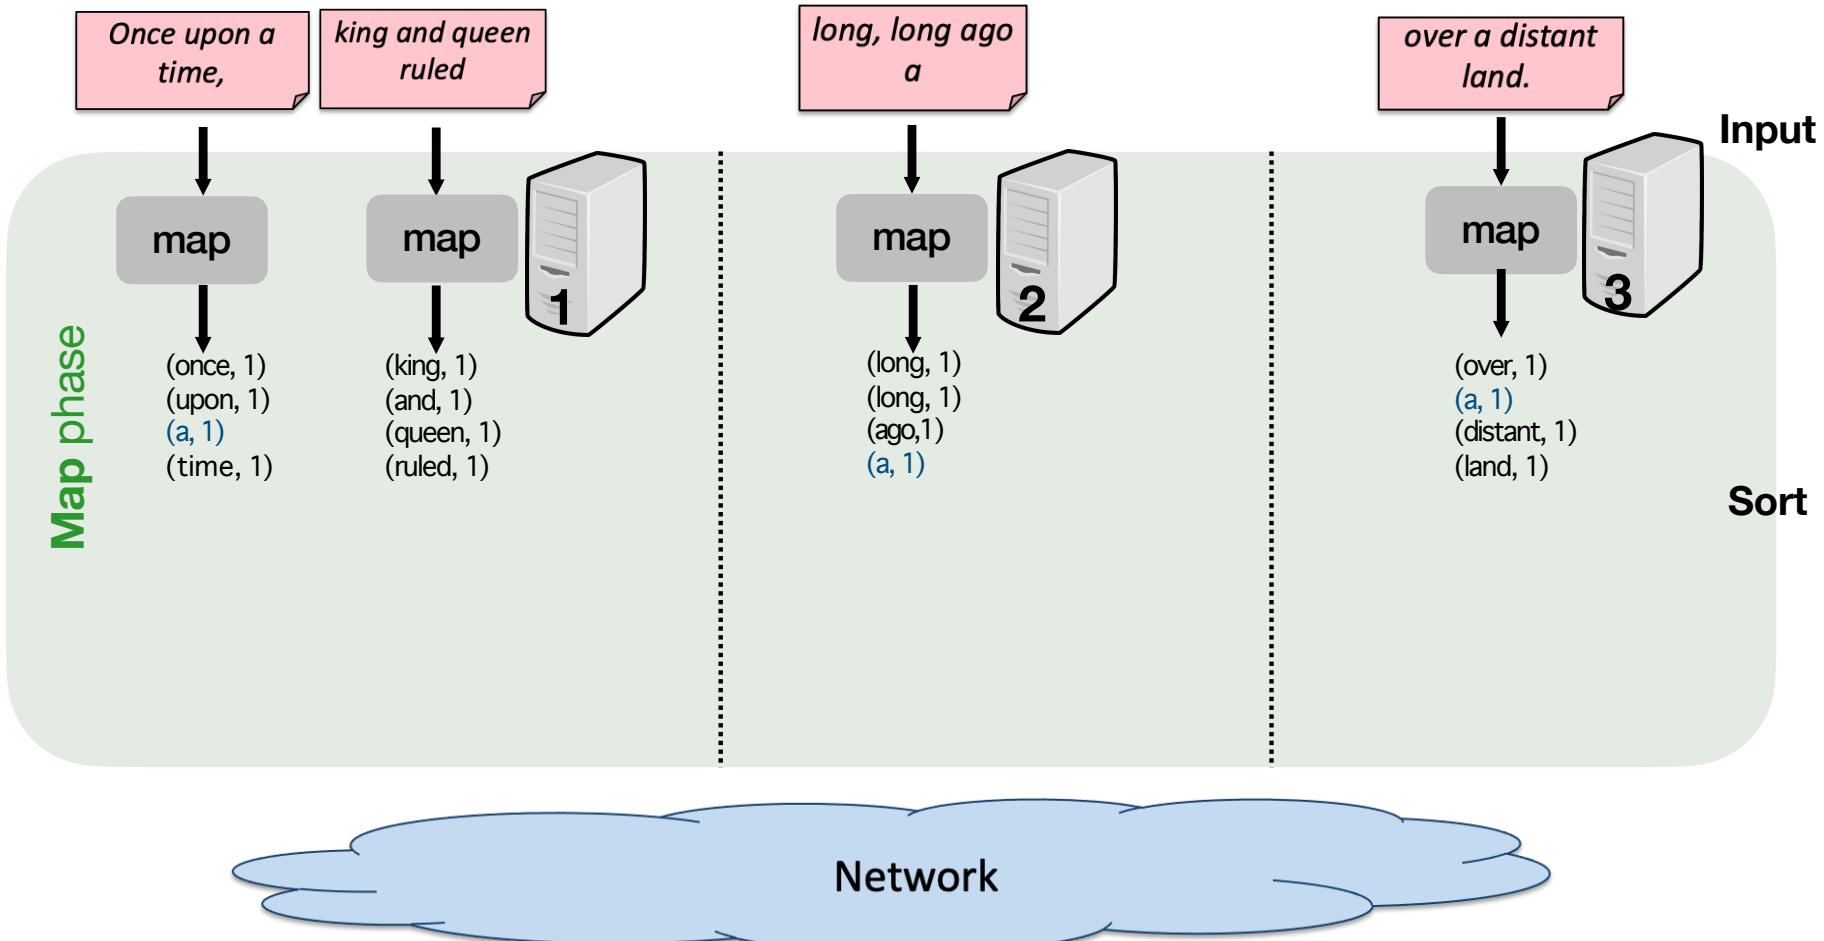
# Wordcount Example (Map Phase)



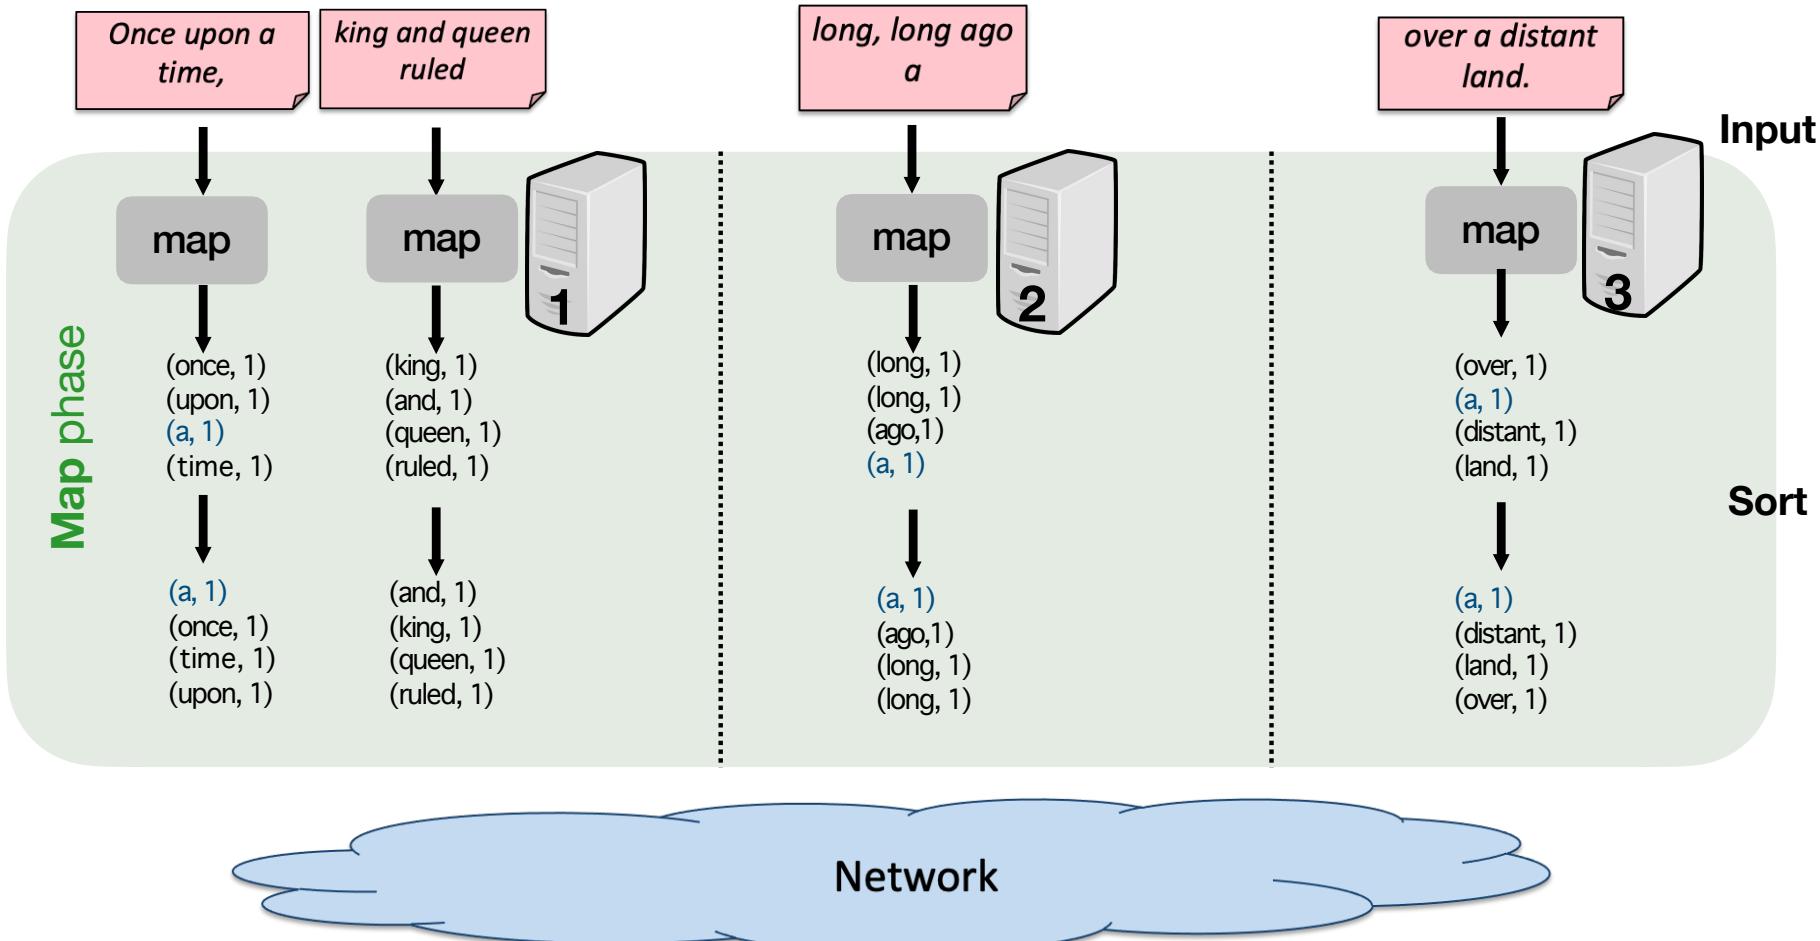
# Wordcount Example (Map Phase)



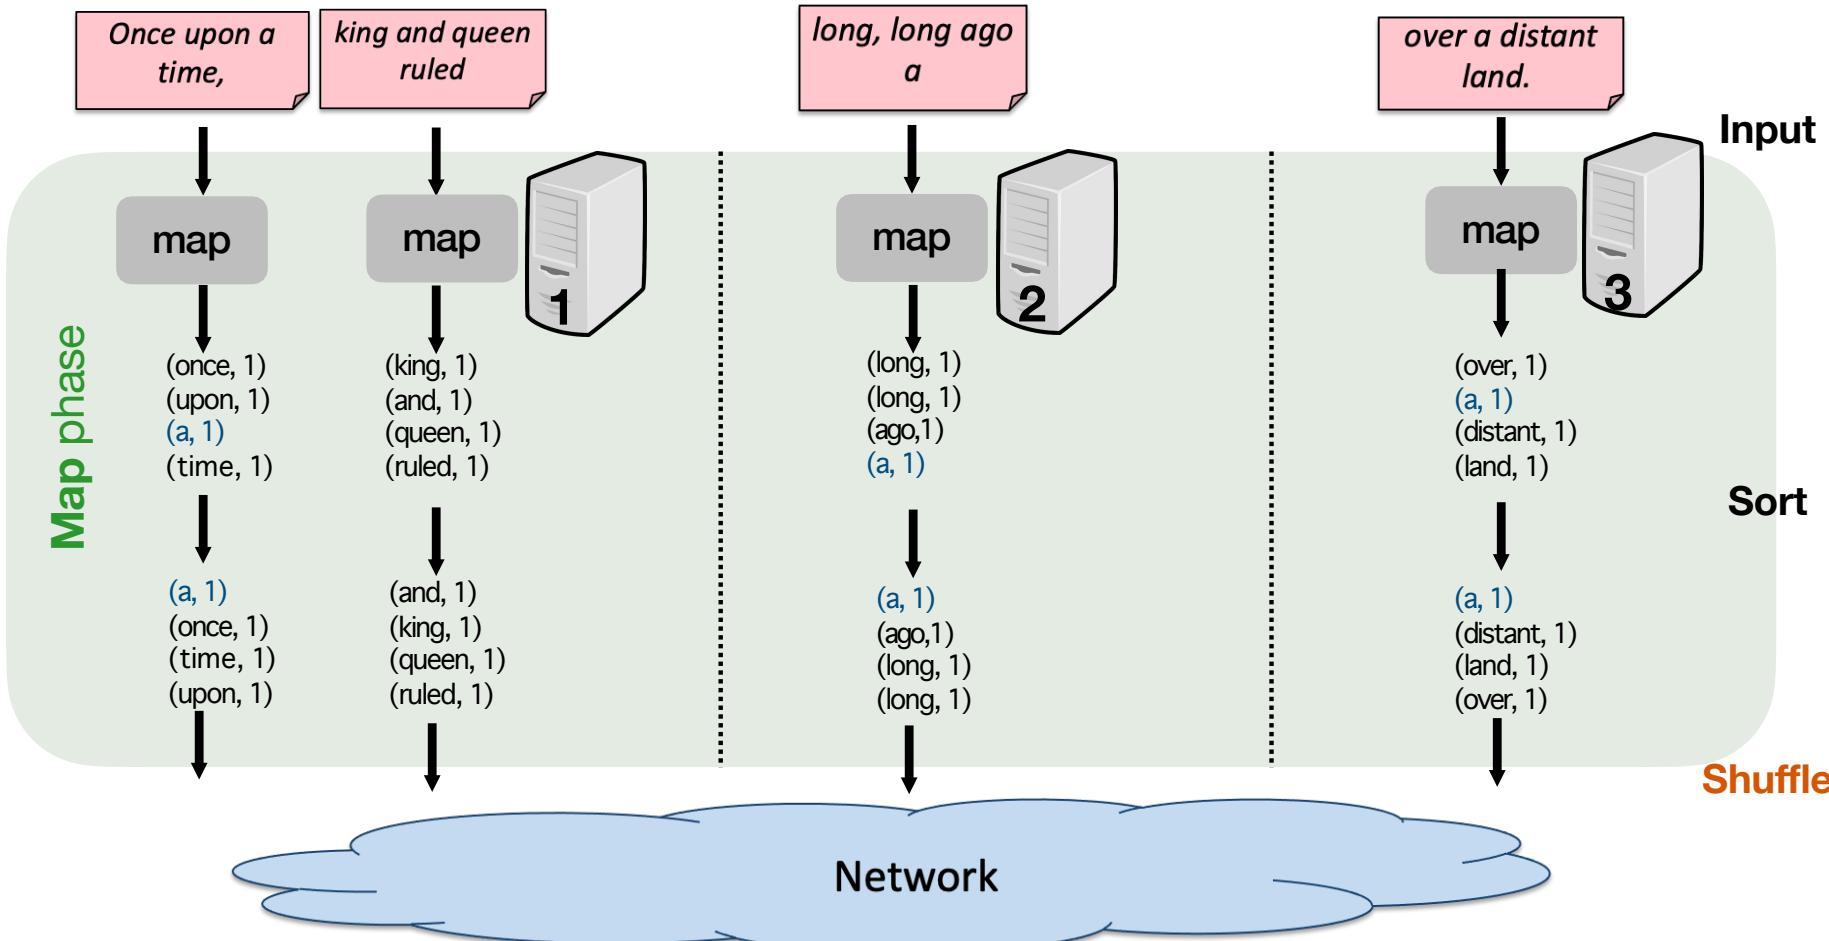
# Wordcount Example (Map Phase)



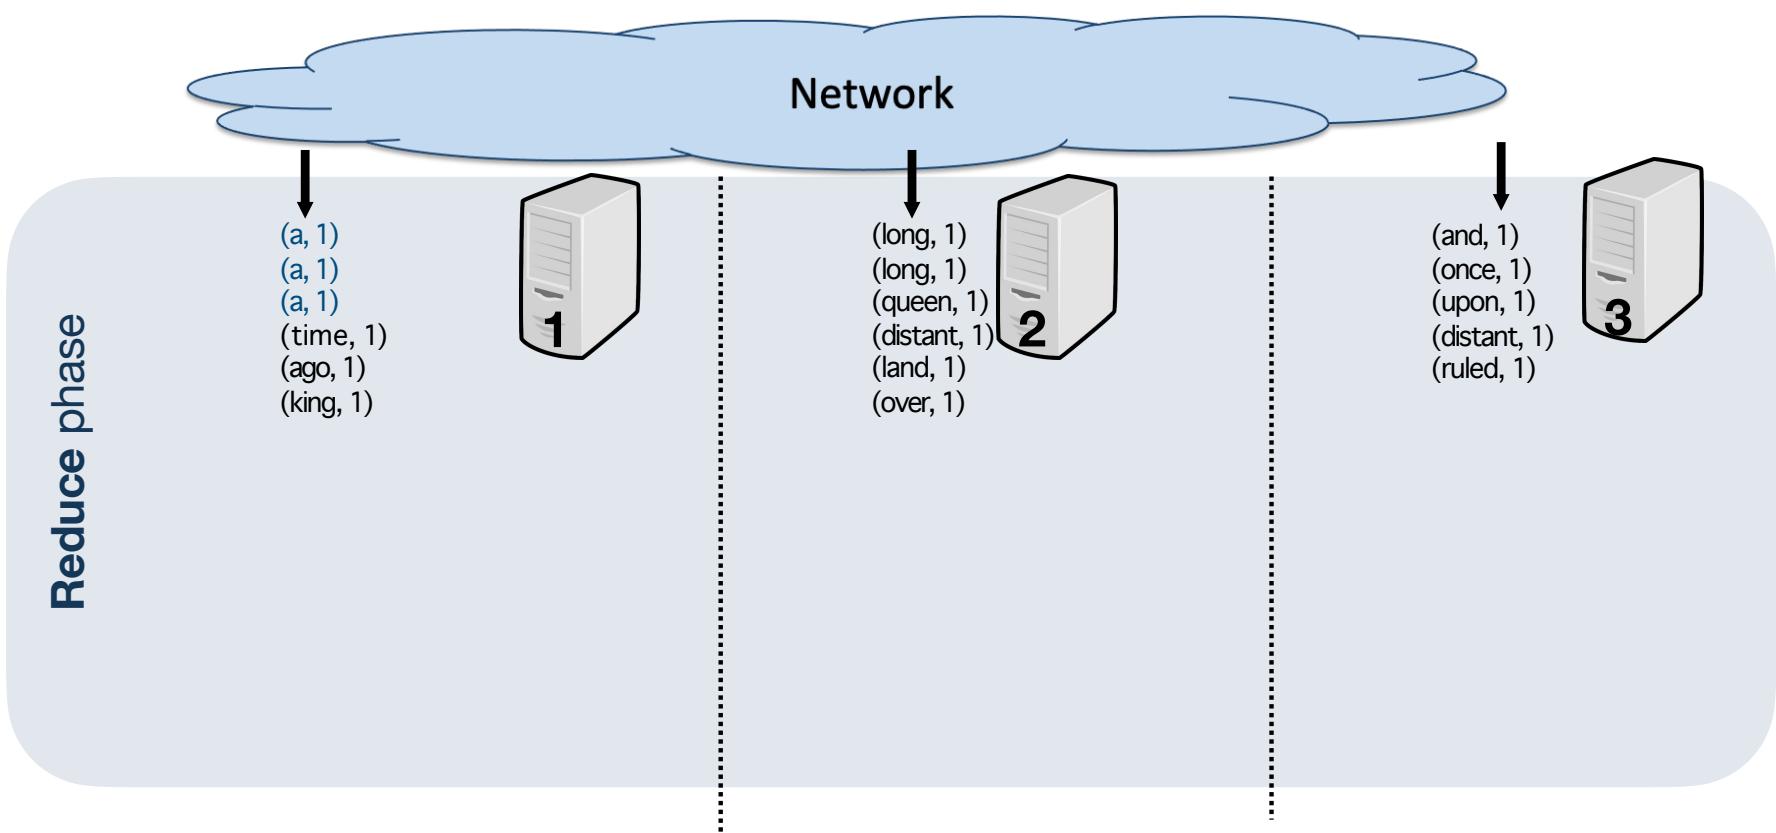
# Wordcount Example (Map Phase)



# Wordcount Example (Map Phase)

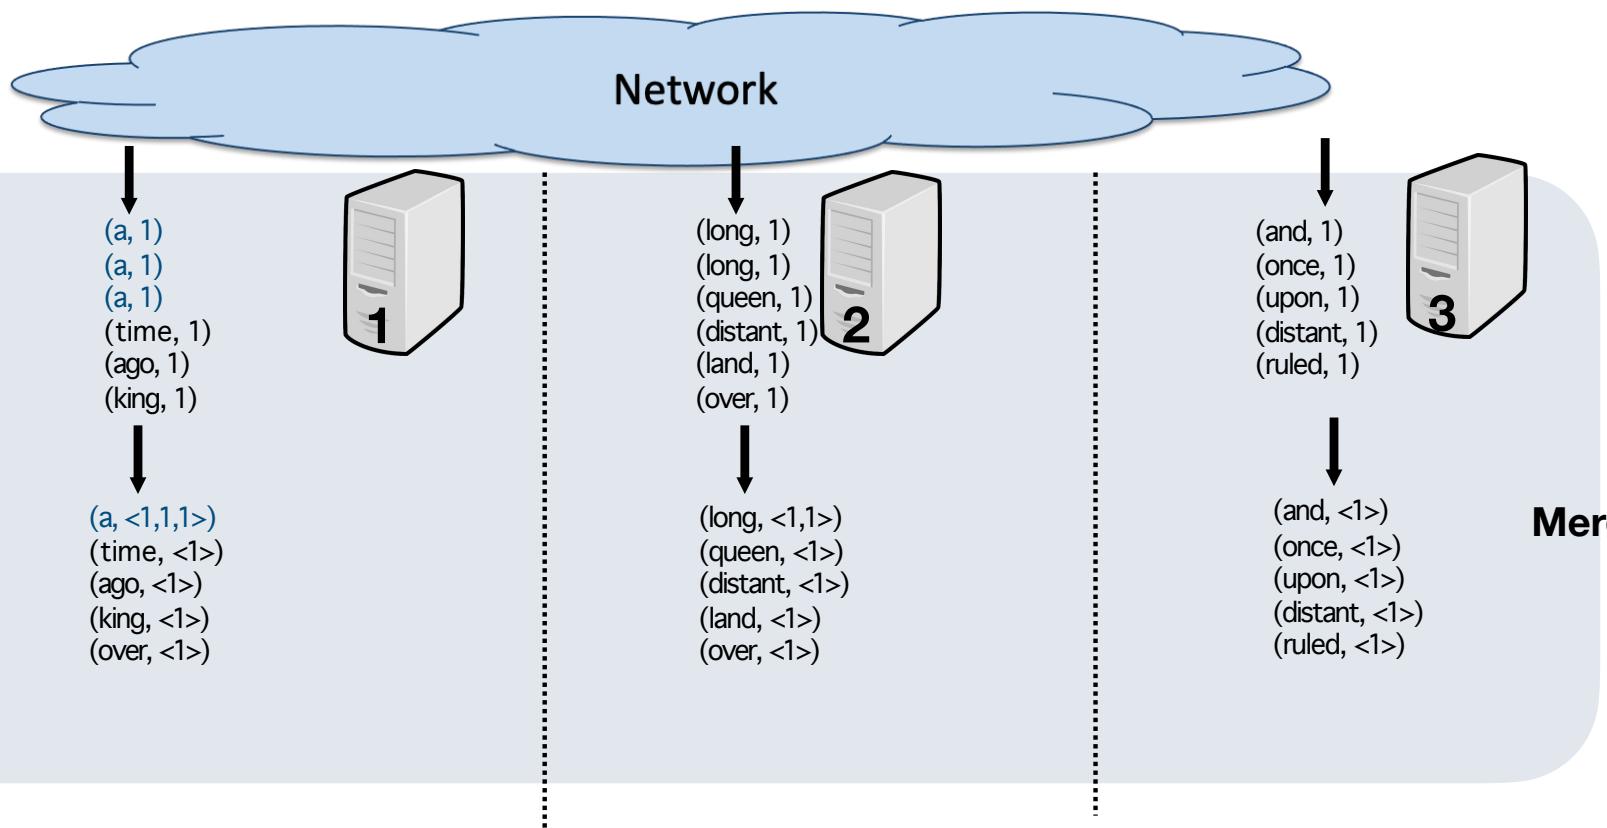


# Wordcount Example (Reduce Phase)

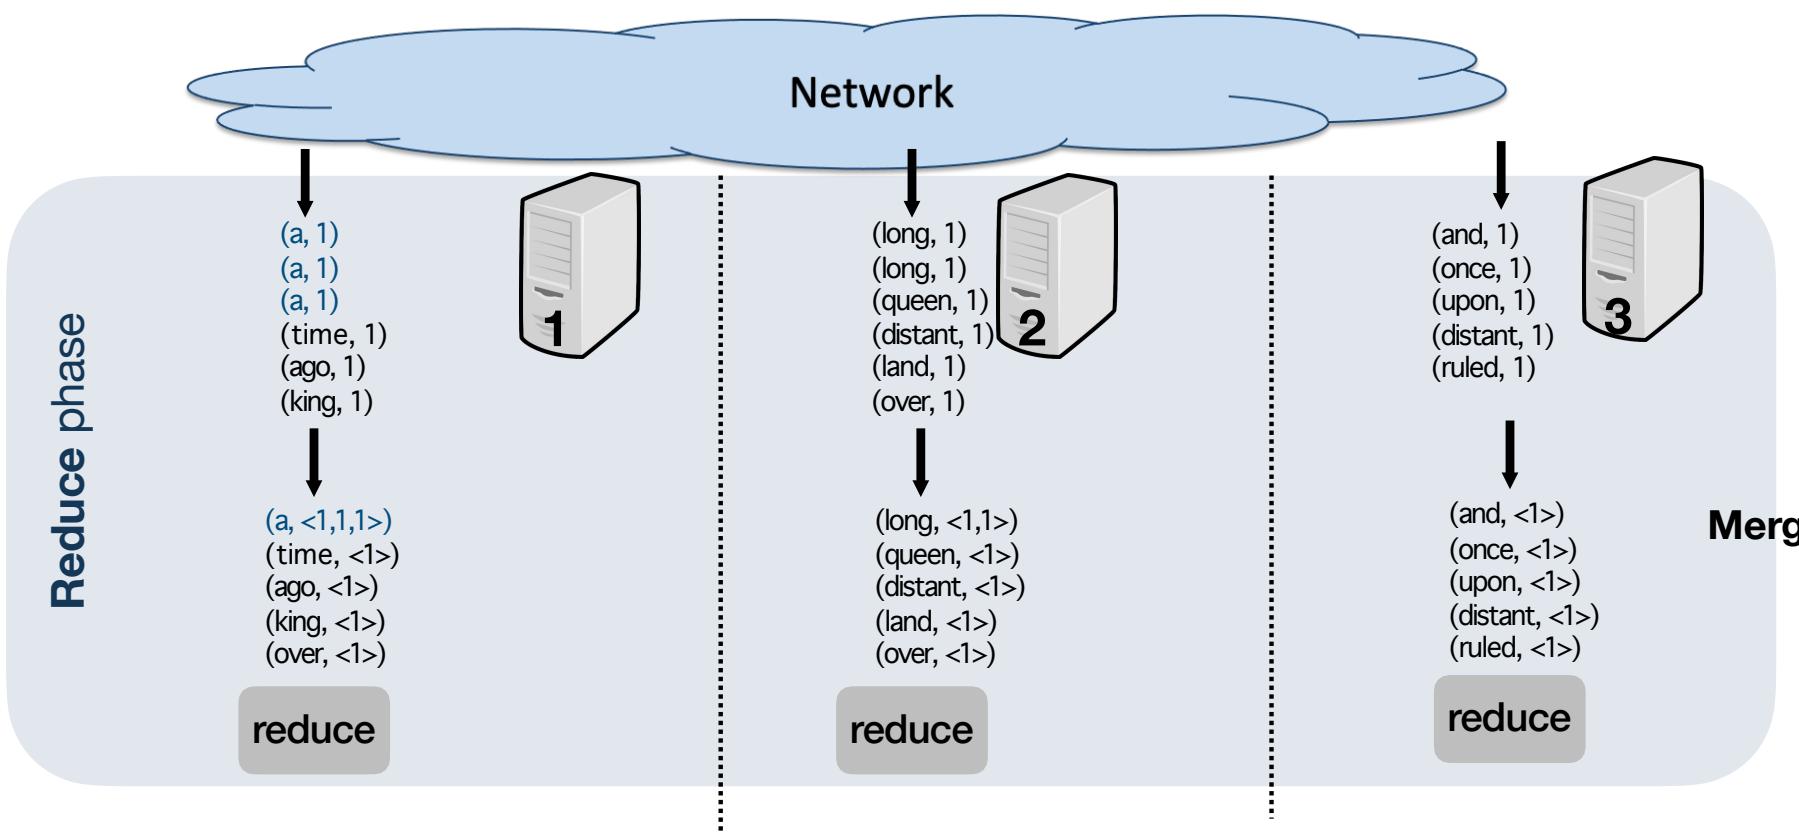


# Wordcount Example (Reduce Phase)

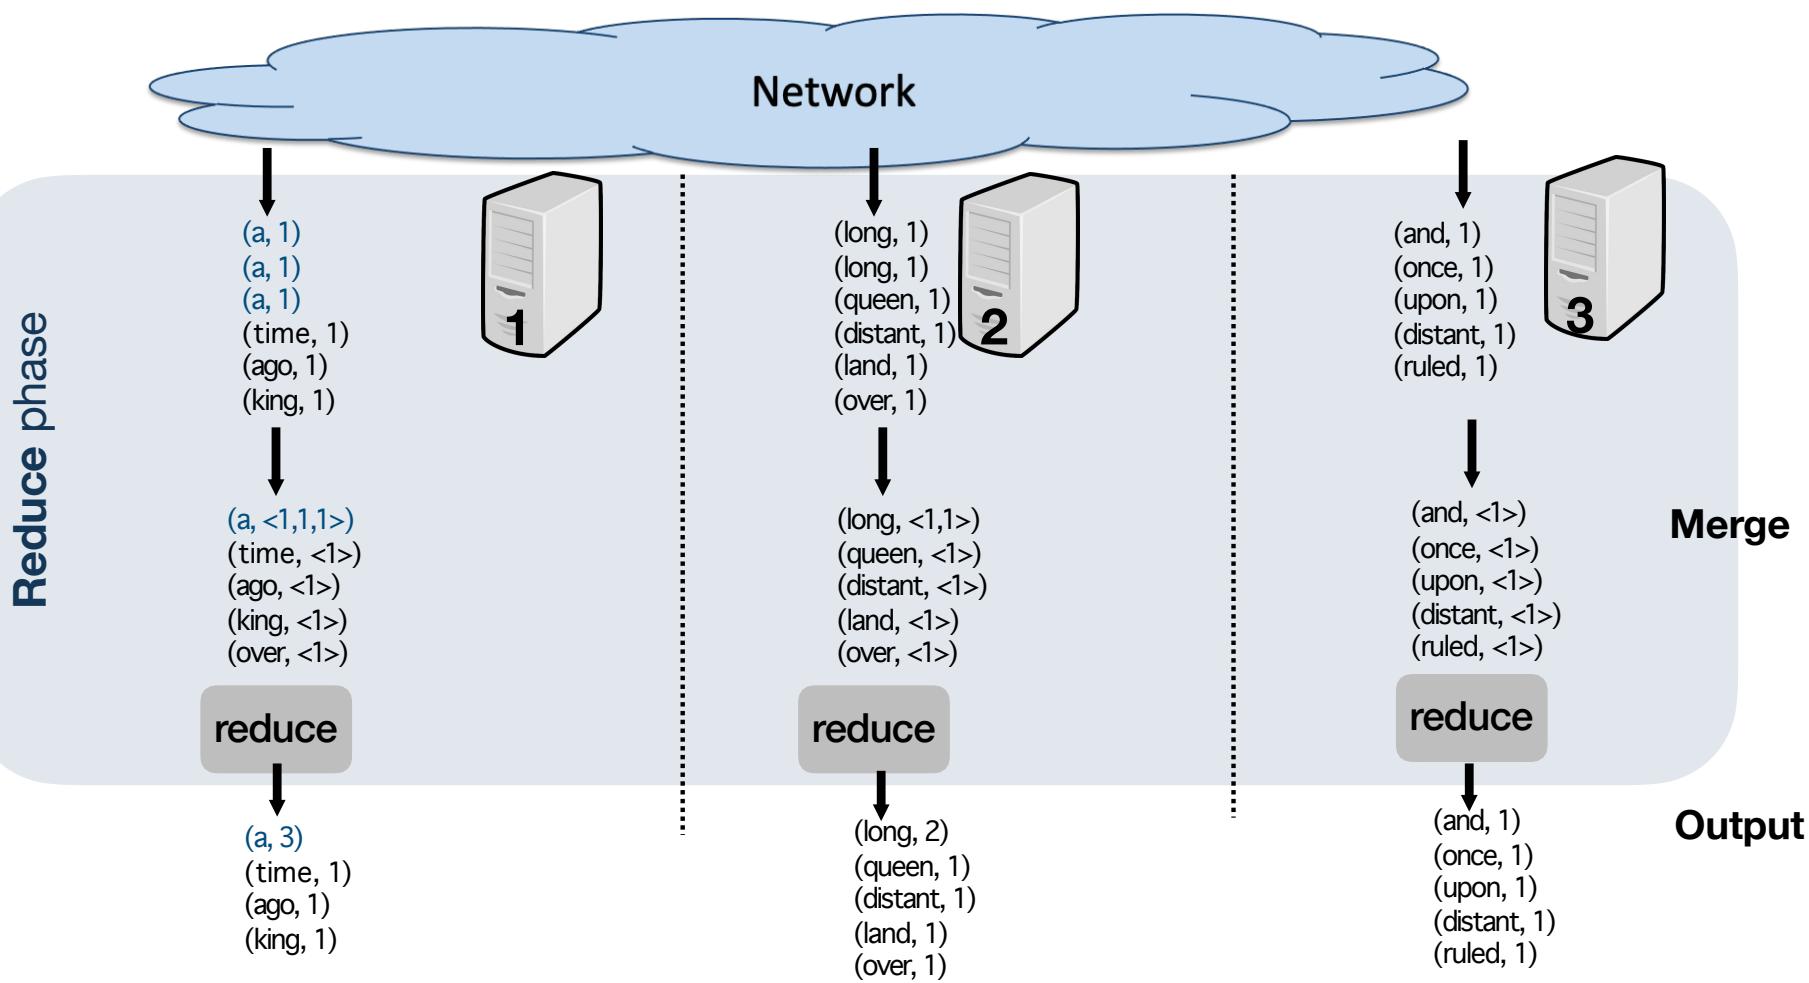
Reduce phase



# Wordcount Example (Reduce Phase)



# Wordcount Example (Reduce Phase)



# Wordcount Example: Pseudocode

**Map**(Integer key, String text)

**foreach** word w **in** text:  
        **emit** (w, 1)

**Reduce**(String word, Iterator<Integer> counts)

**int** result = 0  
    **foreach** count **in** counts:  
        result += count  
    **emit** (word, result)

# Hadoop Problems

- Containers prevent sharing resources
  - Large data structures → one per thread!
- Rigid MapReduce pipeline
  - Low-level abstraction
  - Map → Map → Reduce → Map → Reduce



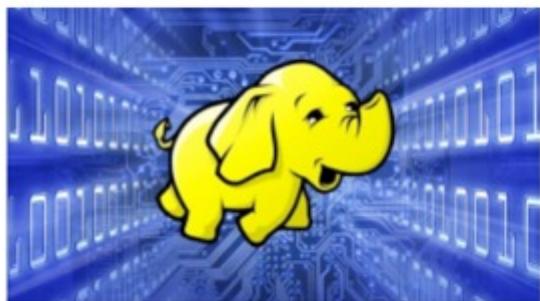
- Limited successes

# Hadoop Critique

March 13, 2017

## Hadoop Has Failed Us, Tech Experts Say

Alex Woodie



The Hadoop dream of unifying data and compute in a distributed manner has all but failed in a smoking heap of cost and complexity, according to technology experts and executives who spoke to *Datanami*.

"I can't find a happy Hadoop customer. It's sort of as simple as that," says Bob Muglia, CEO of [Snowflake Computing](#), which develops and runs a

cloud-based relational data warehouse offering. "It's very clear to me, technologically, that it's not the technology base the world will be built on going forward."

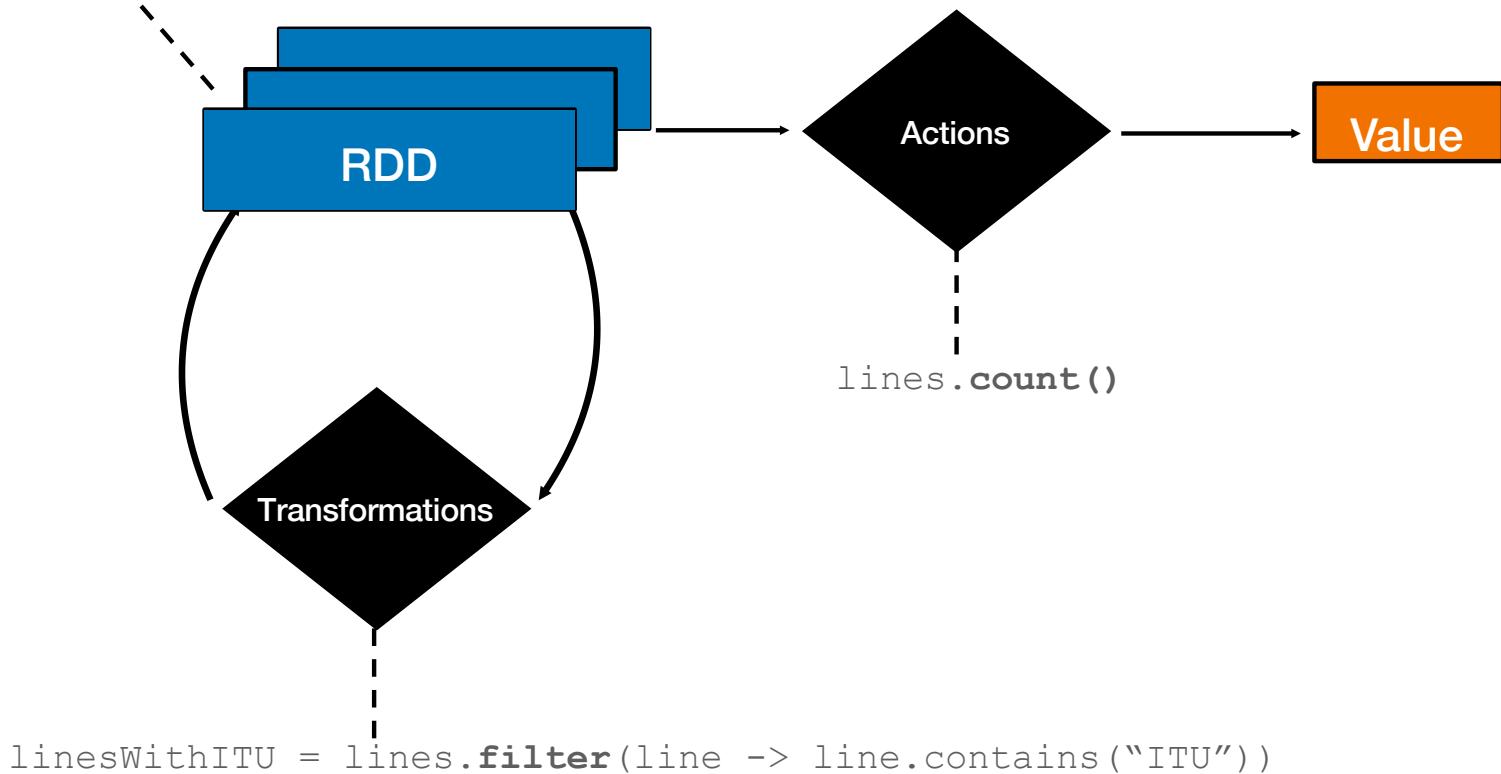
Hadoop's strengths lie in serving as a cheap storage repository and for processing ETL batch workloads, Johnson says. But it's ill-suited for running interactive, user-facing applications, he says.

# Spark RDDs

- Resilient Distributed Datasets (RDDs)
  - Transform one RDD to another via operators
  - Lazy execution – optimizations
- Supports deep pipelines
- Supports worker's memory sharing
- Write programs in terms of distributed datasets and operations on them

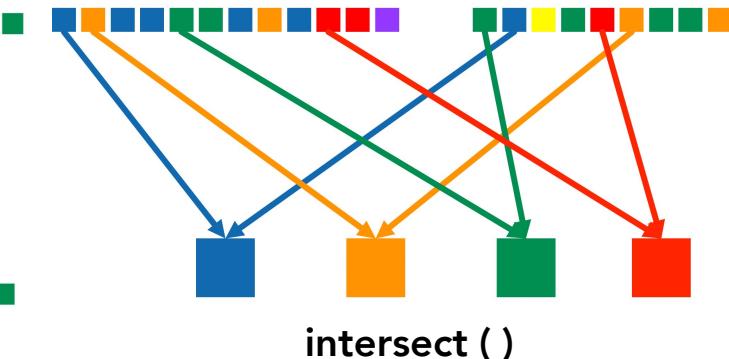
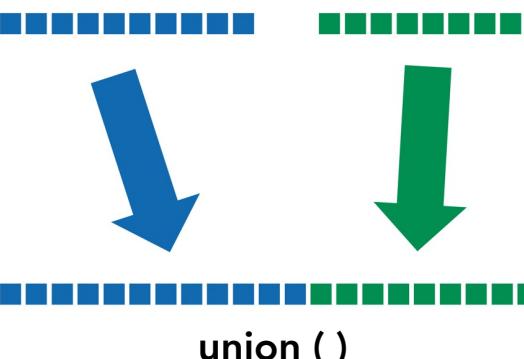
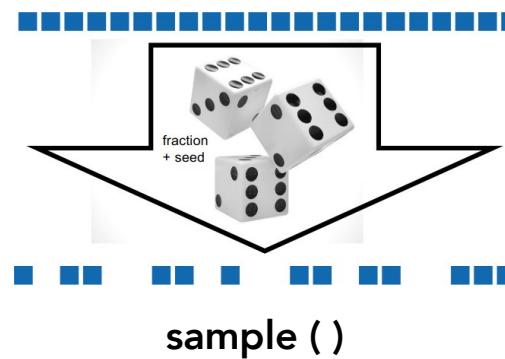
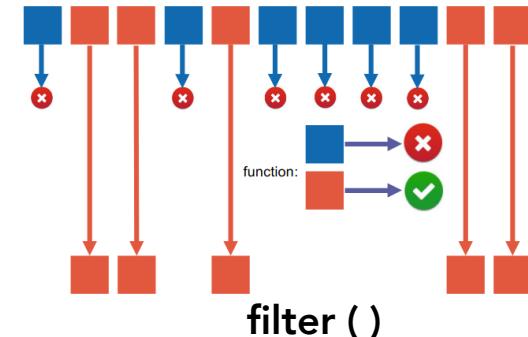
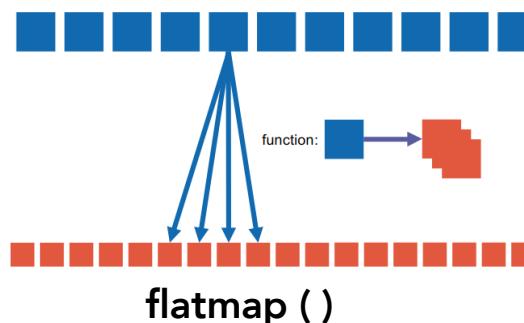
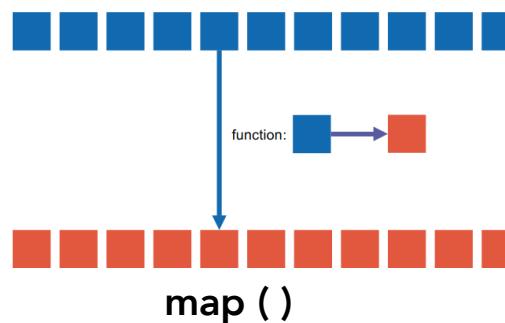
# Working with RDDs (I)

```
lines = sc.textFile("hdfs://data.txt")
```

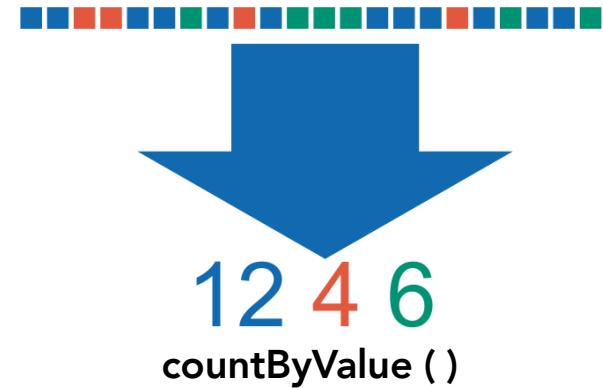
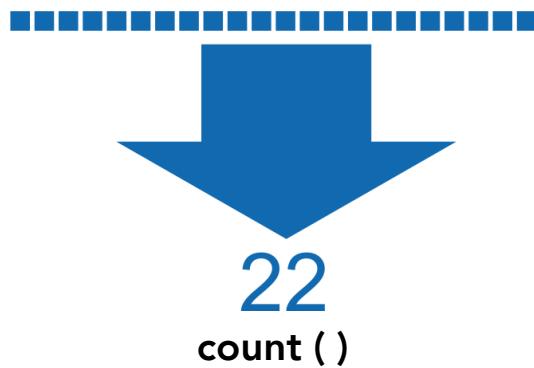
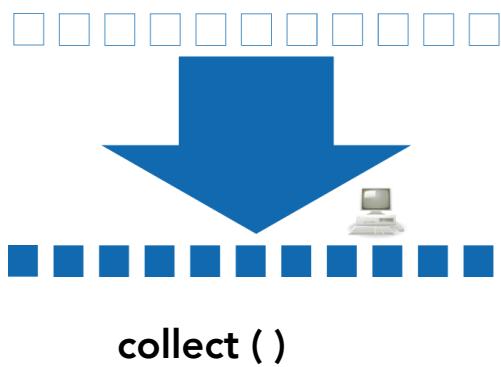


```
linesWithITU = lines.filter(line -> line.contains("ITU"))
```

# Example Transformations



# Example of Actions



# Spark Wordcount Example

```
text_file = sc.textFile("hdfs://...")  
counts = text_file.flatMap(lambda line: line.split(" ")) \\\n    .map(lambda word: (word, 1)) \\  
    .reduceByKey(lambda a, b: a + b)  
counts.saveAsTextFile("hdfs://...")
```

*.flatmap*

*.map*

*.reduceByKey*

# Apache Spark: Performance

2013 Record:  
Hadoop

72 minutes



2014 Record:  
Spark

23 minutes



Also sorted 1PB in 4 hours

<https://www.slideshare.net/rxin/stanford-cs347-guest-lecture-apache-spark>

# Spark SQL

- Structured data within Spark programs
  - DataSet/DataFrame as data representation
  - SQL as query language
  - Improved data parsing

```
import org.apache.spark.sql.types._

// Create an RDD
val peopleRDD = spark.sparkContext.textFile("examples/src/main/resources/people.txt")

// The schema is encoded in a string
val schemaString = "name age"

// Generate the schema based on the string of schema
val fields = schemaString.split(" ")
  .map(fieldName => StructField(fieldName, StringType, nullable = true))
val schema = StructType(fields)

// Convert records of the RDD (people) to Rows
val rowRDD = peopleRDD
  .map(_.split(","))
  .map(attributes => Row(attributes(0), attributes(1).trim))

// Apply the schema to the RDD
val peopleDF = spark.createDataFrame(rowRDD, schema)

// Creates a temporary view using the DataFrame
peopleDF.createOrReplaceTempView("people")

// SQL can be run over a temporary view created using DataFrames
val results = spark.sql("SELECT name FROM people")

// The results of SQL queries are DataFrames and support all the normal RDD operations
// The columns of a row in the result can be accessed by field index or by field name
results.map(attributes => "Name: " + attributes(0)).show()
// +-----+
// |    value|
// +-----+
// |Name: Michael|
// |  Name: Andy|
// | Name: Justin|
// +-----+
```

# Hadoop/Spark

## Strengths

- Good for distributed storage
- Good for complex distributed processing pipelines
  - Especially Spark!
- Excellent throughput!
  - Especially Spark!

## Weaknesses

- Not really data management systems
  - Very high latency of small requests!
- Only suitable for distributed processing pipelines!

# Take Aways

- Analytics applications typically require (repeated) processing of whole datasets of unstructured data
- This access pattern is not well served by SQL or NoSQL data systems
  - RDBMS: No need for transactions, and many analytics processes do not use SQL (but it can be useful for filtering)
  - NoSQL: Distributed storage is important, but individual records are not accessed
  - Both: Distributed (large-scale) processing is required
- MapReduce and (especially) Spark perform better in these applications