

Introduction to Database Systems

IDBS – Spring 2024

- Week 7:
- Storage Hierarchy
- Physical Database Design
- Indexing

Eleni Tzirita Zacharatou

Readings:


PDBM 12.1, 12.2, 12.3.1 - 12.3.7

A large, stylized silhouette of an elephant in a light purple color, facing right, occupies the right half of the slide. It has a thick trunk and large ears.

Storage Hierarchy

Physical Database
Design

Storage Hierarchy



Access Times

1 ns L1 Cache Ref	← 1 sec
4 ns L2 Cache Ref	← 4 sec
100 ns DRAM	← 100 sec
16,000 ns SSD	← 4.4 hours
2,000,000 ns HDD	← 3.3 weeks
~50,000,000 ns Network Storage	← 1.5 years
1,000,000,000 ns Tape Archives	← 31.7 years

Disk-Based Architecture

- The DBMS assumes that the primary storage location of the database is on non-volatile disk.
- The DBMS's components manage the movement of data between non-volatile and volatile storage.

Sequential vs. Random Access


- Random access on non-volatile storage is almost always much slower than sequential access.
- DBMS will want to maximize sequential access.
 - Algorithms try to reduce number of writes to random pages so that data is stored in contiguous blocks.

System Design Goals

- Allow the DBMS to manage databases that exceed the amount of memory available.
- Reading/writing to disk is expensive, so it must be managed carefully to avoid large stalls and performance degradation.
- Random access on disk is usually much slower than sequential access, so the DBMS will want to maximize sequential access.


Why Not Use the OS?

- The DBMS can use memory mapping (*mmap*) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.




Why Not Use the OS?

- The DBMS can use memory mapping (*mmap*) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



Why Not Use the OS?

- The DBMS can use memory mapping (*mmap*) to store the contents of a file into the address space of a program.
- OS is responsible for moving file pages in and out of memory, so the DBMS doesn't need to worry about it.



Memory Mapped I/O Problems

Problem #1: Transaction Safety

→ OS can flush dirty pages at any time.

Problem #2: I/O Stalls

→ DBMS doesn't know which pages are in memory. The OS will stall a thread on page fault.

Problem #3: Error Handling

→ Difficult to validate pages. Any access can cause a **SIGBUS** that the DBMS must handle.

Problem #4: Performance Issues

→ OS data structure contention. TLB shootdowns.

Why Not Use the OS?

- DBMS (almost) always wants to control things itself and can do a better job than the OS.
 - Flushing dirty pages to disk in the correct order.
 - Specialized prefetching.
 - Buffer replacement policy.
 - Thread/process scheduling.
- The OS is *not* your friend.

Storage Manager


- The DBMS stores a database as one or more files on disk typically in a proprietary format.
- The storage manager is responsible for maintaining the files of a database.
- It organizes the files as a collection of *pages*.
 - Tracks data read/written to pages.
 - Tracks the available space.

Database Pages

- A page is a fixed-size block of data.
 - It can contain tuples, meta-data, indexes, log records...
 - Most systems do not mix page types.
 - Some systems require a page to be self-contained.
- Each page is given a unique identifier.
 - The DBMS uses an indirection layer to map page IDs to physical locations.



Indexing



Full Table Scans

When a DBMS sees a query of the form

```
SELECT *  
FROM R  
WHERE <condition>
```

one obviously possible thing to do is read through **all** the tuples of R and report those tuples that satisfy the condition.

This is called a **full table scan**.

Selective Queries

Consider the query from before:

- If we have to report 80% of the tuples in R, it makes sense to do a full table scan.
- On the other hand, if the query is very **selective**, and returns just a small percentage of the tuples, we might hope to do better.

```
SELECT *  
FROM R  
WHERE <condition>
```

Point Queries

- Consider a selection query with a single equality condition:

```
SELECT *  
FROM person  
WHERE birthdate='1975-02-06';
```

- This is a *point query*: We look for a single value of birthdate.
 - We may still return > 1 record!
- Point queries are easy if data is sorted by the attribute used in the condition.
 - How? What algorithms would work?

Range Queries

- Consider a selection query of the form:

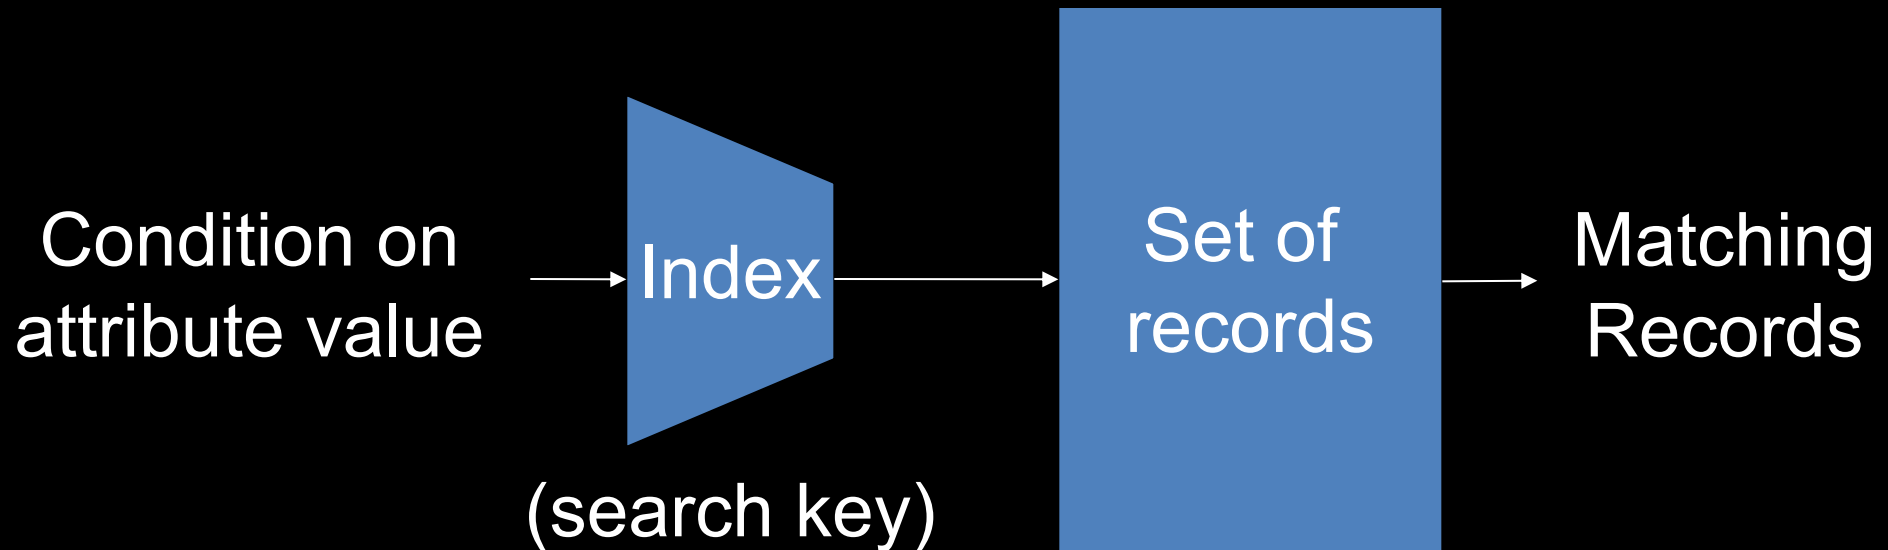
```
SELECT *  
FROM person  
WHERE birthdate BETWEEN  
  '1975-02-01' and '1975-02-28';
```
- This is a *range query*: We look for a range of values of birthdate.
- Range queries are also easy if data is sorted by the right attribute.
 - But often not as **selective** as point queries.

Indexes

- To speed up queries the DBMS may build an index on the `birthdate` attribute.
- A database index is similar to an index in the back of a book:
 - For every piece of data you might be interested in (e.g., the attribute value `1975-02-06`), the index says where to find the row with the actual data!
 - The index itself is organized such that one can quickly do the lookup.
- Looking for information in a relation with the help of an index is called an **index scan** (range) or **index lookup** (point).

Indexing

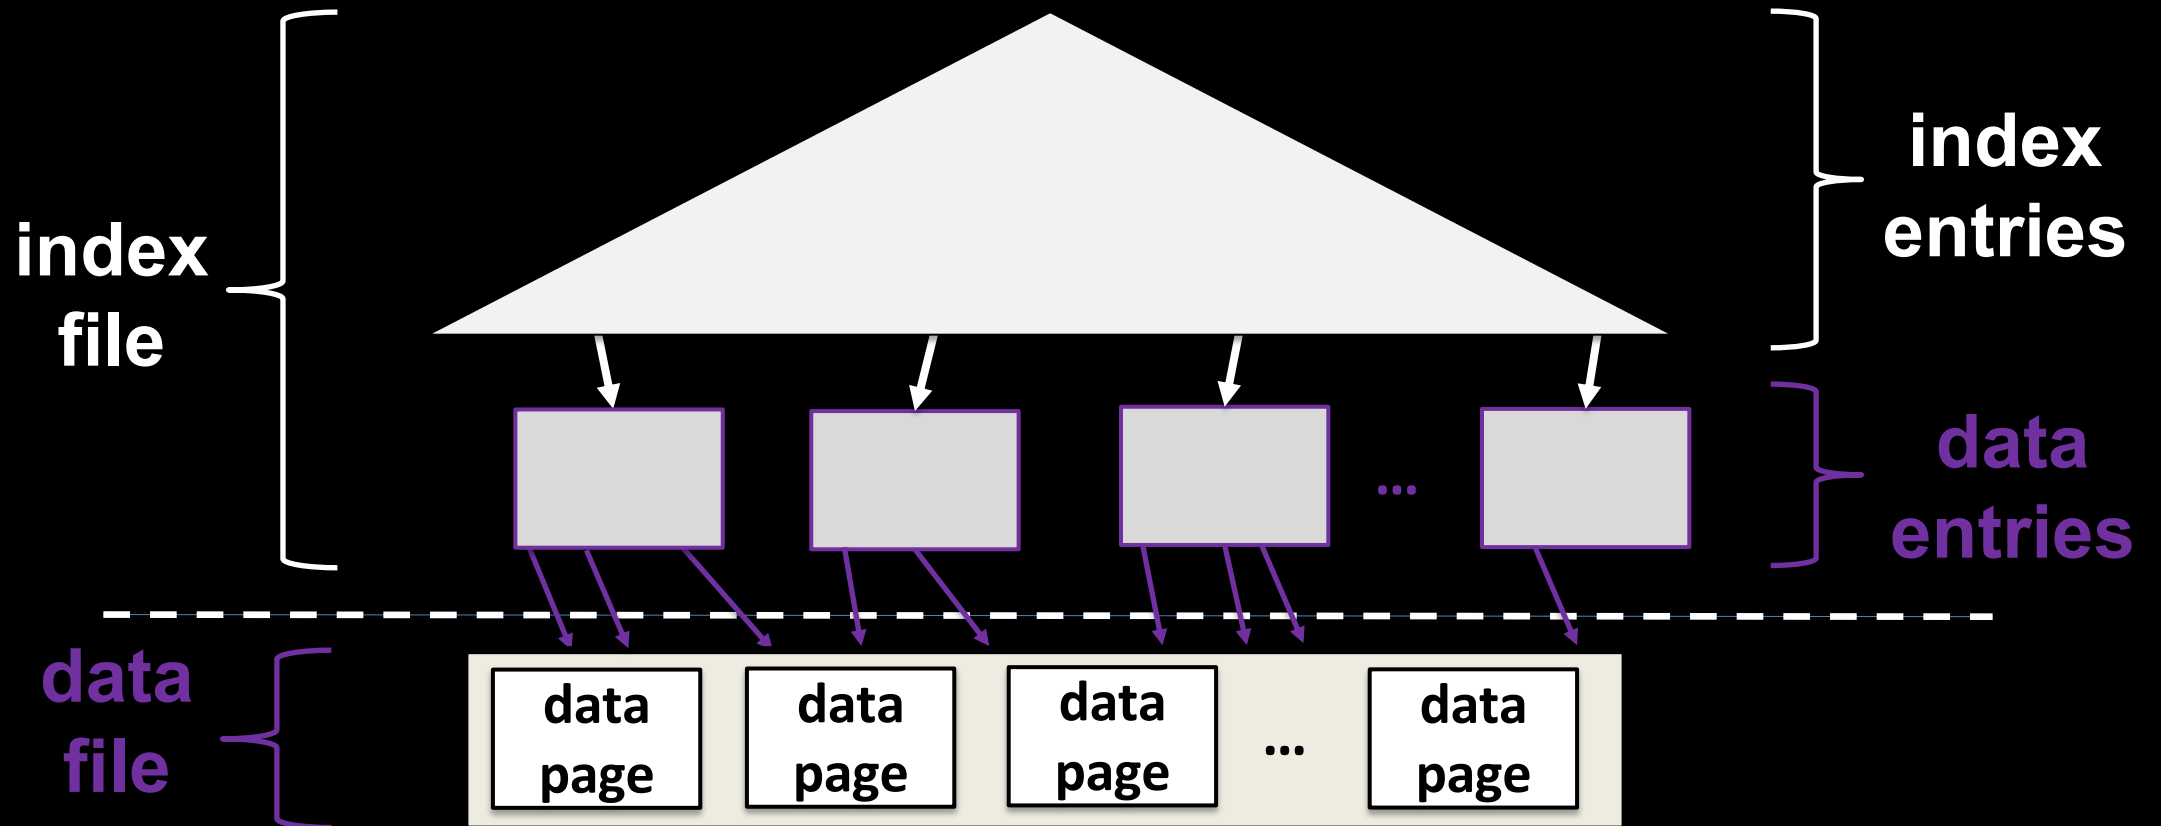
- An index is a data structure that supports efficient access to data
 - In databases, indexes are also stored on disk



Two Techniques

- **Hashing:** Use a fixed transformation algorithm to convert the attribute value into a database address.
- **Tree search:** A dynamic search structure is built that guides the search for a given attribute value to the proper database location.
- Hashing supports equality queries only
 - Typically used dynamically for large-scale joins
 - Rarely available to developers
- Tree search is more versatile and accessible

Clustered Index

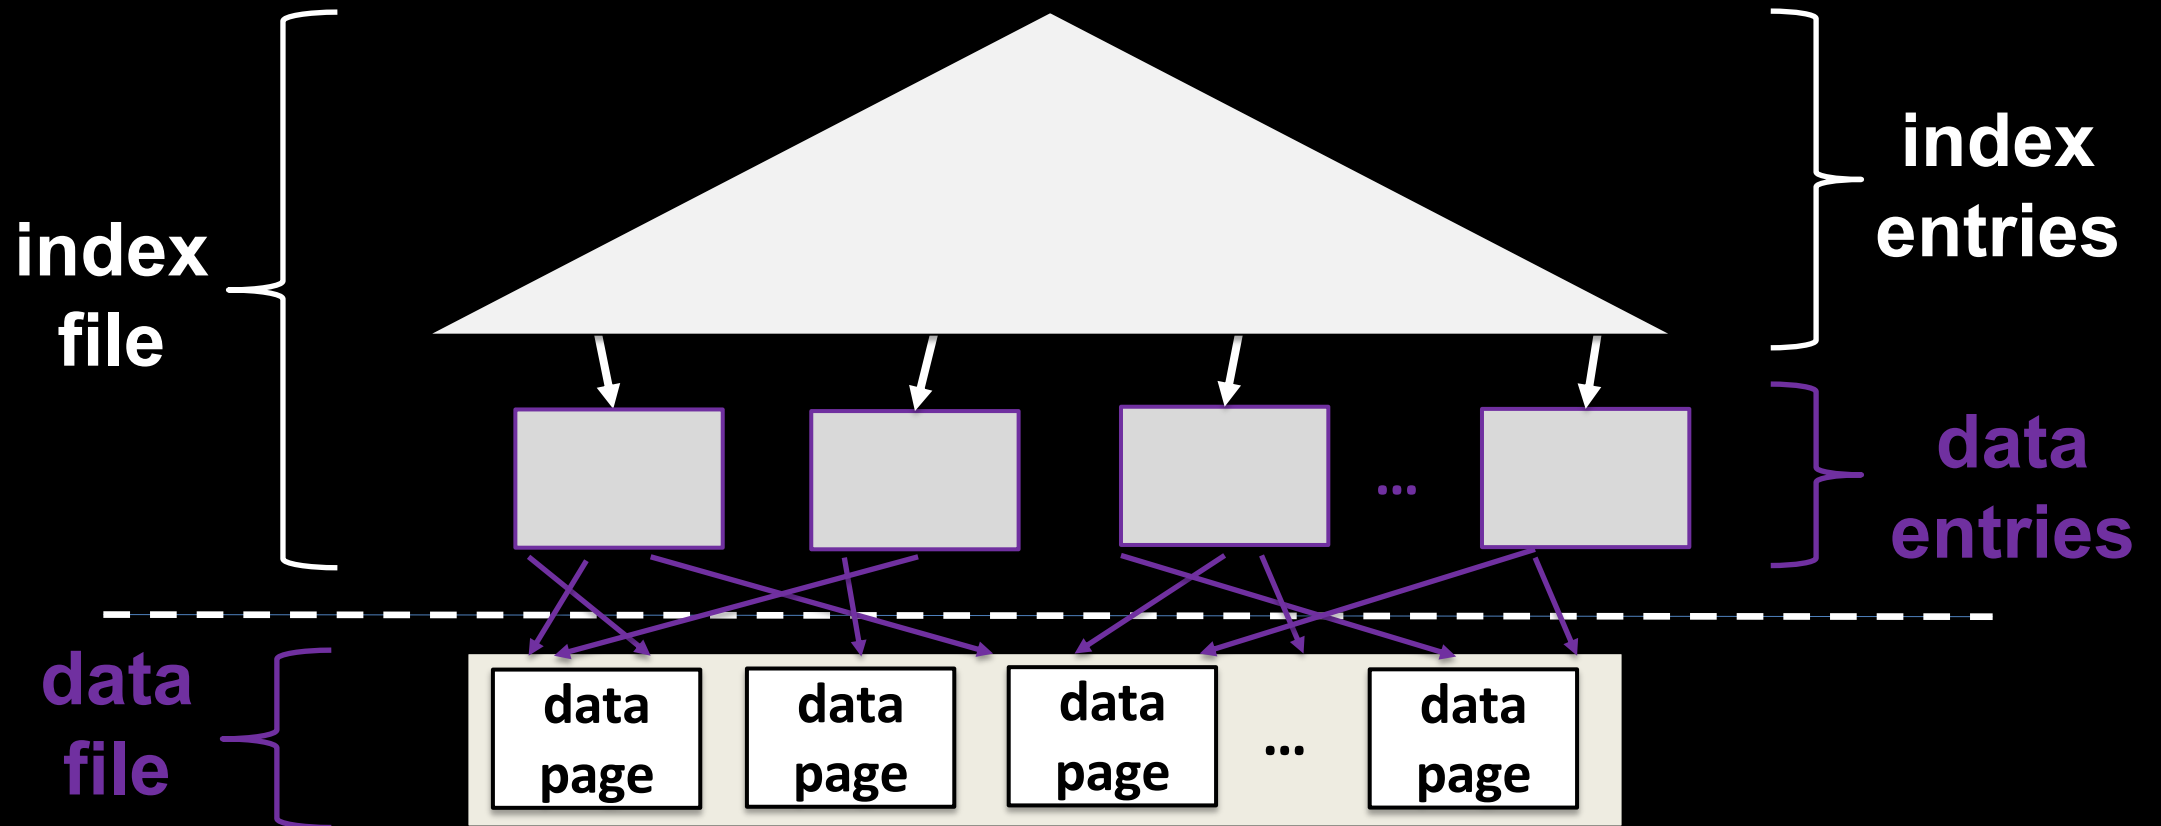


order of data entries matches order of data records

Clustered Index

- If the tuples of a relation are stored sorted according to some attribute, an index on this attribute is called clustered.
 - Clustered indexes make point and range queries on the key *very efficient*
 - Why? Sequential reads + As few reads as possible!
- Many DBMSs automatically build a clustered index on the primary key of each relation.
- A clustered index is sometimes referred to as a primary or sparse index.
 - There can only be one clustered index! Why?

Unclustered Index



order of data entries != order of data records

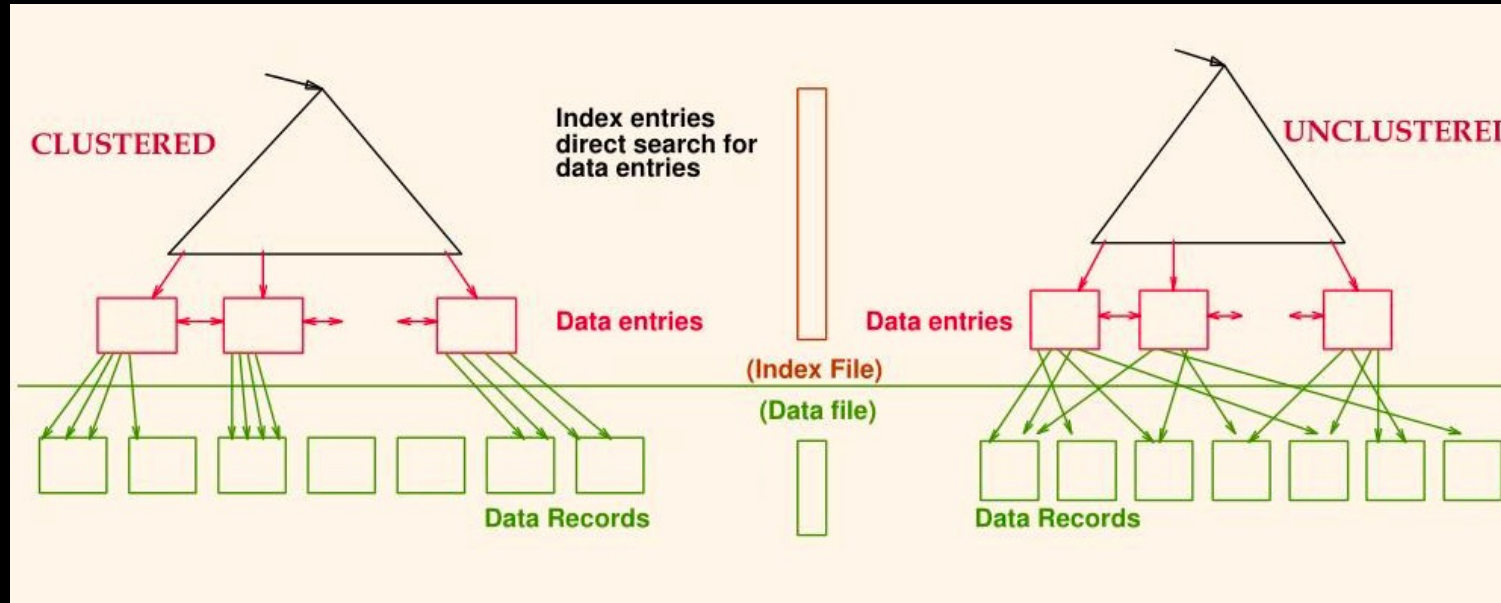
Unclustered Index

- It is possible to create further indexes on a relation. Typical syntax:

```
CREATE INDEX myIndex ON involved(actorId);
```

- The unclustered indexes are sometimes called non-clustered, secondary or dense indexes.
- Unclustered indexes:
 - Make most point queries more efficient.
 - Make some (narrow) range queries more efficient.

Clustered vs. Unclustered Index




- To retrieve M records, where M is small:
 - Clustered: Probably one disk read
 - Unclustered: Probably M random disk reads
- To retrieve M records, where M is large:
 - Clustered: Probably $M/\text{records_per_page}$ sequential disk reads
 - Unclustered: Up to M random disk reads
- We still need to read the index itself – same for both!


Index Scan vs. Full Table Scan

- Point and range queries on the attribute(s) of the clustered index are almost always best performed using an index scan.
- Non-clustered indexes should only be used with high selectivity queries.
- Rule of thumb?

Impact of Clustering on Performance



Covering Index



data entries contain ALL queried attributes

Covering Index

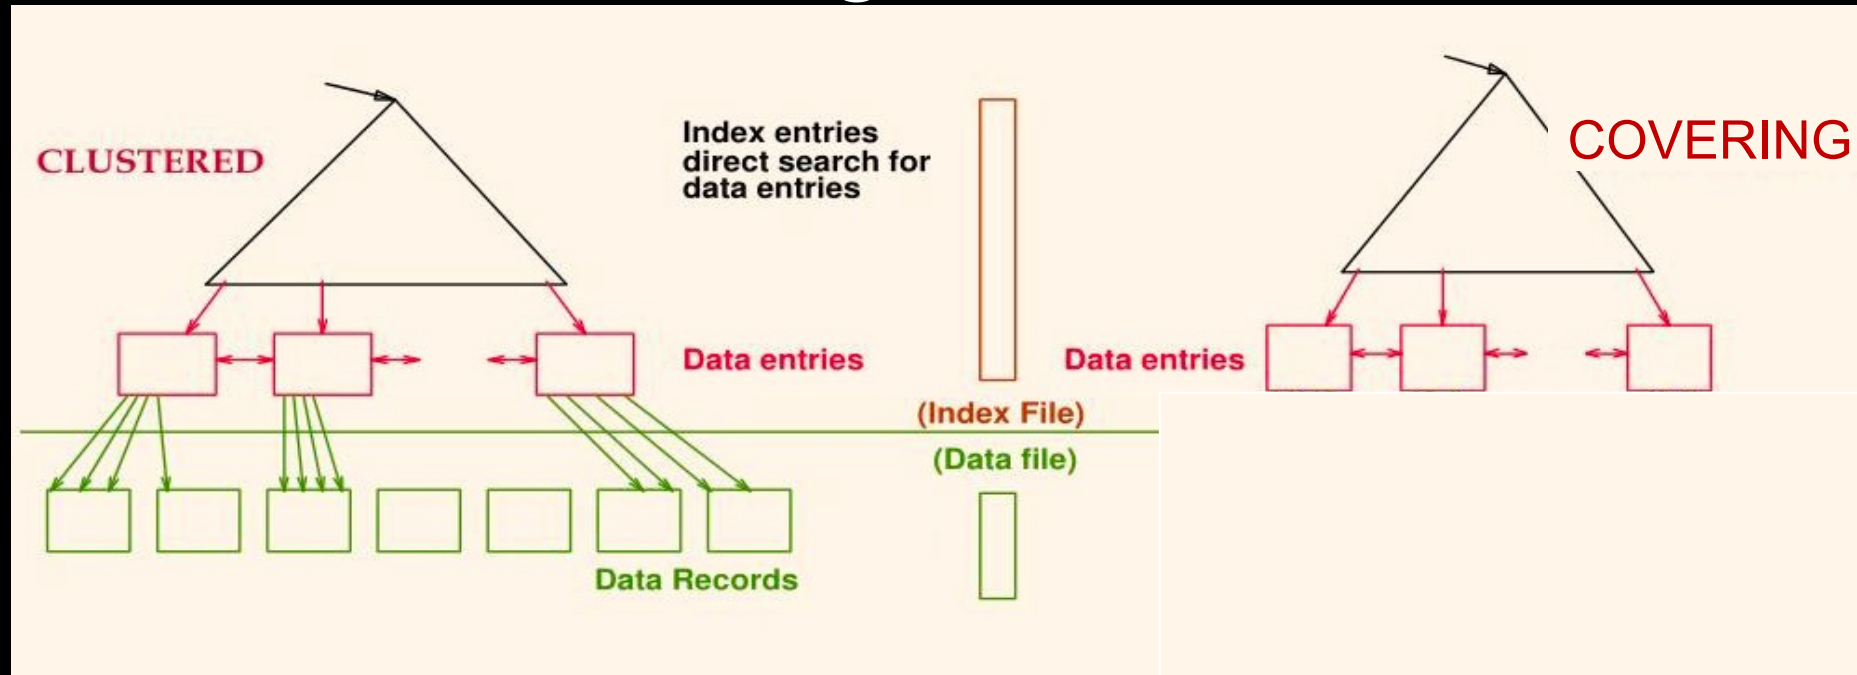
- An index that contains ALL attributes used in a query is called covering
 - The resulting query plans are index-only

```
SELECT COUNT(*) FROM movie WHERE year=1948;  
CREATE INDEX movieyear ON movie(year);
```

```
SELECT name FROM person WHERE height=170;  
CREATE INDEX phn ON person(height, name);
```

- The data from the relation is not needed = no disk reads required to retrieve tuples
 - Should a covering index be clustered or non-clustered?

Clustered vs. Covering Index



- To retrieve M records, where M is small:
 - Clustered: Probably one disk read
 - Covering: Definitely need 0 disk reads
- To retrieve M records, where M is large:
 - Clustered: Probably $M/records_per_page$ sequential disk reads
 - Covering: Definitely need 0 disk reads
- We still need to read the index itself – same for both!

Problem Session

- For each of the following queries:
 - Which index would give the best plan?
 - Would the index be covering?
 - Would you prefer clustered or unclustered index?
 - Based on these queries, if you could choose, which index should be clustered?

1. SELECT * ... WHERE birthdate = '20-02-2002'

2. SELECT * ... WHERE height < 170

3. SELECT * ... WHERE ID = 4564

4. SELECT AVG(birthdate) ...

```
CREATE INDEX bIndex
ON person(birthdate);
CREATE INDEX hIndex
ON person(height);
CREATE INDEX iIndex
ON person(ID);
```

Multi-Attribute Indexes

- Defining an index on several attributes:

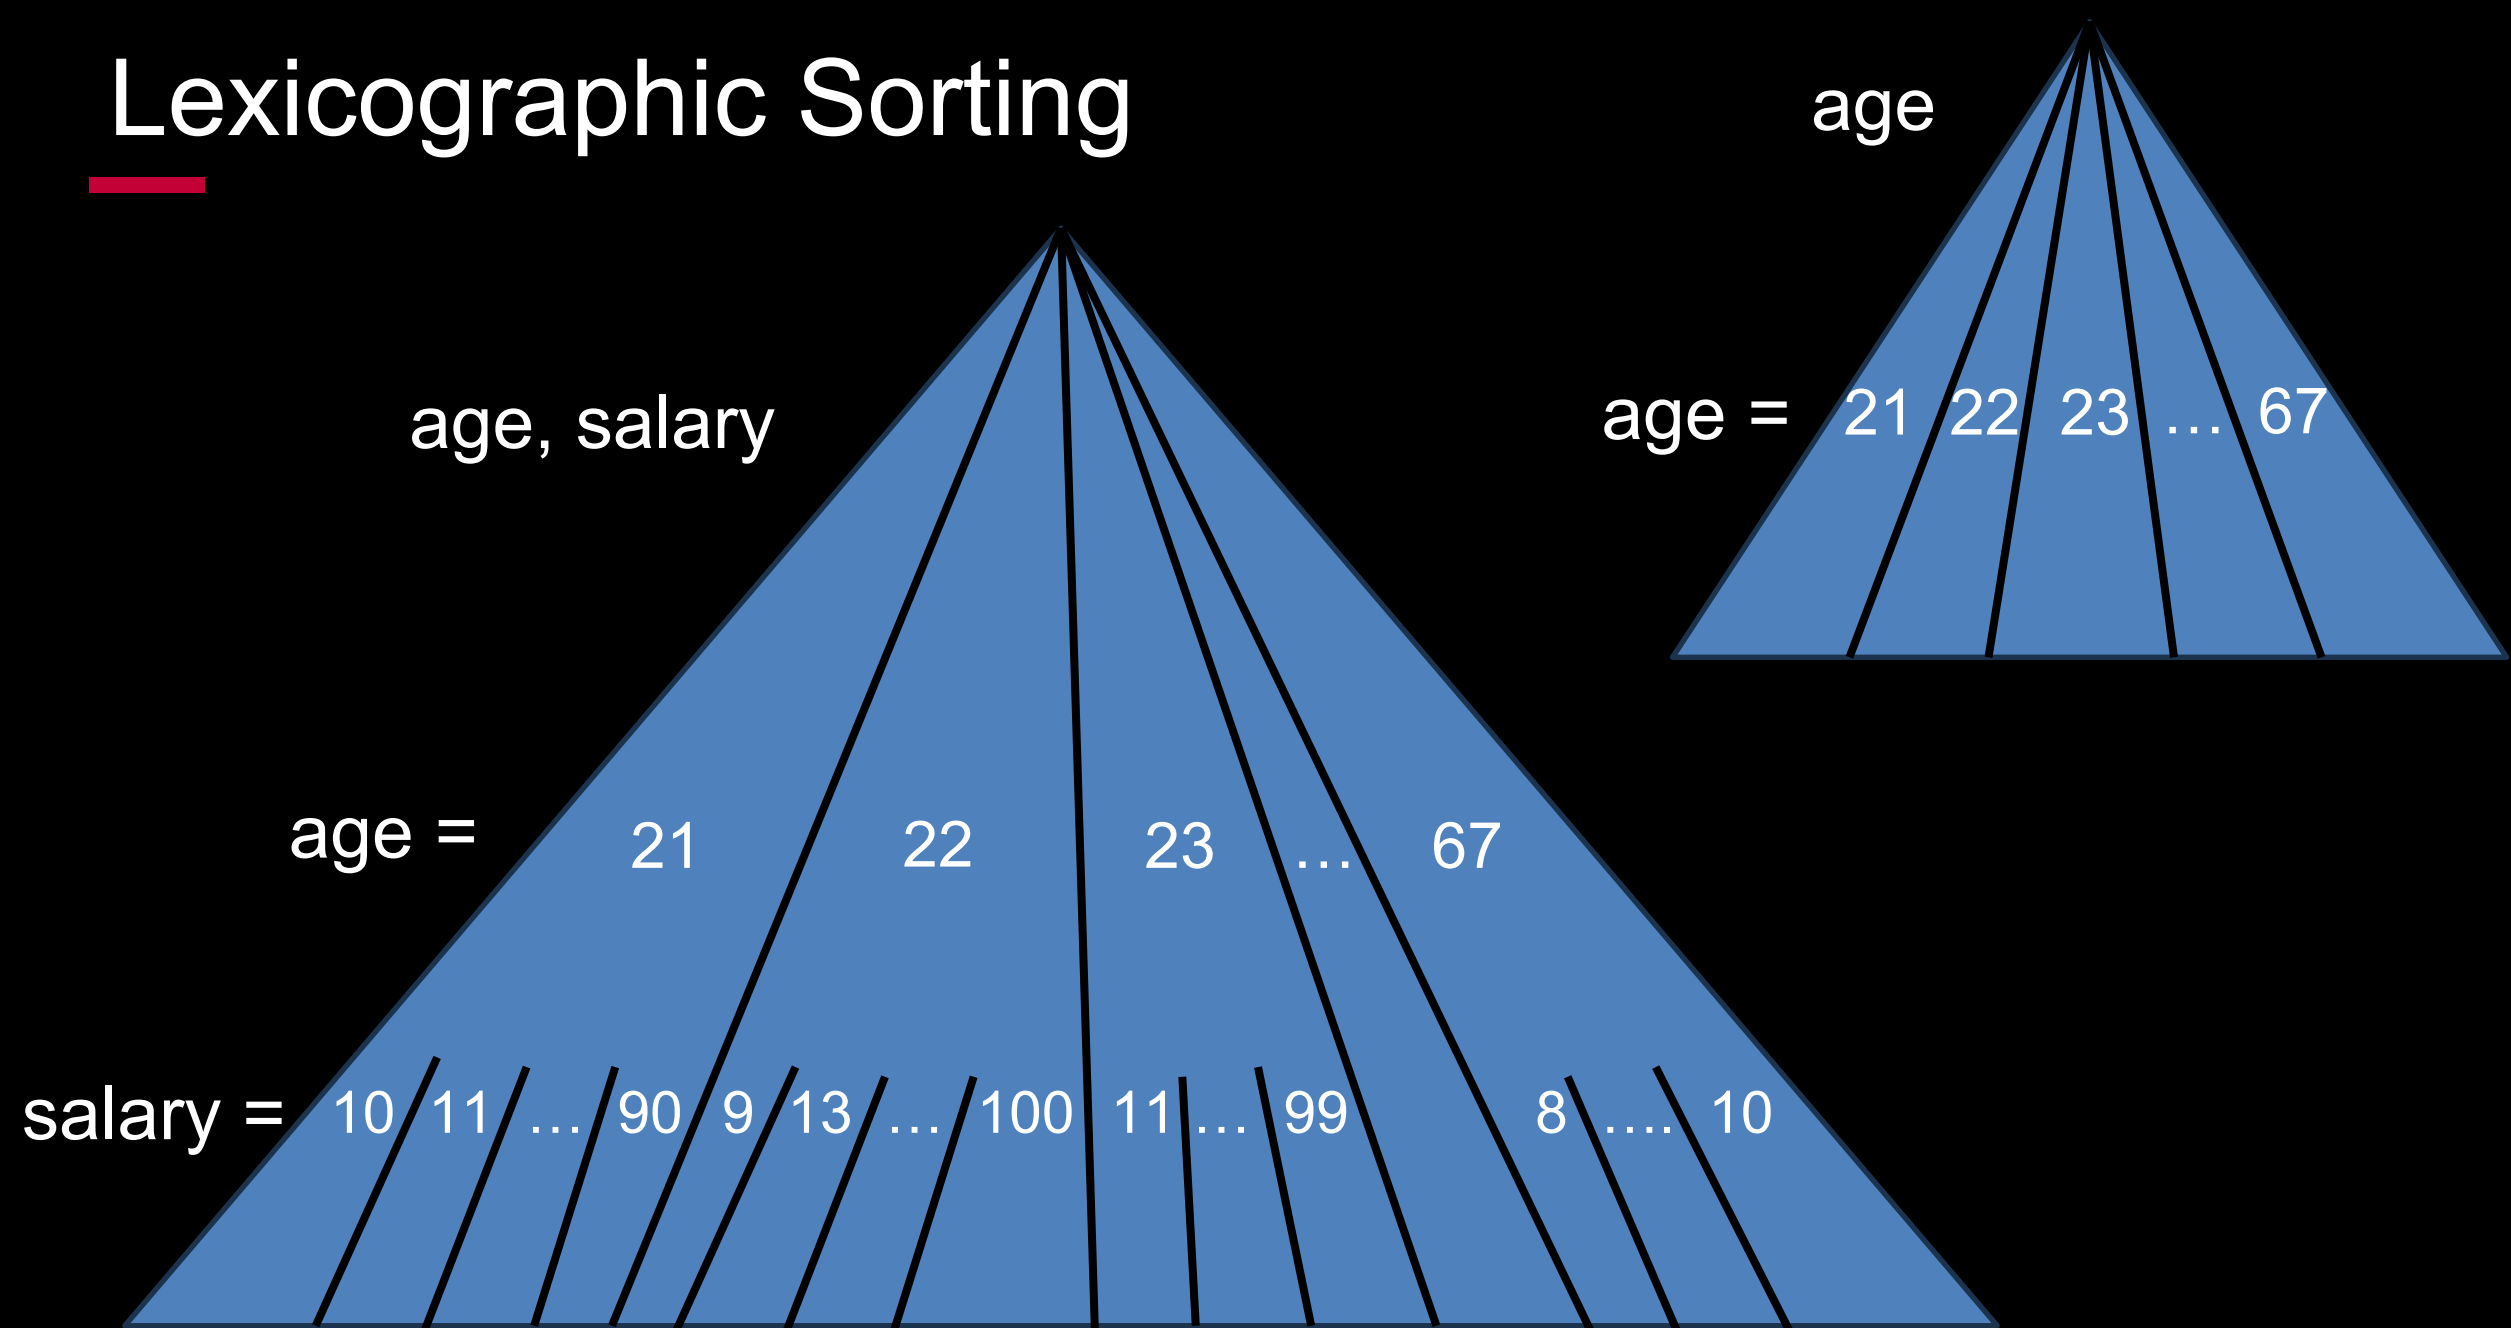
```
CREATE INDEX myIndex  
ON person (height,birthdate);
```

- Speeds up point/range queries such as:

```
SELECT *  
FROM person  
WHERE height=170 and birthdate<'1945-08-08'
```

- An index on several attributes usually gives index for any **prefix** of these attributes, due to lexicographic sorting.

Lexicographic Sorting



Problem Session

- Which point and range queries are "easy" (equality on a prefix, range on one) when the relation is indexed with this two-attribute index?

```
CREATE INDEX myIndex  
ON person (height, birthdate);
```

1. A range query on height?
2. A range query on birthdate?
3. A point query on birthdate?
4. A point query on birthdate combined with a range query on height?
5. A point query on height combined with a range query on birthdate?

Choosing to Use an Index

- The choice of whether to use an index is made by the DBMS *for every instance of a query*
 - May depend on query parameters
 - You do not need to take indexes into account when writing queries
- Estimating selectivity is done using statistics
 - In PostgreSQL, statistics are collected by executing statements such as `ANALYZE table_name;`

Choosing Columns

- Candidates for index search keys
 - Columns in WHERE clauses
 - Columns in GROUP BY clauses
 - Columns in ORDER BY clauses
- Columns that are rarely candidates
 - Large columns (too much space)
 - Frequently updated columns (too much maintenance)
 - Columns in SELECT clauses (not used to find tuples)
 - ... but see **covering indices**!

What Speaks Against Indexing

- Space usage:
 - Similar to size of indexed columns (plus pointer)
 - Most space for leaves, less for tree nodes
 - Not really important!
- Time usage for keeping indexes updated under data insertion/change:
 - Depends on the index architecture
 - This IS important!

Take Away

- Storage hierarchy:
 - Different layers have different characteristics.
 - Sequential access is faster than random access for all layers, but especially for hard disks.
 - DBMSs have various optimizations to minimize data access latency.
- Databases are organized in pages.
- Large databases need to be equipped with suitable indexes
 - Need understanding of what indexes might help a given set of queries.
 - *Key distinction*: Clustered vs. Unclustered.
 - A detailed understanding of various index types is beyond the scope of this course.